

FIT5149 S2 2019 Assessment 1: Predicting the Critical Temperature of a Superconductor

Student information

- Family Name: Vydyula
- Given Name: Shyama Malhari
- Student ID: 29860644
- Student email: svyd0001@student.monash.edu

Programming Language: R 3.5.1 in Jupyter Notebook

R Libraries used:

- psych
- ggplot2
- purrr
- tidyr
- ggcorrplot
- stringr
- dplyr
- data.table
- xgboost
- readr
- stringr
- caret
- usdm
- fmsb
- VIF
- ggpubr

Table of Contents

- [Introduction](#)
- [Data Exploration](#)
- [Model Development](#)
- [Model Comparison](#)
- [Variable Identification and Explanation](#)
- [Conclusion](#)
- [References](#)

In [1]:

```
#install.packages('caret')
#install.packages('ggplot2')
#install.packages('usdm')
#install.packages('fmsb')
#install.packages('VIF')
#install.packages('psych')
#install.packages('ggpubr')
#install.packages('leaps')
```

In [96]:

```
library(tidyverse)
library(caret)
library(leaps)
library(psych)
library(ggplot2)
library(purrr)
library(tidyr)
library(ggcorrplot)
library(stringr)
library(dplyr)
library(data.table)
library(xgboost)
library(readr)
library(stringr)
library(caret)
library(usdm)
library(fmsb)
library(VIF)
library(ggpubr)
```

1. Introduction

According to Wikipedia:

"Superconductivity is the set of physical properties observed in certain materials, wherein electrical resistance vanishes and from which magnetic flux fields are expelled. Any material exhibiting these properties is a superconductor. Unlike an ordinary metallic conductor, whose resistance decreases gradually as its temperature is lowered even down to near absolute zero, a superconductor has a characteristic critical temperature below which the resistance drops abruptly to zero. An electric current through a loop of superconducting wire can persist indefinitely with no power source"

Superconductors have immense applications in the current world such as:

- Used in Magnetic Resonanc Imaging or MRI, which is used in the healthcare industry for physical examination of the human body.
- Used in Particle Acelerators, machines that use EM field to propel charges, at CERN.
- Used in Superconducting Quantum Interference Devices or (SQUIDs)

Since the conductance of a superconductor depends almost entirely on the ability to identify its 'Critical Temperature', or the temperature at which its resistance drops to zero, the need to identify and measure this temperature using the various chemical properties has become a topic of interest for physicians around the world.

In this project, we will be conducting a statistical analysis to predict the 'Critical Temperature' of a superconductor, given some of its chemical properties.

The project is entirely based on : "**A Data Driven Statistical Model for Predicting the Critical Temperature of a Superconductor** at [https://arxiv \(https://arxiv\). org/pdf/1803.10260.pdf](https://arxiv.org/pdf/1803.10260.pdf)"

2. Data Exploration

"However the model builder's main problem is often not how to fit a model but rather how to formulate it in the first place".

Chris Chatfield has rightly said in his publication, 'Exploratory Data Analysis' that is wise to begin a statistical analysis with an informal, exploratory examination of the data called exploratory data analysis (abbreviated EDA). The two main objectives are: data description and model-formulation. It is suggested that it is important to see EDA as an integral part of statistical inference. (Chatfield, 1986)

Loading the dataset:

In [12]:

```
#Reading the Dataset using "read" function in R.  
super <- read.csv('train.csv',header = TRUE, sep = ',')
```

In [13]:

```
#Dimensions of the Data  
dim(super)
```

21263 82

It can be seen from the output that there are 82 columns or variables in this dataset and 21,263 observations.

In [14]:

```
#Structure of the Data  
str(super)
```

```
'data.frame': 21263 obs. of 82 variables:  
 $ number_of_elements : int 4 5 4 4 4 4 4 4 4 4 ...  
 $ mean_atomic_mass : num 88.9 92.7 88.9 88.9 88.9  
 ...  
 $ wtd_mean_atomic_mass : num 57.9 58.5 57.9 57.9 57.8  
 ...
```

```

$ gmean_atomic_mass      : num  66.4  73.1  66.4  66.4  66.4
...
$ wtd_gmean_atomic_mass  : num   36.1  36.4  36.1  36.1  36.1
...
$ entropy_atomic_mass    : num   1.18  1.45  1.18  1.18  1.18
...
$ wtd_entropy_atomic_mass : num   1.062  1.058  0.976  1.022  1
.129 ...
$ range_atomic_mass      : num   123  123  123  123  123  ...
$ wtd_range_atomic_mass  : num   31.8  36.2  35.7  33.8  27.8
...
$ std_atomic_mass        : num    52  47.1  52  52  52  ...
$ wtd_std_atomic_mass    : num   53.6  54  53.7  53.6  53.6  ..
.
$ mean_fie               : num   775  766  775  775  775  ...
$ wtd_mean_fie           : num  1010 1011 1011 1011 1010
...
$ gmean_fie              : num   718  721  718  718  718  ...
$ wtd_gmean_fie          : num   938  939  939  939  937  ...
$ entropy_fie            : num   1.31  1.54  1.31  1.31  1.31
...
$ wtd_entropy_fie        : num   0.791 0.807 0.774 0.783 0
.805 ...
$ range_fie              : num   811  811  811  811  811  ...
$ wtd_range_fie          : num   736  743  743  740  729  ...
$ std_fie                : num   324  290  324  324  324  ...
$ wtd_std_fie            : num   356  355  355  355  356  ...
$ mean_atomic_radius      : num   160  161  160  160  160  ...
$ wtd_mean_atomic_radius  : num   106  105  105  105  106  ...
$ gmean_atomic_radius     : num   136  141  136  136  136  ...
$ wtd_gmean_atomic_radius : num   84.5 84.4 84.2 84.4 84.8
...
$ entropy_atomic_radius   : num   1.26  1.51  1.26  1.26  1.26
...
$ wtd_entropy_atomic_radius : num   1.21  1.2  1.13  1.17  1.26  .
..
$ range_atomic_radius     : int   205  205  205  205  205  205  2
05 171 171 171 ...
$ wtd_range_atomic_radius : num   42.9 50.6 49.3 46.1 36.5
...
$ std_atomic_radius       : num   75.2 67.3 75.2 75.2 75.2
...
$ wtd_std_atomic_radius   : num   69.2 68 67.8 68.5 70.6  ..
.
$ mean_Density            : num  4654 5821 4654 4654 4654
...
$ wtd_mean_Density        : num  2962 3021 2999 2980 2924
...
$ gmean_Density           : num   725 1237 725 725 725  ...
$ wtd_gmean_Density       : num   53.5 54.1 54 53.8 53.1  ..
.
$ entropy_Density         : num   1.03  1.31  1.03  1.03  1.03
...
$ wtd_entropy_Density     : num   0.815 0.915 0.76 0.789 0.
86 ...
$ range_Density           : num  8959 10489 8959 8959 8959
...

```

```

$ wtd_range_Density          : num  1580 1667 1667 1623 1492
...
$ std_Density                : num   3306 3767 3306 3306 3306
...
$ wtd_std_Density            : num   3573 3633 3592 3582 3553
...
$ mean_ElectronAffinity      : num    81.8 90.9 81.8 81.8 81.8
...
$ wtd_mean_ElectronAffinity  : num    112 112 112 112 111 ...
$ gmean_ElectronAffinity     : num    60.1 69.8 60.1 60.1 60.1
...
$ wtd_gmean_ElectronAffinity : num    99.4 101.2 101.1 100.2 97
.8 ...
$ entropy_ElectronAffinity   : num    1.16 1.43 1.16 1.16 1.16
...
$ wtd_entropy_ElectronAffinity : num    0.787 0.839 0.786 0.787 0
.787 ...
$ range_ElectronAffinity     : num    127 127 127 127 127 ...
$ wtd_range_ElectronAffinity : num    81 81.2 81.2 81.1 80.8 ..
.
$ std_ElectronAffinity       : num    51.4 49.4 51.4 51.4 51.4
...
$ wtd_std_ElectronAffinity   : num    42.6 41.7 41.6 42.1 43.5
...
$ mean_FusionHeat            : num    6.91 7.78 6.91 6.91 6.91
...
$ wtd_mean_FusionHeat        : num    3.85 3.8 3.82 3.83 3.87 .
..
$ gmean_FusionHeat           : num    3.48 4.4 3.48 3.48 3.48 .
..
$ wtd_gmean_FusionHeat       : num    1.04 1.04 1.04 1.04 1.04
...
$ entropy_FusionHeat         : num    1.09 1.37 1.09 1.09 1.09
...
$ wtd_entropy_FusionHeat     : num    0.995 1.073 0.927 0.964 1
.045 ...
$ range_FusionHeat           : num    12.9 12.9 12.9 12.9 12.9
...
$ wtd_range_FusionHeat       : num    1.74 1.6 1.76 1.74 1.74 .
..
$ std_FusionHeat             : num    4.6 4.47 4.6 4.6 4.6 ...
$ wtd_std_FusionHeat         : num    4.67 4.6 4.65 4.66 4.68 .
..
$ mean_ThermalConductivity   : num    108 172 108 108 108 ...
$ wtd_mean_ThermalConductivity : num    61 61.4 60.9 61 61.1 ...
$ gmean_ThermalConductivity  : num    7.06 16.06 7.06 7.06 7.06
...
$ wtd_gmean_ThermalConductivity : num    0.622 0.62 0.619 0.621 0.
625 ...
$ entropy_ThermalConductivity : num    0.308 0.847 0.308 0.308 0
.308 ...
$ wtd_entropy_ThermalConductivity : num    0.263 0.568 0.25 0.257 0.
273 ...
$ range_ThermalConductivity  : num    400 430 400 400 400 ...
$ wtd_range_ThermalConductivity : num    57.1 51.4 57.1 57.1 57.1
...
$ std_ThermalConductivity    : num    169 199 169 169 169 ...

```

```

$ wtd_std_ThermalConductivity : num 139 140 139 139 138 ...
$ mean_Valence : num 2.25 2 2.25 2.25 2.25 2.2
5 2.25 2.25 2.25 2.25 ...
$ wtd_mean_Valence : num 2.26 2.26 2.27 2.26 2.24
...
$ gmean_Valence : num 2.21 1.89 2.21 2.21 2.21
...
$ wtd_gmean_Valence : num 2.22 2.21 2.23 2.23 2.21
...
$ entropy_Valence : num 1.37 1.56 1.37 1.37 1.37
...
$ wtd_entropy_Valence : num 1.07 1.05 1.03 1.05 1.1 .
..
$ range_Valence : int 1 2 1 1 1 1 1 1 1 1 ...
$ wtd_range_Valence : num 1.09 1.13 1.11 1.1 1.06 .
..
$ std_Valence : num 0.433 0.632 0.433 0.433 0
.433 ...
$ wtd_std_Valence : num 0.437 0.469 0.445 0.441 0
.429 ...
$ critical_temp : num 29 26 19 22 23 23 11 33 3
6 31 ...

```

From the above output, it is clear that all the columns that are in the dataset are numerical in type. The first column tells us how many elements are present in the super conductor and the last column is the critical temperature of that super conductor. The rest of the 80 columns are derived from the 8 basic properties of the super conductor as mentioned in Hamidieh, K. (2018), the foundation reference for this Assignment.

The 8 properties can be described as follows:

Data Description:

- 1. Atomic Mass** : It is the total proton and neutron rest masses measured in Atomic Mass Units (AMU).
- 2. First Ionization Energy** : It is the energy required to remove a valence electron from an element. Measured in kilo-Joules per mole (kJ/mol).
- 3. Atomic Radius** : It is the calculated atomic radius of the element. Measured in picometer (pm).
- 4. Density** : It is the density at standard temperature and pressure. Measured in kilograms per meters cubed (kg/m³).
- 5. Electron Affinity** : The energy required to add an electron to a neutral atom. Measured in kilo-Joules per mole (kJ/mol).
- 6. Fusion Heat** : The energy to change from solid to liquid state without temperature change. Measured in kilo-Joules per mole (kJ/mol).
- 7. Thermal Conductivity** : The thermal conductivity coefficient measured in watts per meter-Kelvin (W/(m x K)).
- 8. Valence** : The number of chemical bonds formed by the elements. It is a numerical value and doesn't have a unit of measurement.

In [15]:

```
#Setting the kernel to allow wide datasets to be displayed without breaks
options(repr.matrix.max.cols=1000, repr.matrix.max.rows=2000)
```

Renaming the columns for lucidity:

In [16]:

```
super1 <- super
```

In [17]:

```
super1 <- super1 %>% rename_at(vars(starts_with("mean_")), funs(str_replace(.,  
"mean_", "m_"))) #mean  
super1 <- super1 %>% rename_at(vars(starts_with("wtd_mean_")), funs(str_replac  
e(., "wtd_mean_", "wm_"))) #weighted mean  
super1 <- super1 %>% rename_at(vars(starts_with("gmean_")), funs(str_replace(.,  
, "gmean_", "gm_"))) #geometric mean  
super1 <- super1 %>% rename_at(vars(starts_with("wtd_gmean_")), funs(str_repla  
ce(., "wtd_gmean_", "wgm_"))) #weighted geometric mean  
super1 <- super1 %>% rename_at(vars(starts_with("entropy_")), funs(str_replace  
(., "entropy_", "e_"))) #entropy  
super1 <- super1 %>% rename_at(vars(starts_with("wtd_entropy_")), funs(str_rep  
lace(., "wtd_entropy_", "we_"))) #weighted entropy  
super1 <- super1 %>% rename_at(vars(starts_with("range_")), funs(str_replace(.,  
, "range_", "r_"))) #range  
super1 <- super1 %>% rename_at(vars(starts_with("wtd_range_")), funs(str_repla  
ce(., "wtd_range_", "wr_"))) #weighted range  
super1 <- super1 %>% rename_at(vars(starts_with("wtd_std_")), funs(str_replace  
(., "wtd_std_", "wstd_"))) #weighted standard deviation
```

Warning message:

"`is_lang()` is deprecated as of rlang 0.2.0.

Please use `is_call()` instead.

This warning is displayed once per session."Warning message:

"`lang_modify()` is deprecated as of rlang 0.2.0.

Please use `call_modify()` instead.

This warning is displayed once per session."

In [18]:

```
super1 <- super1 %>% rename_at(vars(contains("atomic_mass")), funs(str_replace  
(., "atomic_mass", "am"))) #Atomic Mass  
super1 <- super1 %>% rename_at(vars(contains("atomic_radius")), funs(str_repla  
ce(., "atomic_radius", "ar"))) #Atomic Radius  
super1 <- super1 %>% rename_at(vars(contains("Density")), funs(str_replace(.,  
"Density", "d"))) #Density  
super1 <- super1 %>% rename_at(vars(contains("ElectronAffinity")), funs(str_re  
place(., "ElectronAffinity", "ea"))) #Electron affinity  
super1 <- super1 %>% rename_at(vars(contains("FusionHeat")), funs(str_replace(  
., "FusionHeat", "fh"))) #Fusion Heat  
super1 <- super1 %>% rename_at(vars(contains("ThermalConductivity")), funs(str  
_replace(., "ThermalConductivity", "tc"))) #Thermal Conductivity  
super1 <- super1 %>% rename_at(vars(contains("Valence")), funs(str_replace(.,  
"Valence", "v"))) #Valence
```

Lets check out the first and the last few rows of the dataset using the `head()` and `tail()` functions.

In [19]:

```
head(super1)
```

	number_of_elements	m_am	wm_am	gm_am	wgm_am	e_am	we_am	r_an
	4	88.94447	57.86269	66.36159	36.11661	1.181795	1.0623955	122.906
	5	92.72921	58.51842	73.13279	36.39660	1.449309	1.0577551	122.906
	4	88.94447	57.88524	66.36159	36.12251	1.181795	0.9759805	122.906
	4	88.94447	57.87397	66.36159	36.11956	1.181795	1.0222909	122.906
	4	88.94447	57.84014	66.36159	36.11072	1.181795	1.1292237	122.906
	4	88.94447	57.79504	66.36159	36.09893	1.181795	1.2252028	122.906

In [20]:

```
tail(super1)
```

	number_of_elements	m_am	wm_am	gm_am	wgm_am	e_am	we_am
21258	3	89.38983	89.38983	63.69471	63.69471	0.7825737	0.7825737
21259	4	106.95788	53.09577	82.51538	43.13556	1.1771448	1.2541187
21260	5	92.26674	49.02137	64.81266	32.86775	1.3232866	1.5716301
21261	2	99.66319	95.60910	99.43388	95.46432	0.6908472	0.5301975
21262	2	99.66319	97.09560	99.43388	96.90108	0.6908472	0.6408830
21263	3	87.46833	86.85850	82.55576	80.45872	1.0412701	0.8952292

Summary Statistics:

Since the given dataset is an extremely wide dataset, we will compute the summaries of the variables block by block. i.e., Atomic Mass, Entropy, Atomic Radius and so on. This will help us see the output clearly and deduce some information from the same. Let us begin with the first block, which is **"Atomic Mass"** and also includes **Number of Elements**:

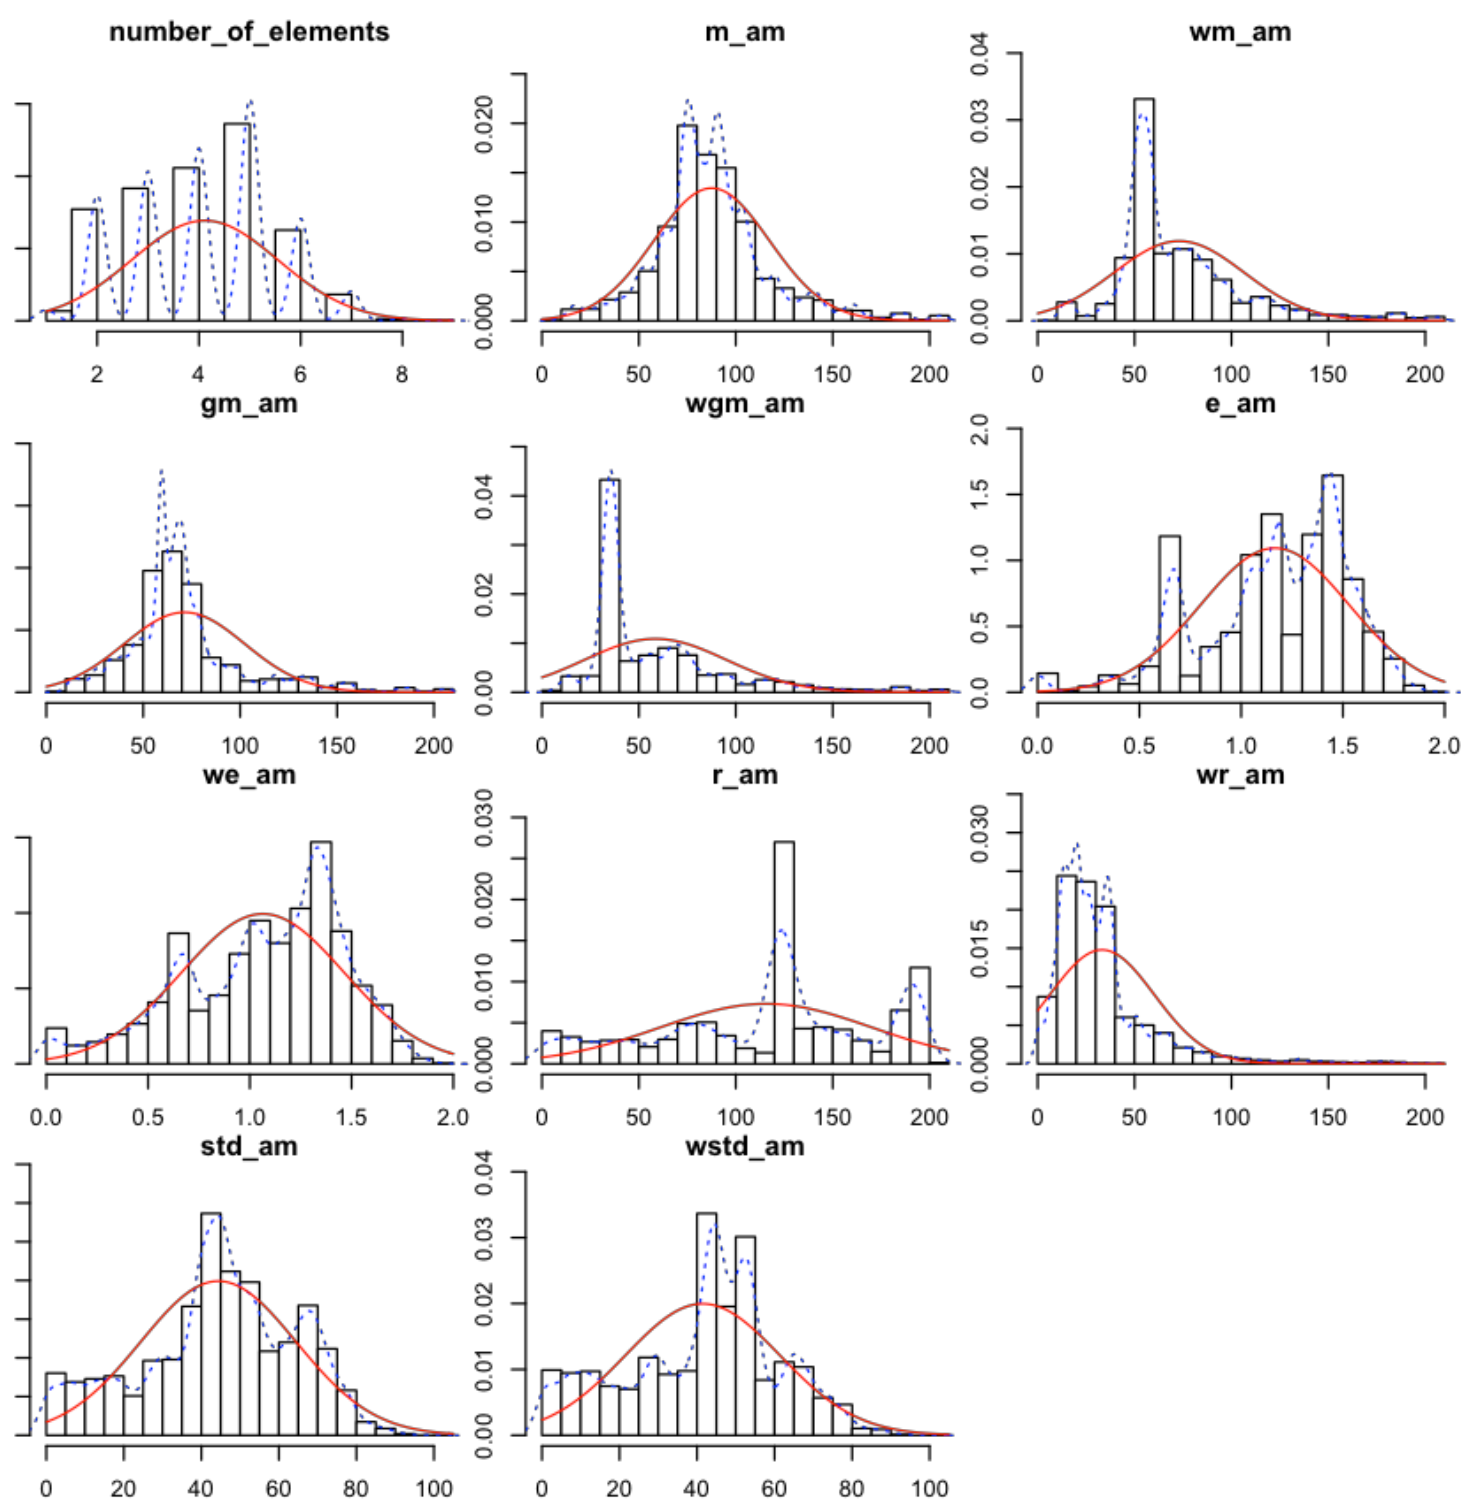
In [21]:

```
summary(super1[1:11])
multi.hist(super1[1:11],dcol= c("blue","red"),dltty=c("dotted","solid"))
```

number_of_elements	m_am	wm_am	gm_am
Min. :1.000	Min. : 6.941	Min. : 6.423	Min. :
5.321			
1st Qu.:3.000	1st Qu.: 72.458	1st Qu.: 52.144	1st Qu.: 5
8.041			
Median :4.000	Median : 84.923	Median : 60.697	Median : 6
6.362			
Mean :4.115	Mean : 87.558	Mean : 72.988	Mean : 7
1.291			
3rd Qu.:5.000	3rd Qu.:100.404	3rd Qu.: 86.104	3rd Qu.: 7
8.117			
Max. :9.000	Max. :208.980	Max. :208.980	Max. :20
8.980			

wgm_am	e_am	we_am	r_am
Min. : 1.961	Min. :0.0000	Min. :0.0000	Min. : 0.0
0			
1st Qu.: 35.249	1st Qu.:0.9667	1st Qu.:0.7754	1st Qu.: 78.5
1			
Median : 39.918	Median :1.1995	Median :1.1468	Median :122.9
1			
Mean : 58.540	Mean :1.1656	Mean :1.0639	Mean :115.6
0			
3rd Qu.: 73.113	3rd Qu.:1.4445	3rd Qu.:1.3594	3rd Qu.:154.1
2			
Max. :208.980	Max. :1.9838	Max. :1.9582	Max. :207.9
7			

wr_am	std_am	wstd_am
Min. : 0.00	Min. : 0.00	Min. : 0.00
1st Qu.: 16.82	1st Qu.: 32.89	1st Qu.: 28.54
Median : 26.64	Median : 45.12	Median : 44.29
Mean : 33.23	Mean : 44.39	Mean : 41.45
3rd Qu.: 38.36	3rd Qu.: 59.32	3rd Qu.: 53.63
Max. :205.59	Max. :101.02	Max. :101.02



In [22]:

```
table(super[1])
```

1	2	3	4	5	6	7	8	9
285	3280	3895	4496	5792	2666	774	61	14

- From the summary of "number_of_elements" we can see that the minimum number of elements that a superconductor contains is 1 and the maximum number of elements that a superconductor contains is 9. On an average, a superconductor contains 4 elements in its chemical composition. And this statement can be backed up in the next point.
- We can see from the output of running the **table()** function on this column that the count of the superconductors which have 5 elements are the highest. The next highest count belongs to superconductors with 4 number of elements. While superconductors with 2 or 3 elements have close numbers. Superconductors with 1 or 7 or 8 or 9 elements are highly variable, the least being superconductors with 9 elements with a count of just 14.

The variables extracted for the "Atomic Mass" of a superconductor are observed to have the following properties:

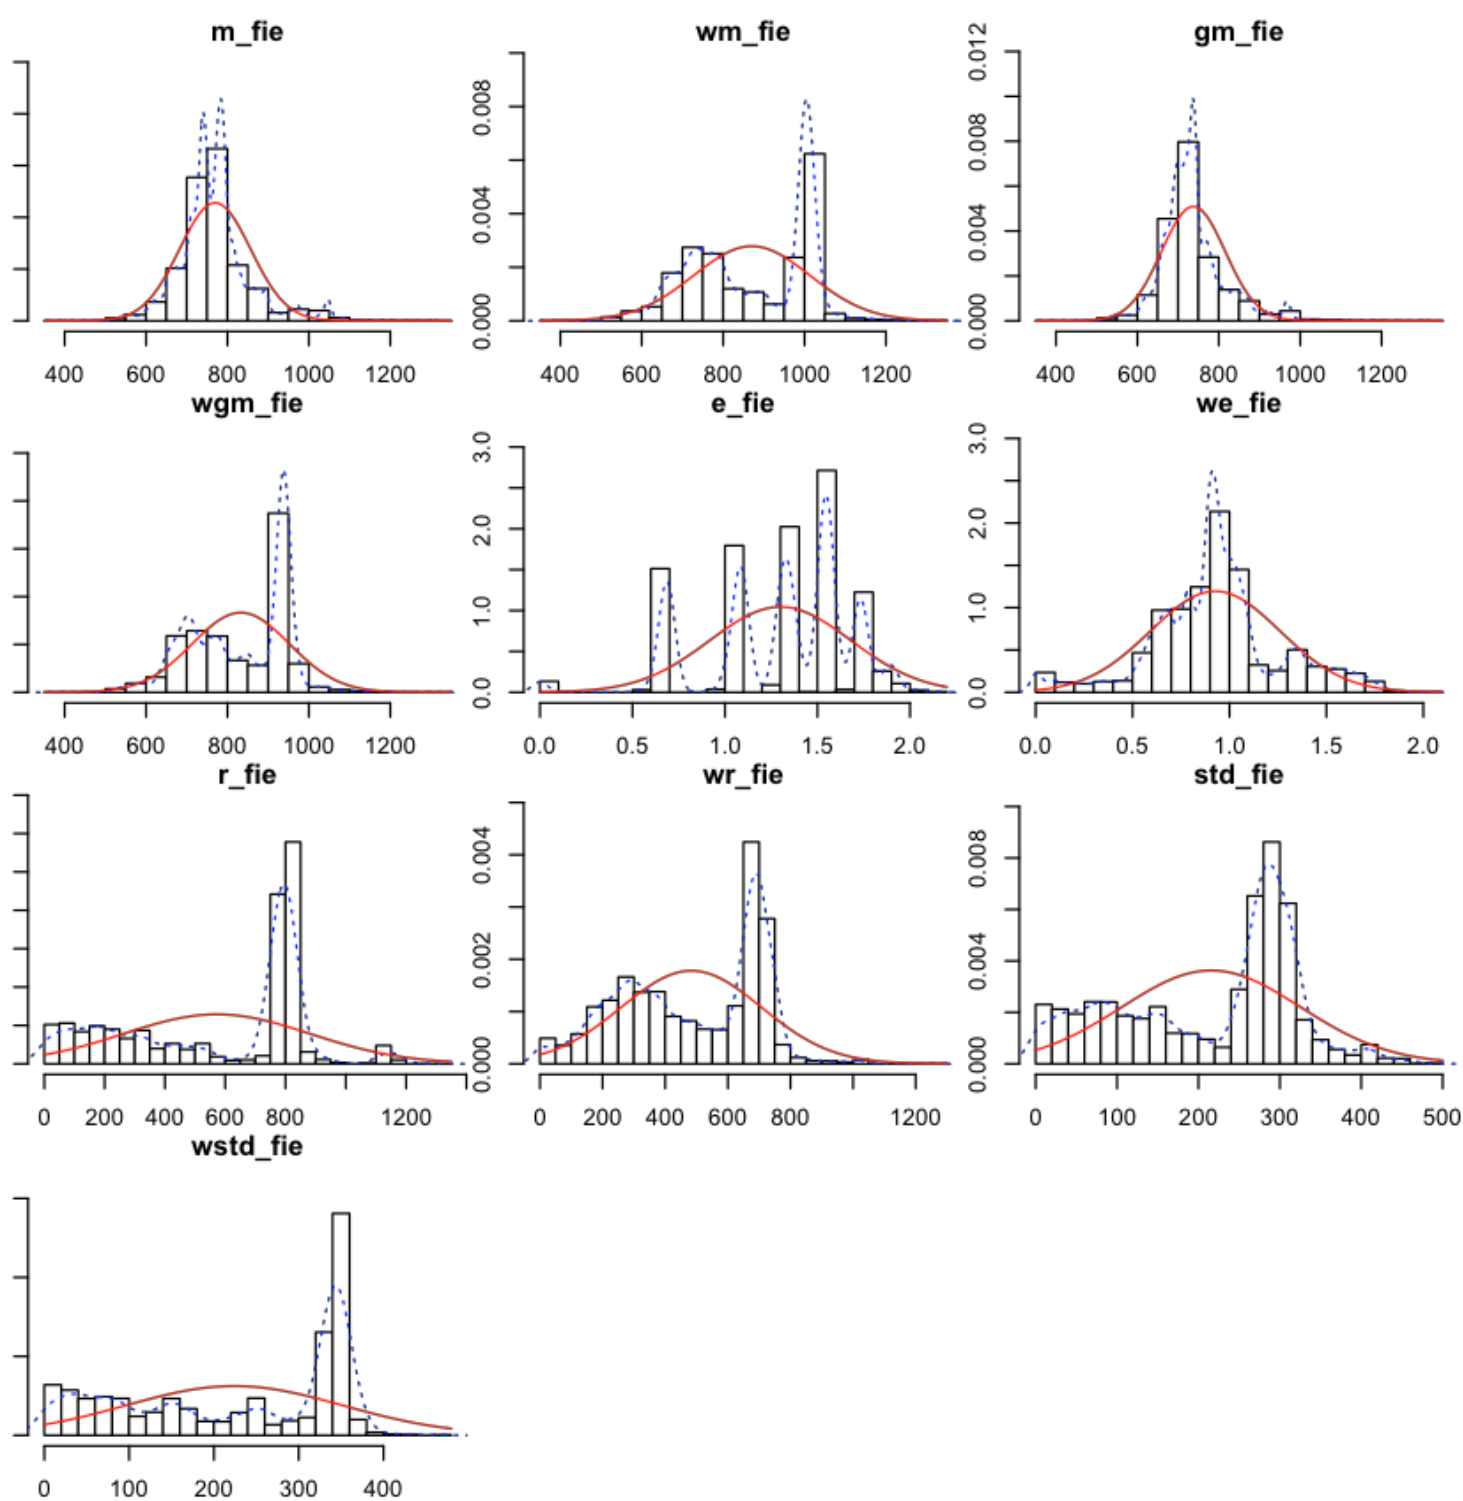
- The mean, weighted mean, geometric mean and weighted geometric mean of atomic masses have the same upper limit or max value of 208.980 AMU.
- While the minimum values of mean, weighted mean and geometric mean lie close to each other; the minimum value of the weighted geometric mean is far lesser than its counterparts.
- Since entropy is a logarithmic value, the range of that value is much lesser than the rest of the derivatives.
- The range variables are pretty much self explanatory as they are the difference between the atomic masses of the elements.
- Where as the standard deviation and weighted standard deviation of the atomic mass has a mean value between 41-45 AMU. Almost half as much as the mean of mean atomic mass of the superconductor.
- We can also see from the histograms that mean atomic mass seems to have a somewhat normal distribution. the geometric mean seems to be skewed to the right.
- The entropy and range based derived functions also seem to have a very variable distribution.
- The highly variant data could help the model perform better.

The next block of stats is **"First Ionization Energy"**:

In [23]:

```
summary(super1[12:21])  
multi.hist(super1[12:21],dcol= c("blue","red"),dltty=c("dotted","solid"))
```

m_fie		wm_fie		gm_fie		wgm_fie	
Min.	: 375.5	Min.	: 375.5	Min.	: 375.5	Min.	: 375.5
1st Qu.:	723.7	1st Qu.:	738.9	1st Qu.:	692.5	1st Qu.:	720.1
Median :	764.9	Median :	890.0	Median :	728.0	Median :	856.2
Mean :	769.6	Mean :	870.4	Mean :	737.5	Mean :	832.8
3rd Qu.:	796.3	3rd Qu.:	1004.1	3rd Qu.:	765.7	3rd Qu.:	937.6
Max.	:1313.1	Max.	:1348.0	Max.	:1313.1	Max.	:1327.6
e_fie		we_fie		r_fie		wr_fie	
Min.	:0.000	Min.	:0.0000	Min.	: 0.0	Min.	: 0.0
1st Qu.:	1.086	1st Qu.:	0.7538	1st Qu.:	262.4	1st Qu.:	291.1
Median :	1.356	Median :	0.9168	Median :	764.1	Median :	510.4
Mean :	1.299	Mean :	0.9267	Mean :	572.2	Mean :	483.5
3rd Qu.:	1.551	3rd Qu.:	1.0618	3rd Qu.:	810.6	3rd Qu.:	690.7
Max.	:2.158	Max.	:2.0386	Max.	:1304.5	Max.	:1251.9
std_fie		wstd_fie					
Min.	: 0.0	Min.	: 0.00				
1st Qu.:	114.1	1st Qu.:	92.99				
Median :	266.4	Median :	258.45				
Mean :	215.6	Mean :	224.05				
3rd Qu.:	297.7	3rd Qu.:	342.66				
Max.	:499.7	Max.	:479.16				



- Since FIE refers to the ionization energy which is measured in Kilo Joules per mole, the values in general, are high.
- While the mean, weighted mean, geometric mean, weighted geometric mean, range and weighted range all lie on a similar plane with values > 1300 or close to 1300 kJ/mole.
- But the mean of all these variables lies between 730 to 870 kJ/mole.
- Entropy variables, are again logarithmic and hence they have a small range of 0-2.1 kJ/mole.
- The mean standard deviation lies around 215 kJ/mole.
- While the mean FIE histogram does seem like a decent data distribution, it is striking to see the how the variability of the other extracted features is like.
- For example, the spread of we_fie is possibly ideal. Where as the distribution of entropy_fie does not look promising at all. Meaning that, there are several redundant values which might or might not be useful for us during model development.

The next block of stats is **"Atomic Radius"**:

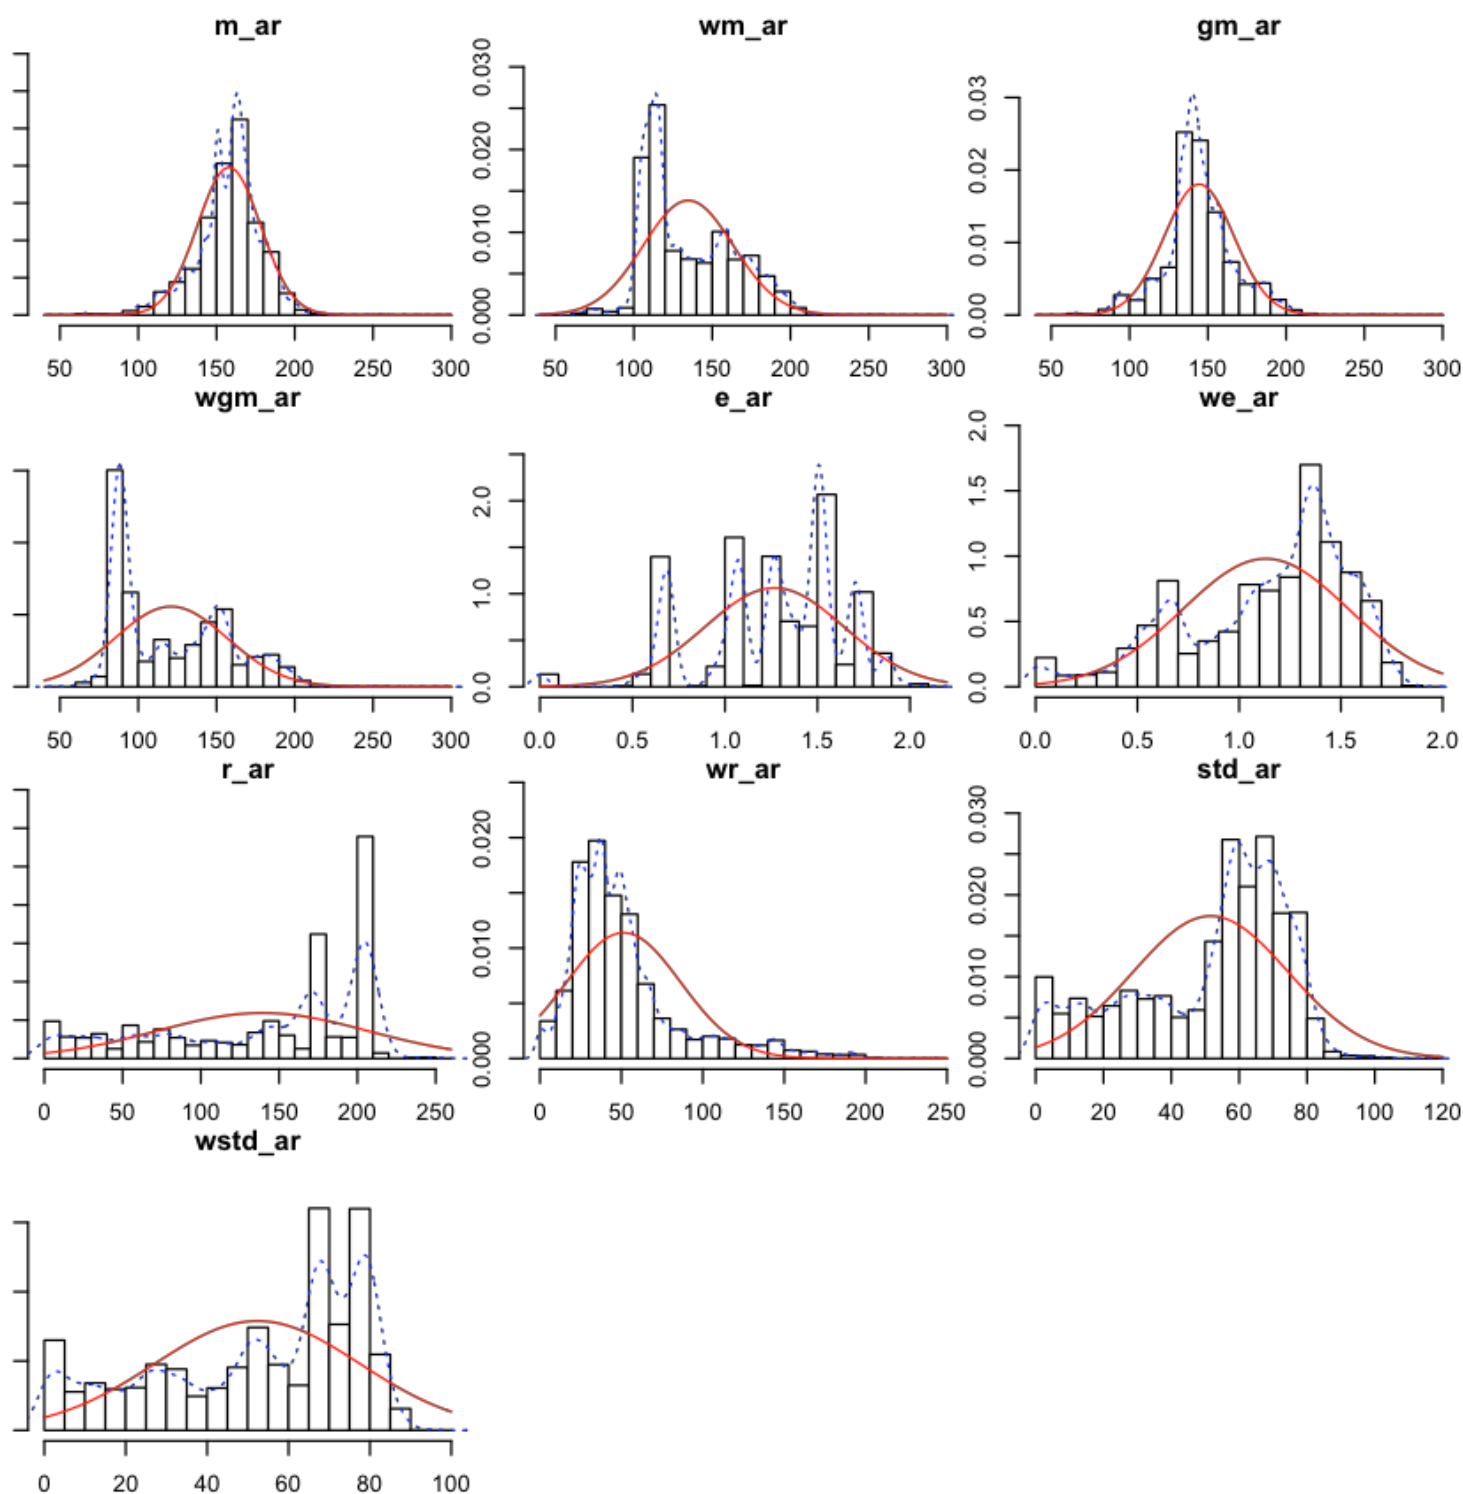
In [24]:

```
summary(super1[22:31])  
multi.hist(super1[22:31],dcol= c("blue","red"),dltty=c("dotted","solid"))
```

m_ar	wm_ar	gm_ar	wgm_ar
Min. : 48.0	Min. : 48.0	Min. : 48.0	Min. : 48.00
1st Qu.:149.3	1st Qu.:112.1	1st Qu.:133.5	1st Qu.: 89.21
Median :160.2	Median :126.0	Median :142.8	Median :113.18
Mean :158.0	Mean :134.7	Mean :144.4	Mean :120.99
3rd Qu.:169.9	3rd Qu.:158.3	3rd Qu.:155.9	3rd Qu.:150.99
Max. :298.0	Max. :298.0	Max. :298.0	Max. :298.00

e_ar	we_ar	r_ar	wr_ar
Min. :0.000	Min. :0.0000	Min. : 0.0	Min. : 0.00
1st Qu.:1.066	1st Qu.:0.8522	1st Qu.: 80.0	1st Qu.: 28.60
Median :1.331	Median :1.2429	Median :171.0	Median : 43.00
Mean :1.268	Mean :1.1311	Mean :139.3	Mean : 51.37
3rd Qu.:1.512	3rd Qu.:1.4257	3rd Qu.:205.0	3rd Qu.: 60.22
Max. :2.142	Max. :1.9037	Max. :256.0	Max. :240.16

std_ar	wstd_ar
Min. : 0.00	Min. : 0.00
1st Qu.: 35.11	1st Qu.:32.02
Median : 58.66	Median :59.93
Mean : 51.60	Mean :52.34
3rd Qu.: 69.42	3rd Qu.:73.78
Max. :115.50	Max. :97.14

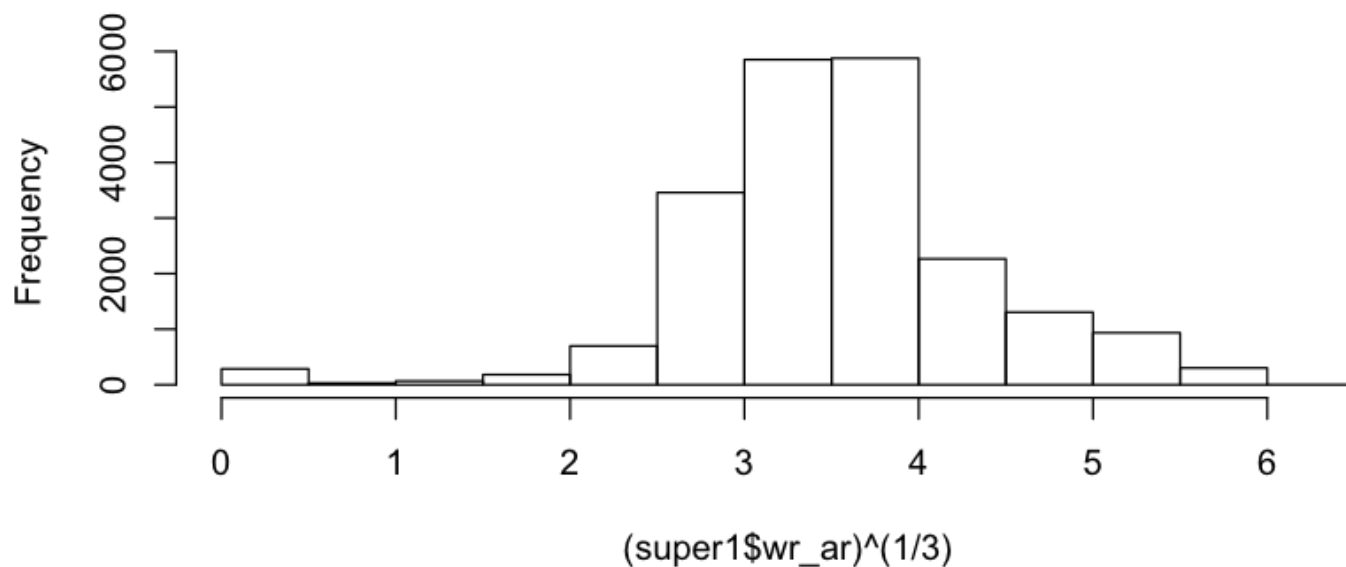


- The average atomic radius of the superconductors is around 160 picometers.
- The minimum and maximum of mean, weighted mean, geometric mean and weighted geometric mean atomic radii all lie between 48 - 298 picometers.
- The standard deviation the atomic radii has a mean of 50~ picometers.
- The distribution of the variables themselves looks extremely promising. The mean atomic radius has an extremely normal distribution.
- The weighted range atomic radius is extremely right skewed.

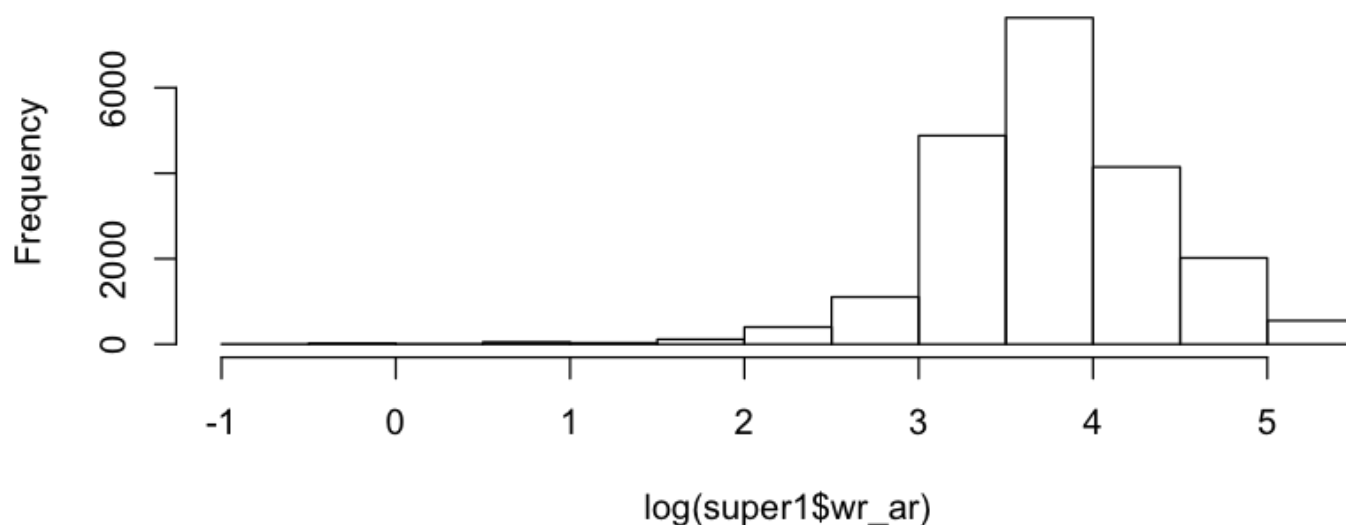
In [25]:

```
#Checking the distribution of weighted range atomic radius after cube root transformation
par(mfrow = c(2,1))
hist((super1$wr_ar)^(1/3))
hist(log(super1$wr_ar))
```

Histogram of $(\text{super1\$wr_ar})^{1/3}$



Histogram of $\log(\text{super1\$wr_ar})$



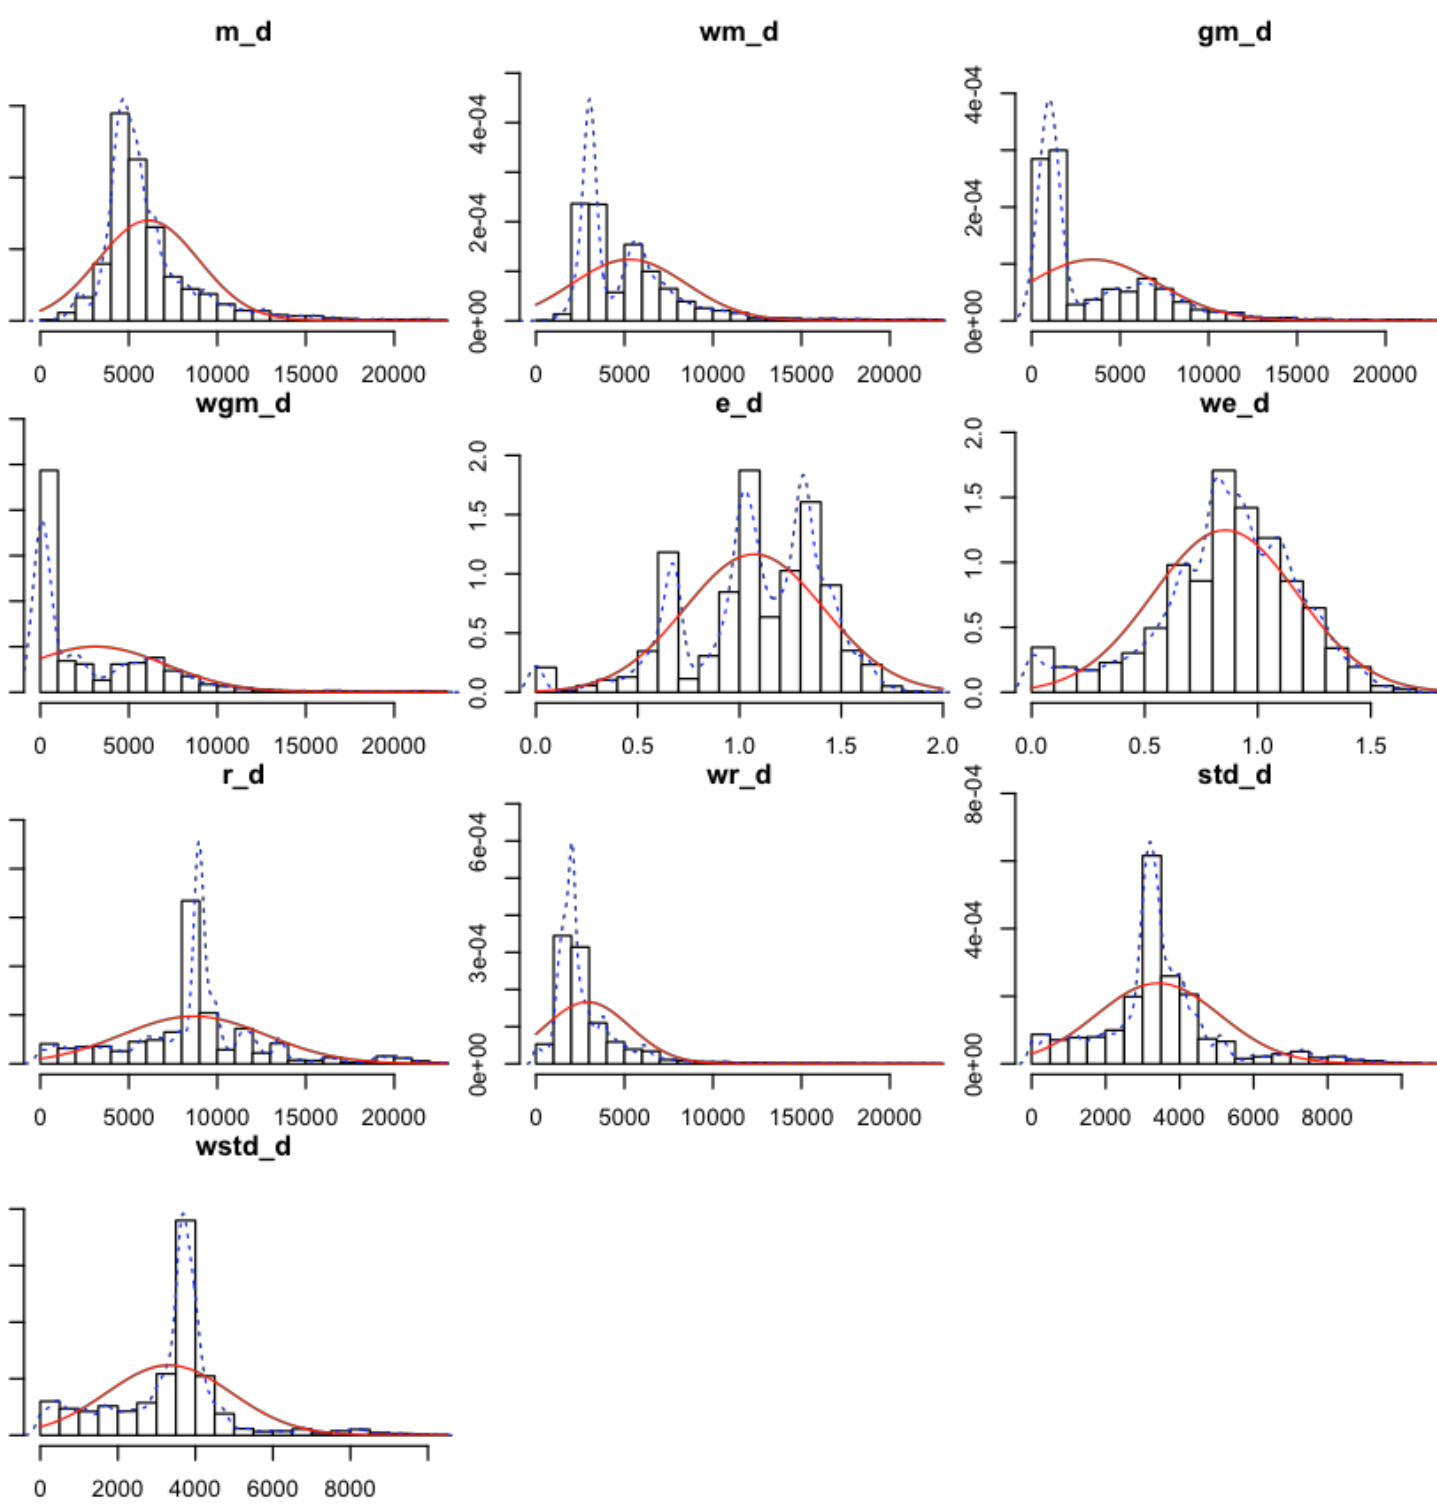
It can be seen from above that, the log transformation of the variable isn't as effective here as much as the cube root transformation.

The next block we will be looking into is **"Density"**:

In [26]:

```
summary(super1[32:41])
multi.hist(super1[32:41],dcol= c("blue","red"),dltty=c("dotted", "solid"))
```

m_d		wm_d		gm_d			
Min.	: 1.429	Min.	: 1.429	Min.	: 1.429		
1st Qu.:	4513.500	1st Qu.:	2999.158	1st Qu.:	883.117		
Median :	5329.086	Median :	4303.422	Median :	1339.975		
Mean :	6111.465	Mean :	5267.189	Mean :	3460.692		
3rd Qu.:	6728.000	3rd Qu.:	6416.333	3rd Qu.:	5794.965		
Max.	:22590.000	Max.	:22590.000	Max.	:22590.000		
wgm_d		e_d		we_d		r_d	
Min.	: 0.686	Min.	:0.000	Min.	:0.0000	Min.	:
1st Qu.:	66.747	1st Qu.:	0.914	1st Qu.:	0.6887	1st Qu.:	664
Median :	1515.365	Median :	1.091	Median :	0.8827	Median :	895
Mean :	3117.241	Mean :	1.072	Mean :	0.8560	Mean :	866
3rd Qu.:	5766.015	3rd Qu.:	1.324	3rd Qu.:	1.0809	3rd Qu.:	977
Max.	:22590.000	Max.	:1.954	Max.	:1.7034	Max.	:2258
wr_d		std_d		wstd_d			
Min.	: 0	Min.	: 0	Min.	: 0		
1st Qu.:	1657	1st Qu.:	2819	1st Qu.:	2564		
Median :	2083	Median :	3302	Median :	3626		
Mean :	2903	Mean :	3417	Mean :	3319		
3rd Qu.:	3409	3rd Qu.:	4004	3rd Qu.:	3959		
Max.	:22434	Max.	:10724	Max.	:10411		

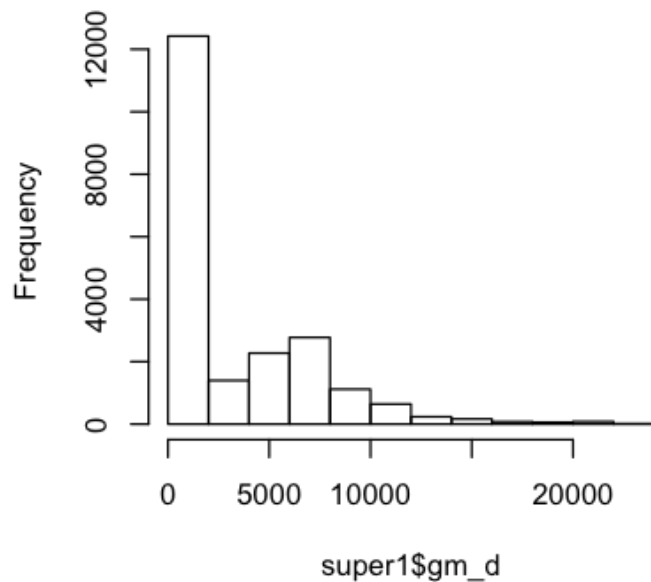


- Since density is measured per cubed meter, again, the values will be large.
- The mean density of a superconductor being around 6000~ kg/m³.
- The most dense element has a density of 22590 kg/m³.
- From the plots, it seems to appear that most of the features extracted for density seem to be normal except for geometric mean, which has a steep skew.
- We can try to apply some transformation to see if it help normalize it.

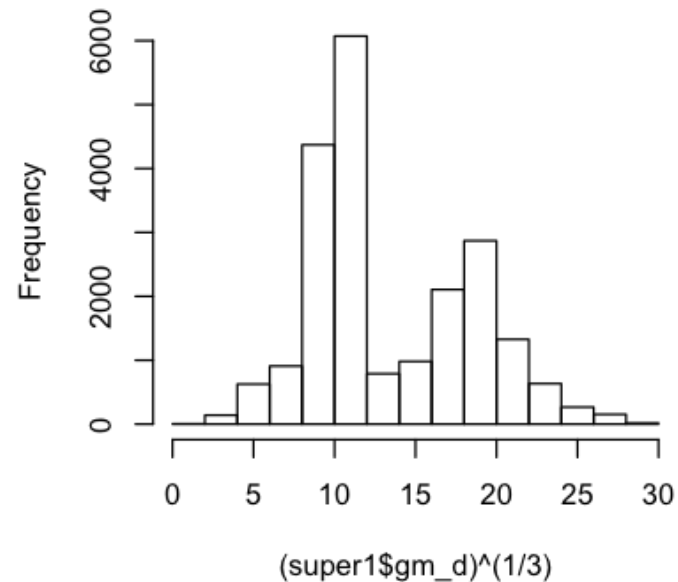
In [27]:

```
#Checking the distribution of weighted geometric mean density after transformation
par(mfrow =c(2,2))
hist(super1$gm_d) #actual data
hist((super1$gm_d)^(1/3)) #cube root transformation
hist((super1$gm_d)^(1/2)) #square root transformation
hist(log(super1$wgm_d)) #logarithmic transformation
```

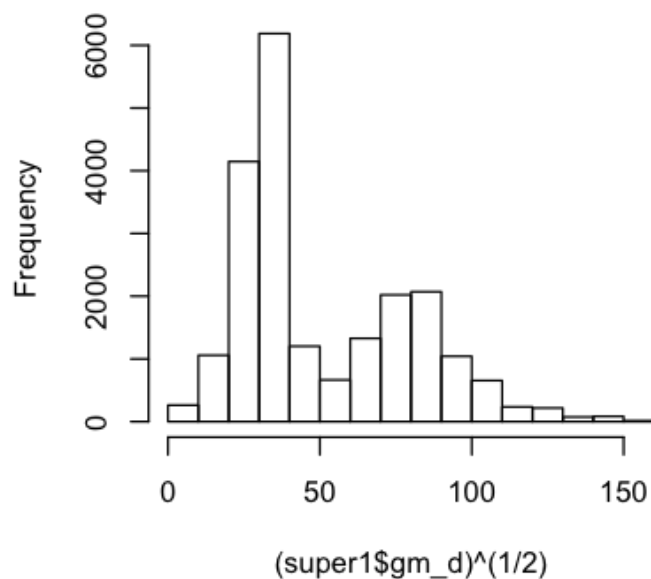
Histogram of super1\$gm_d



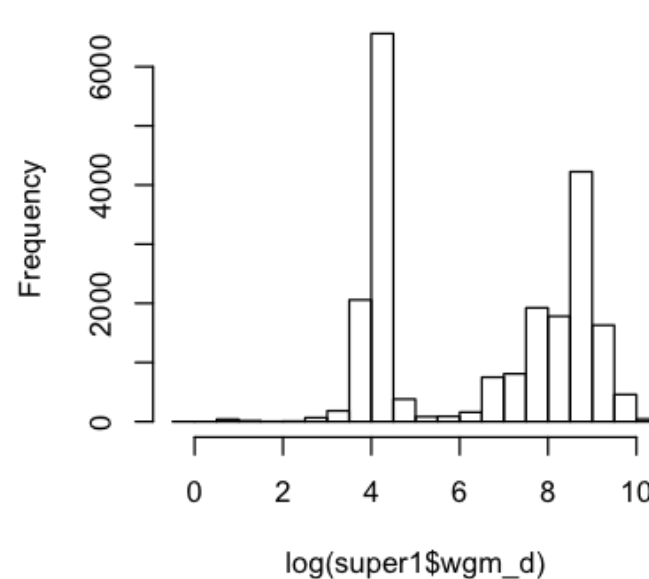
Histogram of (super1\$gm_d)^(1/3)



Histogram of (super1\$gm_d)^(1/2)



Histogram of log(super1\$wgm_d)



From the above histograms, we can see that the transformations do not seem to have any useful outcome .

The next block we will be looking into is **"Electron Affinity"**:

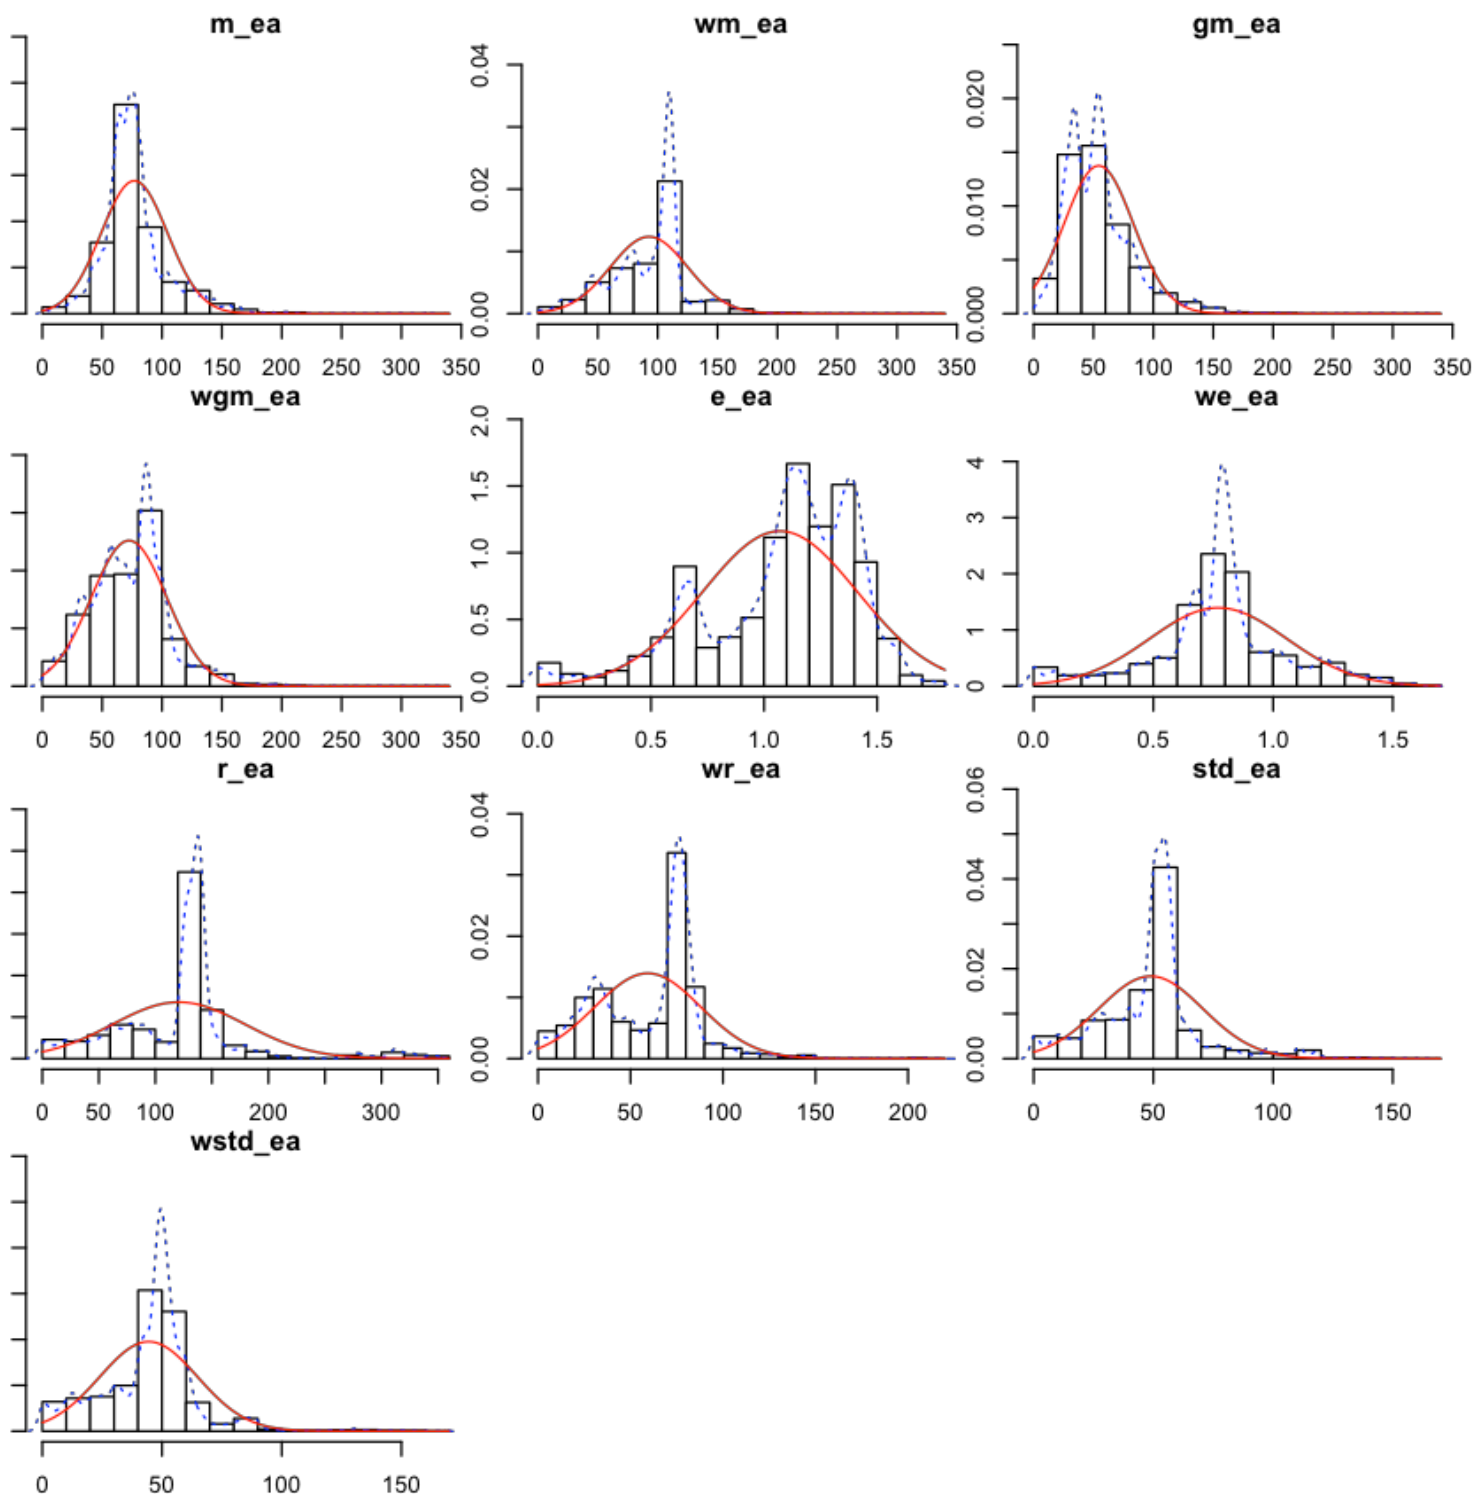
In [28]:

```
summary(super1[42:51])
multi.hist(super1[42:51],dcol= c("blue", "red"),dltty=c("dotted", "solid"))
```

m_ea	wm_ea	gm_ea	wgm_ea
Min. : 1.50	Min. : 1.50	Min. : 1.50	Min. : 1.50
1st Qu.: 62.09	1st Qu.: 73.35	1st Qu.: 33.70	1st Qu.: 50.77
Median : 73.10	Median :102.86	Median : 51.47	Median : 73.17
Mean : 76.88	Mean : 92.72	Mean : 54.36	Mean : 72.42
3rd Qu.: 85.50	3rd Qu.:110.74	3rd Qu.: 67.51	3rd Qu.: 89.98
Max. :326.10	Max. :326.10	Max. :326.10	Max. :326.10

e_ea	we_ea	r_ea	wr_ea
Min. :0.0000	Min. :0.0000	Min. : 0.0	Min. : 0.00
1st Qu.:0.8906	1st Qu.:0.6607	1st Qu.: 86.7	1st Qu.: 34.04
Median :1.1383	Median :0.7812	Median :127.0	Median : 71.16
Mean :1.0702	Mean :0.7708	Mean :120.7	Mean : 59.33
3rd Qu.:1.3459	3rd Qu.:0.8775	3rd Qu.:138.6	3rd Qu.: 76.71
Max. :1.7677	Max. :1.6754	Max. :349.0	Max. :218.70

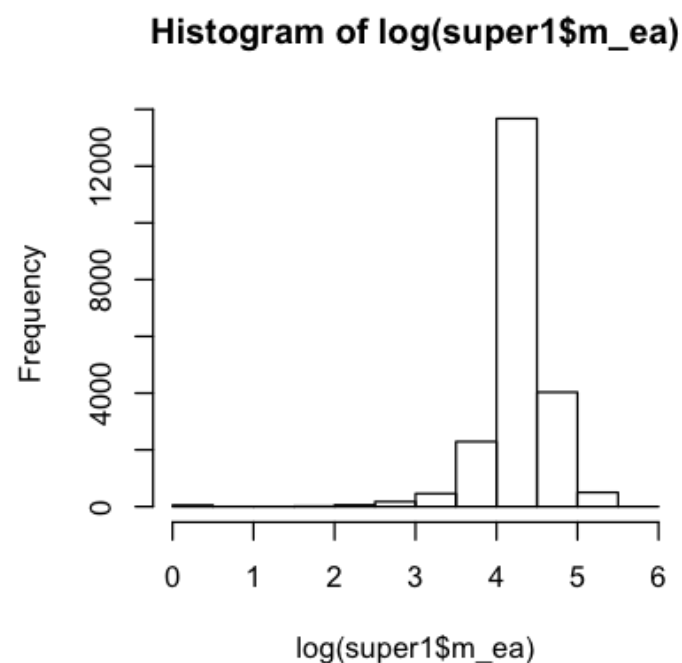
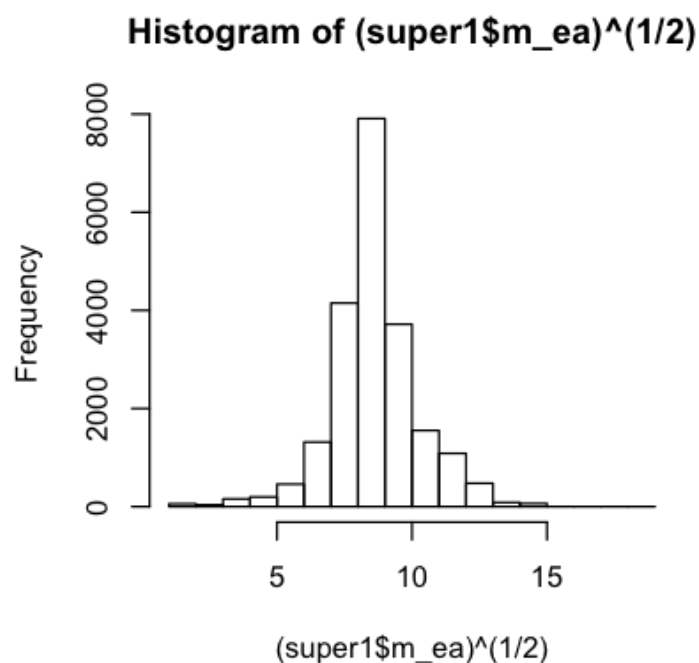
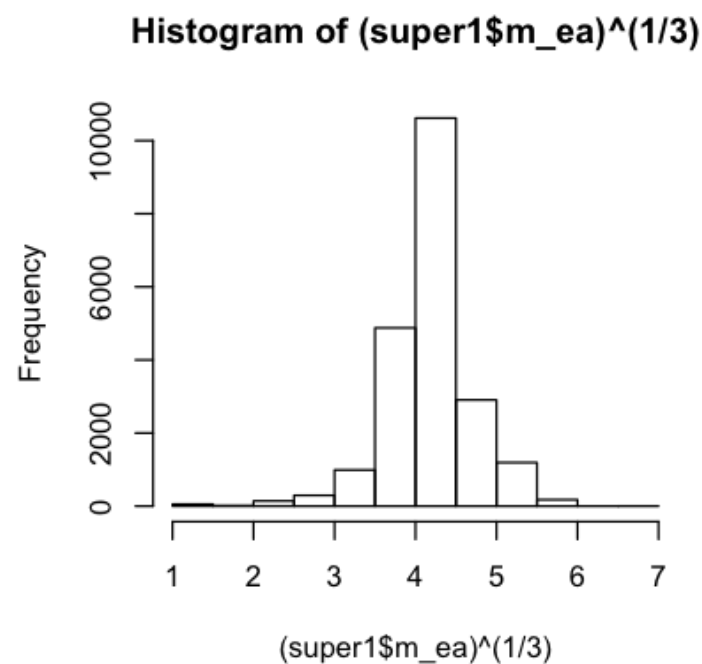
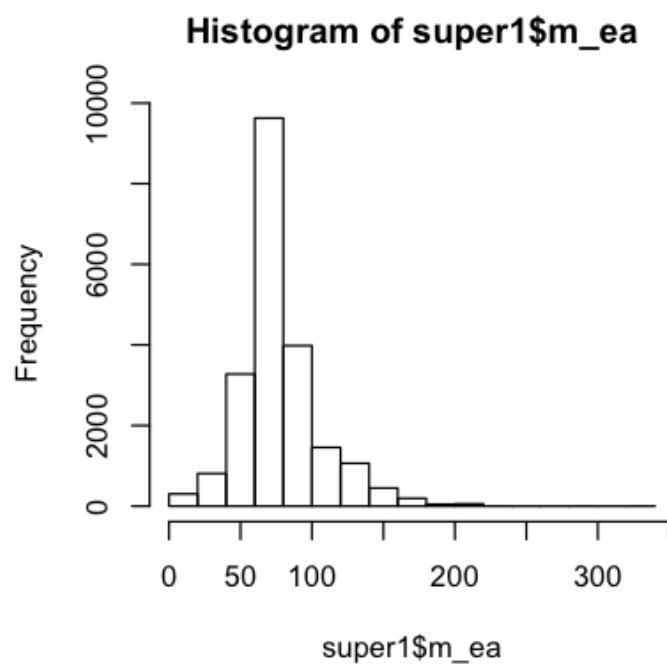
std_ea	wstd_ea
Min. : 0.00	Min. : 0.00
1st Qu.: 38.37	1st Qu.: 33.44
Median : 51.13	Median : 48.03
Mean : 48.91	Mean : 44.41
3rd Qu.: 56.22	3rd Qu.: 53.32
Max. :162.90	Max. :169.08



- Electron affinity seems to be lying between 1.50 to 326.10 kJ/mole.
- While the mean of mean, weighted mean, geometric mean and weighted geometric mean columns are highly variant.
- Most of the data is highly positively skewed. Thus, we can try to fix positive skew by applying root transformations on the data.
- Let us explore the possibilities below:

In [29]:

```
#Checking the distribution of weighted geometric mean density after transformation
par(mfrow = c(2,2))
hist(super1$m_ea) #actual data
hist((super1$m_ea)^(1/3)) #cube root transformation
hist((super1$m_ea)^(1/2)) #square root transformation
hist(log(super1$m_ea)) #logarithmic transformation
```



At LAST, we were able to transform a variable, in this case using square root, and cause a significant impact on it. <insert emoji>

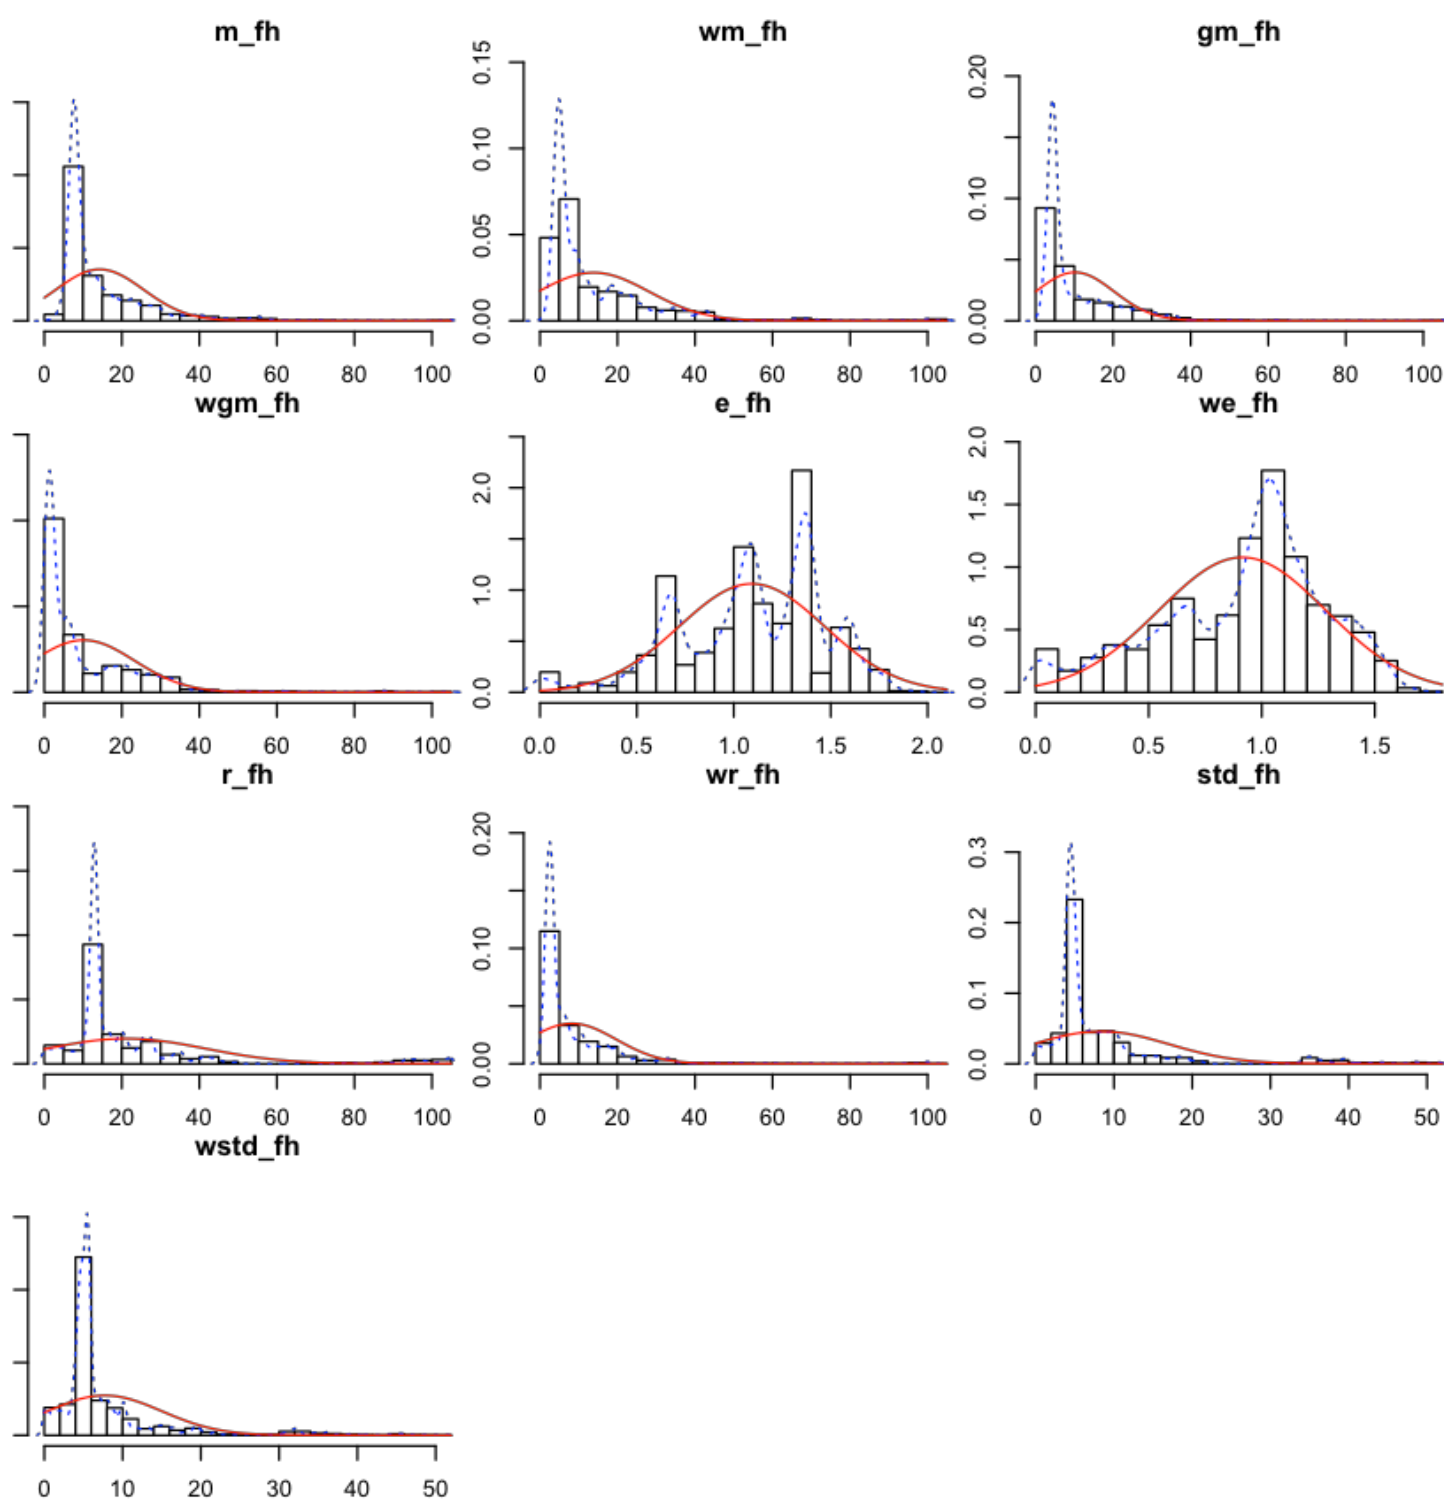
If this variable, 'm_ea' does happen to be an important predictor, we can try to check the accuracy of the model using the true value and compare it with using the transformed value.

The next block we will be looking into is "Fusion Heat":

In [30]:

```
summary(super1[52:61])
multi.hist(super1[52:61],dcol= c("blue","red"),dltty=c("dotted","solid"))
```

m_fh		wm_fh		gm_fh		wgm_fh	
Min.	: 0.222	Min.	: 0.222	Min.	: 0.222	Min.	: 0
1st Qu.:	7.589	1st Qu.:	5.033	1st Qu.:	4.110	1st Qu.:	1
Median :	9.304	Median :	8.331	Median :	5.253	Median :	4
Mean :	14.296	Mean :	13.848	Mean :	10.137	Mean :	10
3rd Qu.:	17.114	3rd Qu.:	18.514	3rd Qu.:	13.600	3rd Qu.:	16
Max.	:105.000	Max.	:105.000	Max.	:105.000	Max.	:105
e_fh		we_fh		r_fh		wr_fh	
Min.	:0.0000	Min.	:0.0000	Min.	: 0.00	Min.	: 0.00
1st Qu.:	0.8333	1st Qu.:	0.6727	1st Qu.:	12.88	1st Qu.:	2.32
Median :	1.1121	Median :	0.9950	Median :	12.88	Median :	3.43
Mean :	1.0933	Mean :	0.9141	Mean :	21.14	Mean :	8.21
3rd Qu.:	1.3781	3rd Qu.:	1.1574	3rd Qu.:	23.20	3rd Qu.:	10.49
Max.	:2.0344	Max.	:1.7472	Max.	:104.78	Max.	:102.67
std_fh		wstd_fh					
Min.	: 0.000	Min.	: 0.000				
1st Qu.:	4.261	1st Qu.:	4.603				
Median :	4.948	Median :	5.501				
Mean :	8.323	Mean :	7.718				
3rd Qu.:	9.041	3rd Qu.:	8.018				
Max.	:51.635	Max.	:51.680				

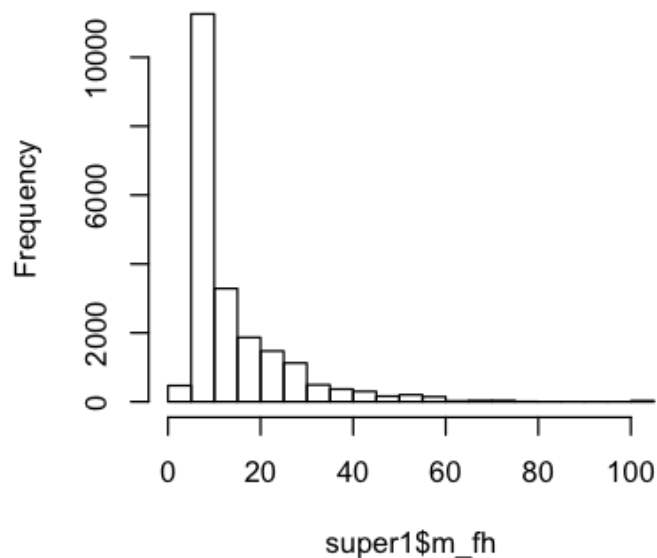


- The fusion heat values also do not seem to be large.
- They have a very low mean lurking between 10-15 kJ/mole.
- While the maximum range is around 104 kJ/mole.
- The standard deviation and weighted standard deviation also seem to be extremely identical in their values.
- From the plots, its very evident that most of the fusion heat variables also seem to be extremely skewed to the right.
- We can check if we are able to eliminate this by using transformation.

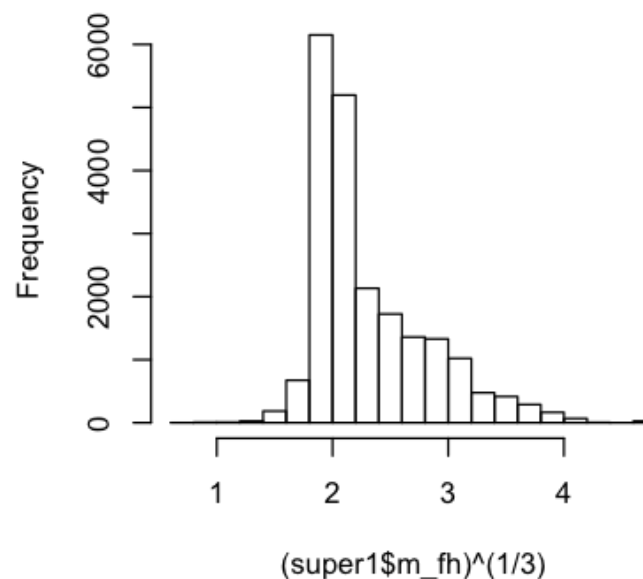
In [31]:

```
#Checking the distribution of mean fusion heat after transformation
par(mfrow = c(2,2))
hist(super1$m_fh) #actual data
hist((super1$m_fh)^(1/3)) #cube root transformation
hist((super1$m_fh)^(1/2)) #square root transformation
hist(log(super1$m_fh)) #logarithmic transformation
```

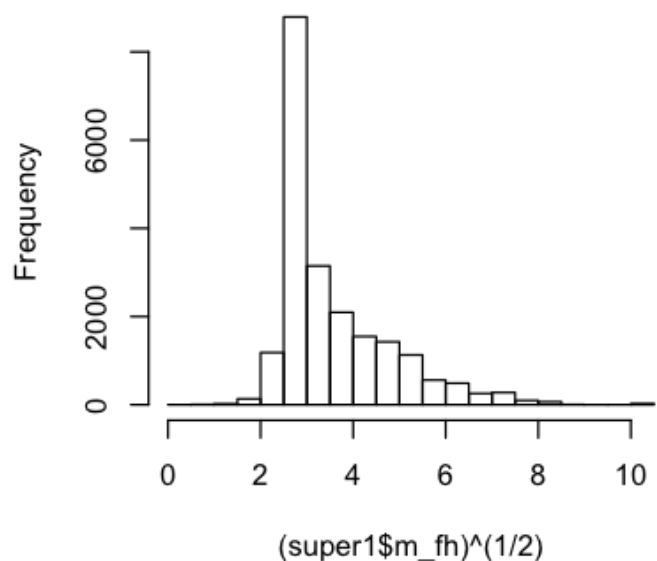
Histogram of super1\$m_fh



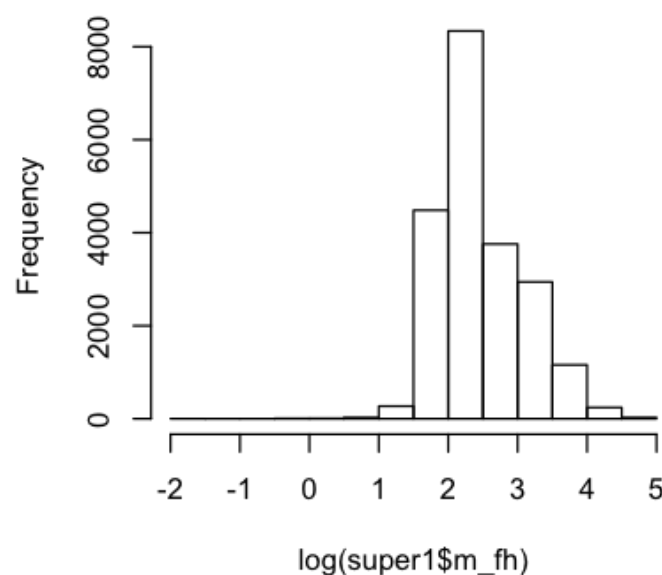
Histogram of (super1\$m_fh)^(1/3)



Histogram of (super1\$m_fh)^(1/2)



Histogram of log(super1\$m_fh)



Out of all the transformations applied, the log transformation seems to be of a *close to normal* fit than any other. Thus, if this variable is also considered to be of significant importance from later exploration, we can apply transformations as necessary.

Next block we would be looking into is **"Thermal Conductivity"**:

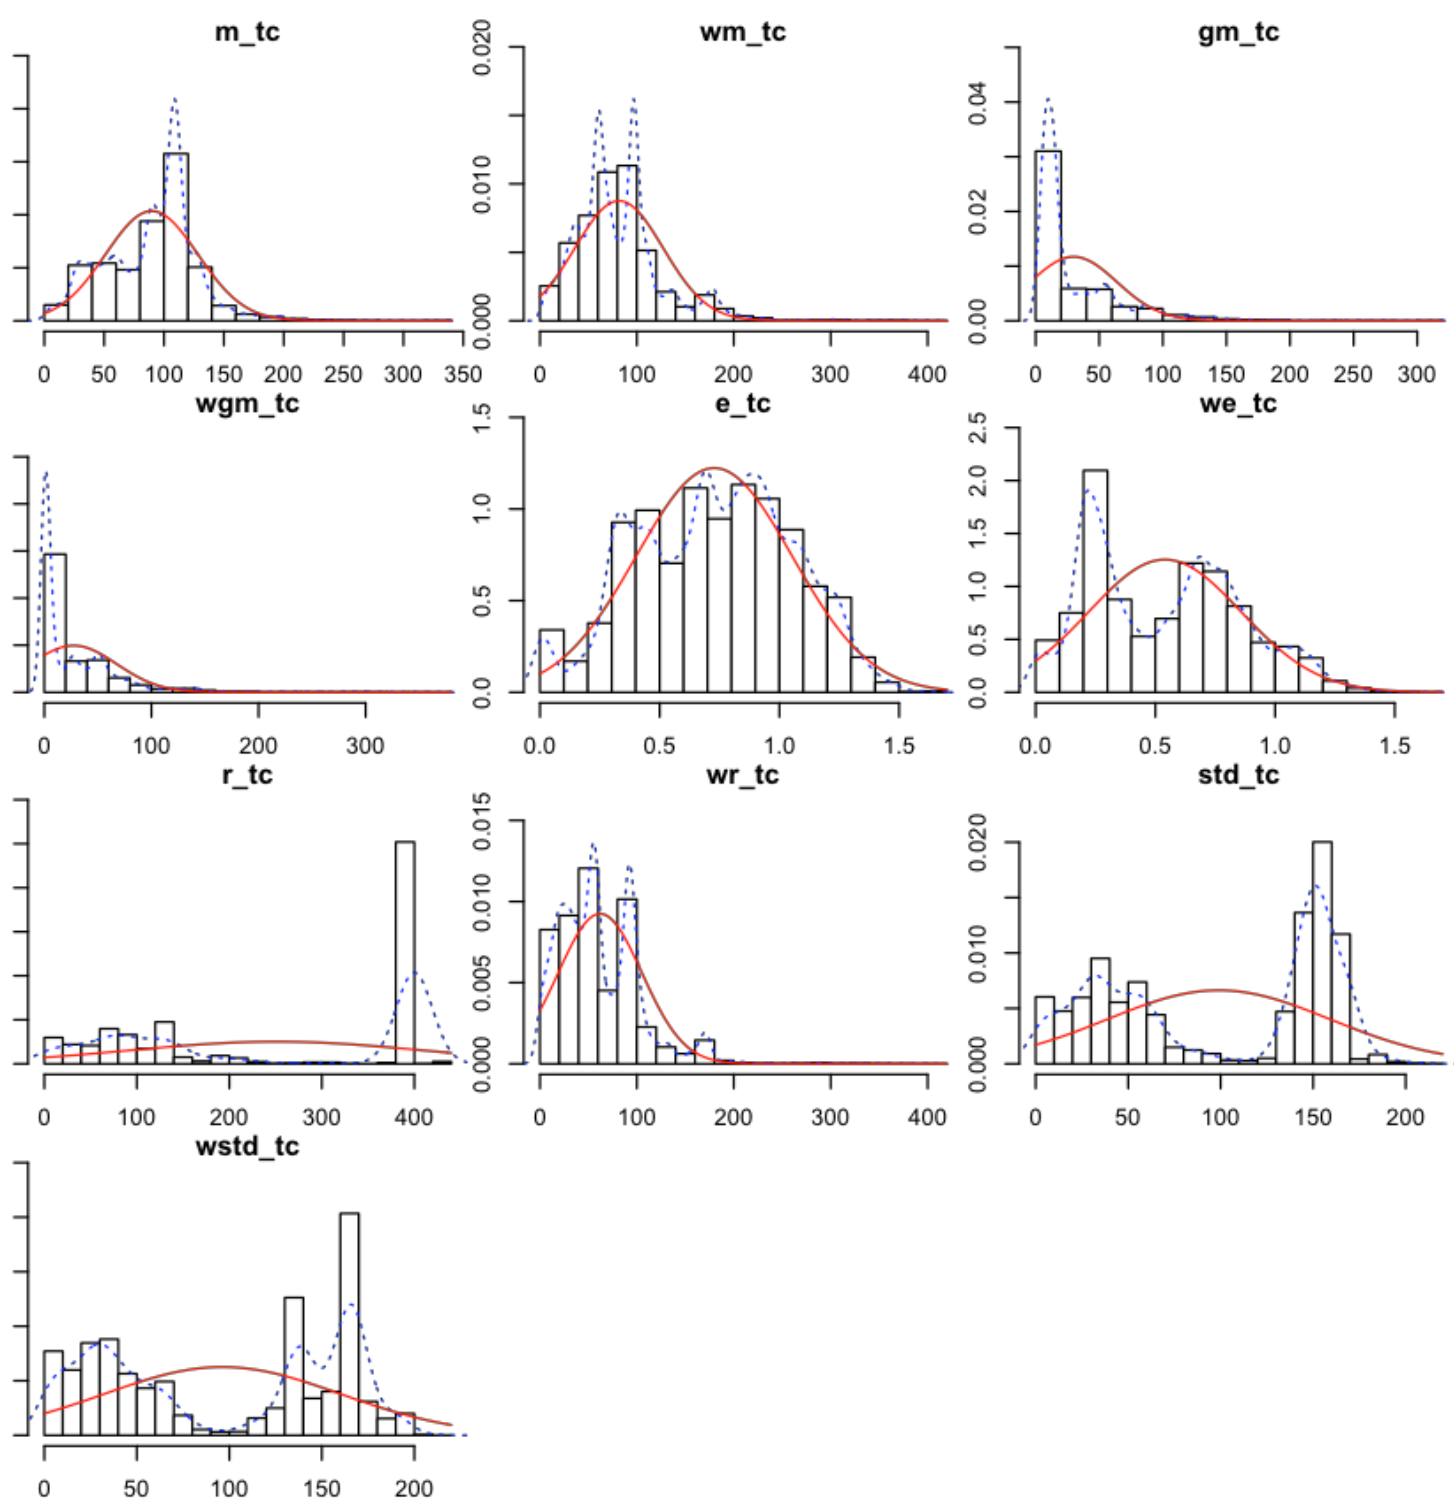
In [32]:

```
summary(super1[62:71])
multi.hist(super1[62:71], dcol= c("blue", "red"), dlty=c("dotted", "solid"))
```


	m_tc	wm_tc	gm_tc	wgm_tc
tc				
Min.	: 0.0266	Min. : 0.0266	Min. : 0.0266	Min. :
0.023				
1st Qu.:	61.0000	1st Qu.: 54.1810	1st Qu.: 8.3398	1st Qu.:
1.087				
Median :	96.5044	Median : 73.3333	Median : 14.2876	Median :
6.096				
Mean :	89.7069	Mean : 81.5491	Mean : 29.8417	Mean :
27.308				
3rd Qu.:	111.0053	3rd Qu.: 99.0629	3rd Qu.: 42.3713	3rd Qu.:
47.308				
Max.	:332.5000	Max. :406.9600	Max. :317.8836	Max. :
376.033				

	e_tc	we_tc	r_tc	wr_tc
Min.	:0.0000	Min. :0.0000	Min. : 0.00	Min. : 0.00
1st Qu.:	0.4578	1st Qu.:0.2507	1st Qu.: 86.38	1st Qu.: 29.35
Median :	0.7387	Median :0.5458	Median :399.80	Median : 56.56
Mean :	0.7276	Mean :0.5400	Mean :250.89	Mean : 62.03
3rd Qu.:	0.9622	3rd Qu.:0.7774	3rd Qu.:399.97	3rd Qu.: 91.87
Max.	:1.6340	Max. :1.6130	Max. :429.97	Max. :401.44

	std_tc	wstd_tc
Min.	: 0.00	Min. : 0.00
1st Qu.:	37.93	1st Qu.: 31.99
Median :	135.76	Median :113.56
Mean :	98.94	Mean : 96.23
3rd Qu.:	153.81	3rd Qu.:162.71
Max.	:214.99	Max. :213.30



- Thermal conductivity also has low values with 75% of its mean values lying below 111 watts per meter-K.
- The lowest value being 0.0266 watts per meter-K and the high values of 332.50 and 406.96 watts per meter-K for the mean and weighted mean thermal conductivities respectively.
- The distributions of thermal conductivity based features are extremely varied.
- For example, the **std_tc** and **wstd_tc** variables seem to be **bimodal** in nature.
- Each of these individual features needs to be checked/transformed.
- We will not be going into it now as it can get extremely complex. But we can always come back to this for insights after estimating their relevance to the target variable.

And the final block we would be looking into is "**Valence**":

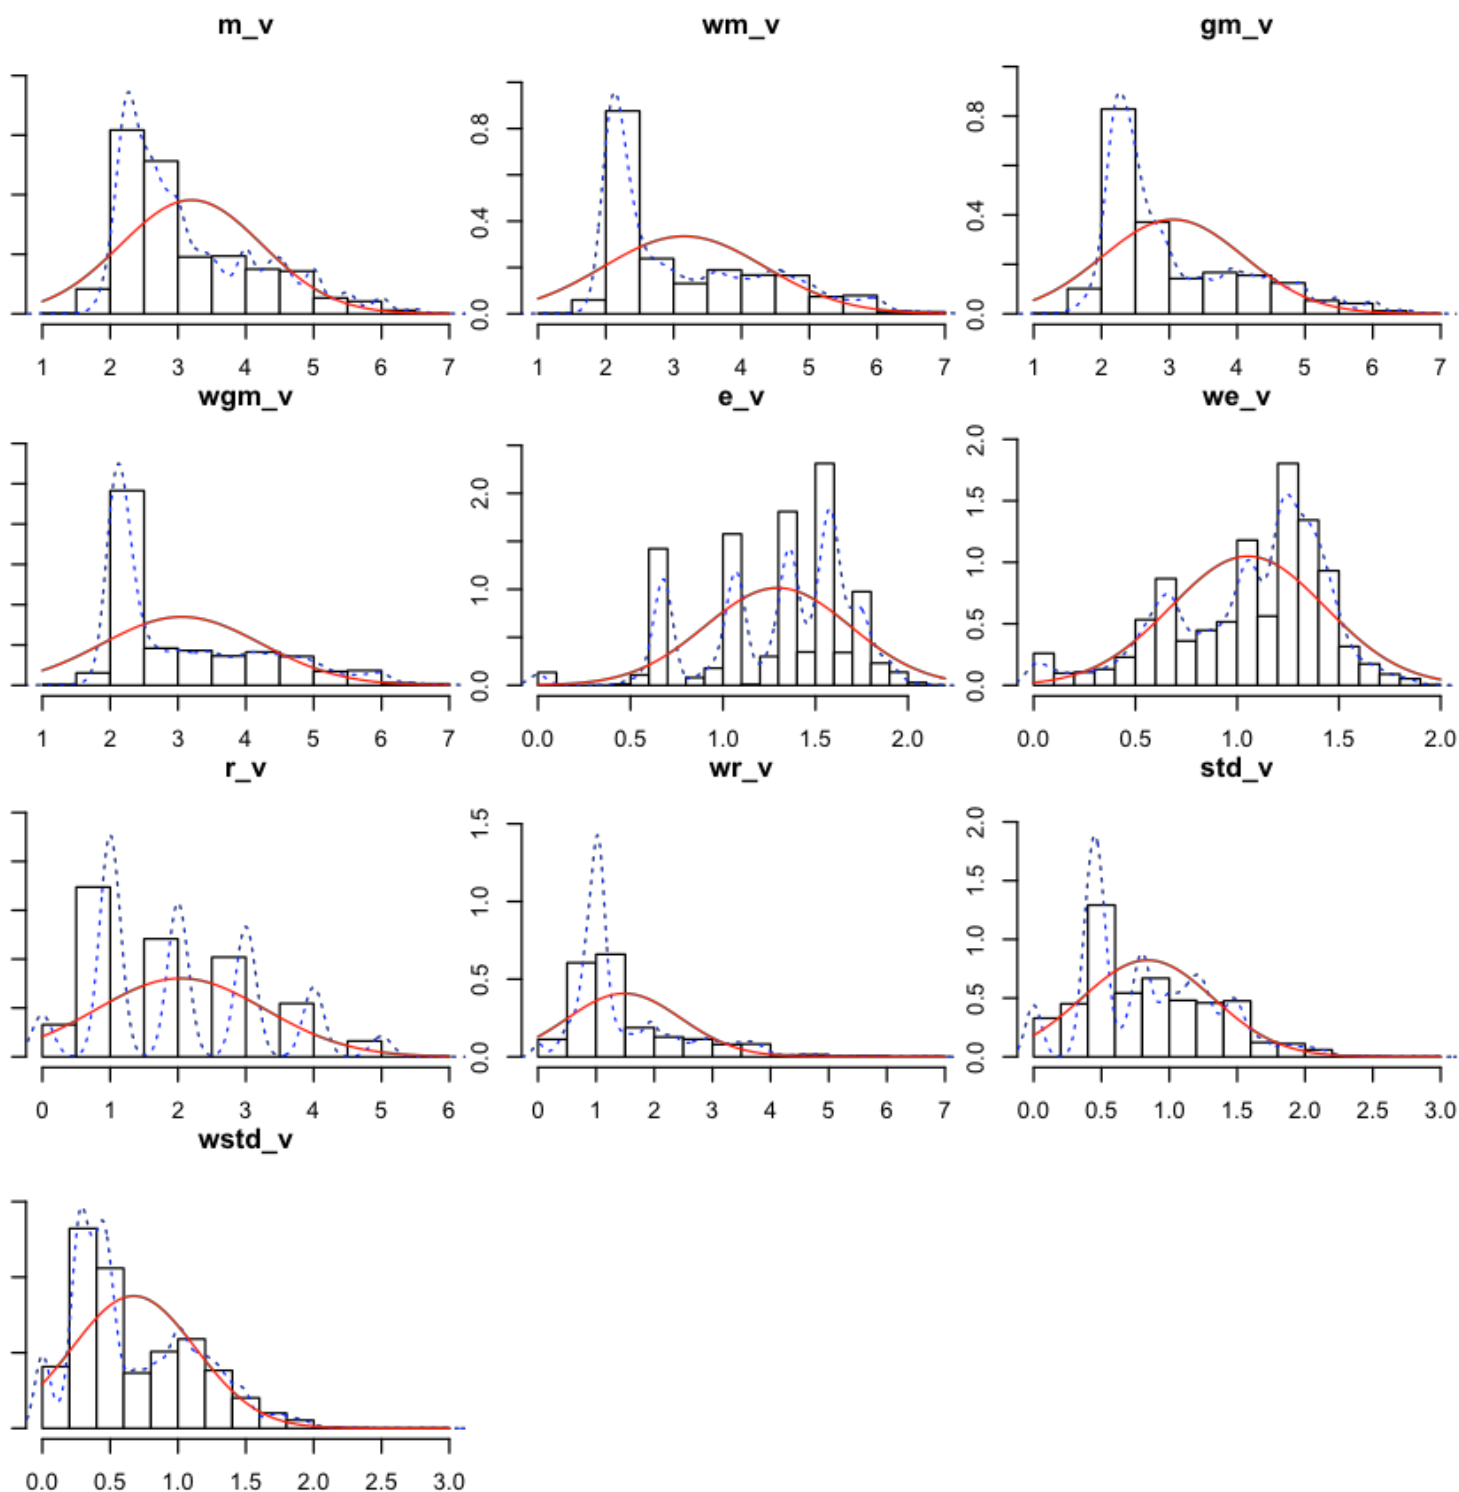
In [33]:

```
summary(super1[72:81])
multi.hist(super1[72:81],dcol= c("blue","red"),dltty=c("dotted","solid"))
```

m_v	wm_v	gm_v	wgm_v
Min. :1.000	Min. :1.000	Min. :1.000	Min. :1.000
1st Qu.:2.333	1st Qu.:2.117	1st Qu.:2.280	1st Qu.:2.091
Median :2.833	Median :2.618	Median :2.615	Median :2.434
Mean :3.198	Mean :3.153	Mean :3.057	Mean :3.056
3rd Qu.:4.000	3rd Qu.:4.026	3rd Qu.:3.728	3rd Qu.:3.915
Max. :7.000	Max. :7.000	Max. :7.000	Max. :7.000

e_v	we_v	r_v	wr_v
Min. :0.000	Min. :0.0000	Min. :0.000	Min. :0.0000
1st Qu.:1.061	1st Qu.:0.7757	1st Qu.:1.000	1st Qu.:0.9215
Median :1.369	Median :1.1665	Median :2.000	Median :1.0631
Mean :1.296	Mean :1.0528	Mean :2.041	Mean :1.4830
3rd Qu.:1.589	3rd Qu.:1.3308	3rd Qu.:3.000	3rd Qu.:1.9184
Max. :2.142	Max. :1.9497	Max. :6.000	Max. :6.9922

std_v	wstd_v
Min. :0.0000	Min. :0.0000
1st Qu.:0.4518	1st Qu.:0.3069
Median :0.8000	Median :0.5000
Mean :0.8393	Mean :0.6740
3rd Qu.:1.2000	3rd Qu.:1.0204
Max. :3.0000	Max. :3.0000



- Valence is just the number of chemical bonds made by the element.
- Although the values appear to be continuous, they could be discrete, lying within a range. As an element cannot form bonds in decimal values.
- The minimum and maximum number of bonds formed by a superconductor are 1 and 7 respectively.

Unique Materials:

In [34]:

```
unique <- read.csv('unique_m.csv')
```

In [35]:

```
head(unique)
```

H	He	Li	Be	B	C	N	O	F	Ne	Na	Mg	Al	Si	P	S	Cl	Ar	K	Ca	Sc	Ti	V	Cr
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

In [36]:

```
#Converting the dataset into a better format for analysis and visualisation.
un <- unique
un[ un > 0 ] <- 1
un$critic_temp <- unique$critical_temp
un$critical_temp <- NULL
```

Warning message in Ops.factor(left, right):
 “’>’ not meaningful for factors”

In [37]:

```
#Converting the values into factor format
un[1:86] <- lapply(un[1:86], factor)
un1 <- data.frame(apply(un[1:86], 2, table))
```

In [38]:

```
#Renaming the columns
un2<- un1[,grepl("*Freq",names(un1))]
un2 <- un2 %>% rename_at(vars(contains(".Freq")), funs(str_replace(., ".Freq",
"")))
un2 <- un2[c(-1),]
rownames(un2) <- NULL
un2
```

H	He	Li	Be	B	C	N	O	F	Ne	Na	Mg	Al	Si	P	S
299	21263	311	96	1205	1274	306	11964	669	21263	322	522	731	725	355	694

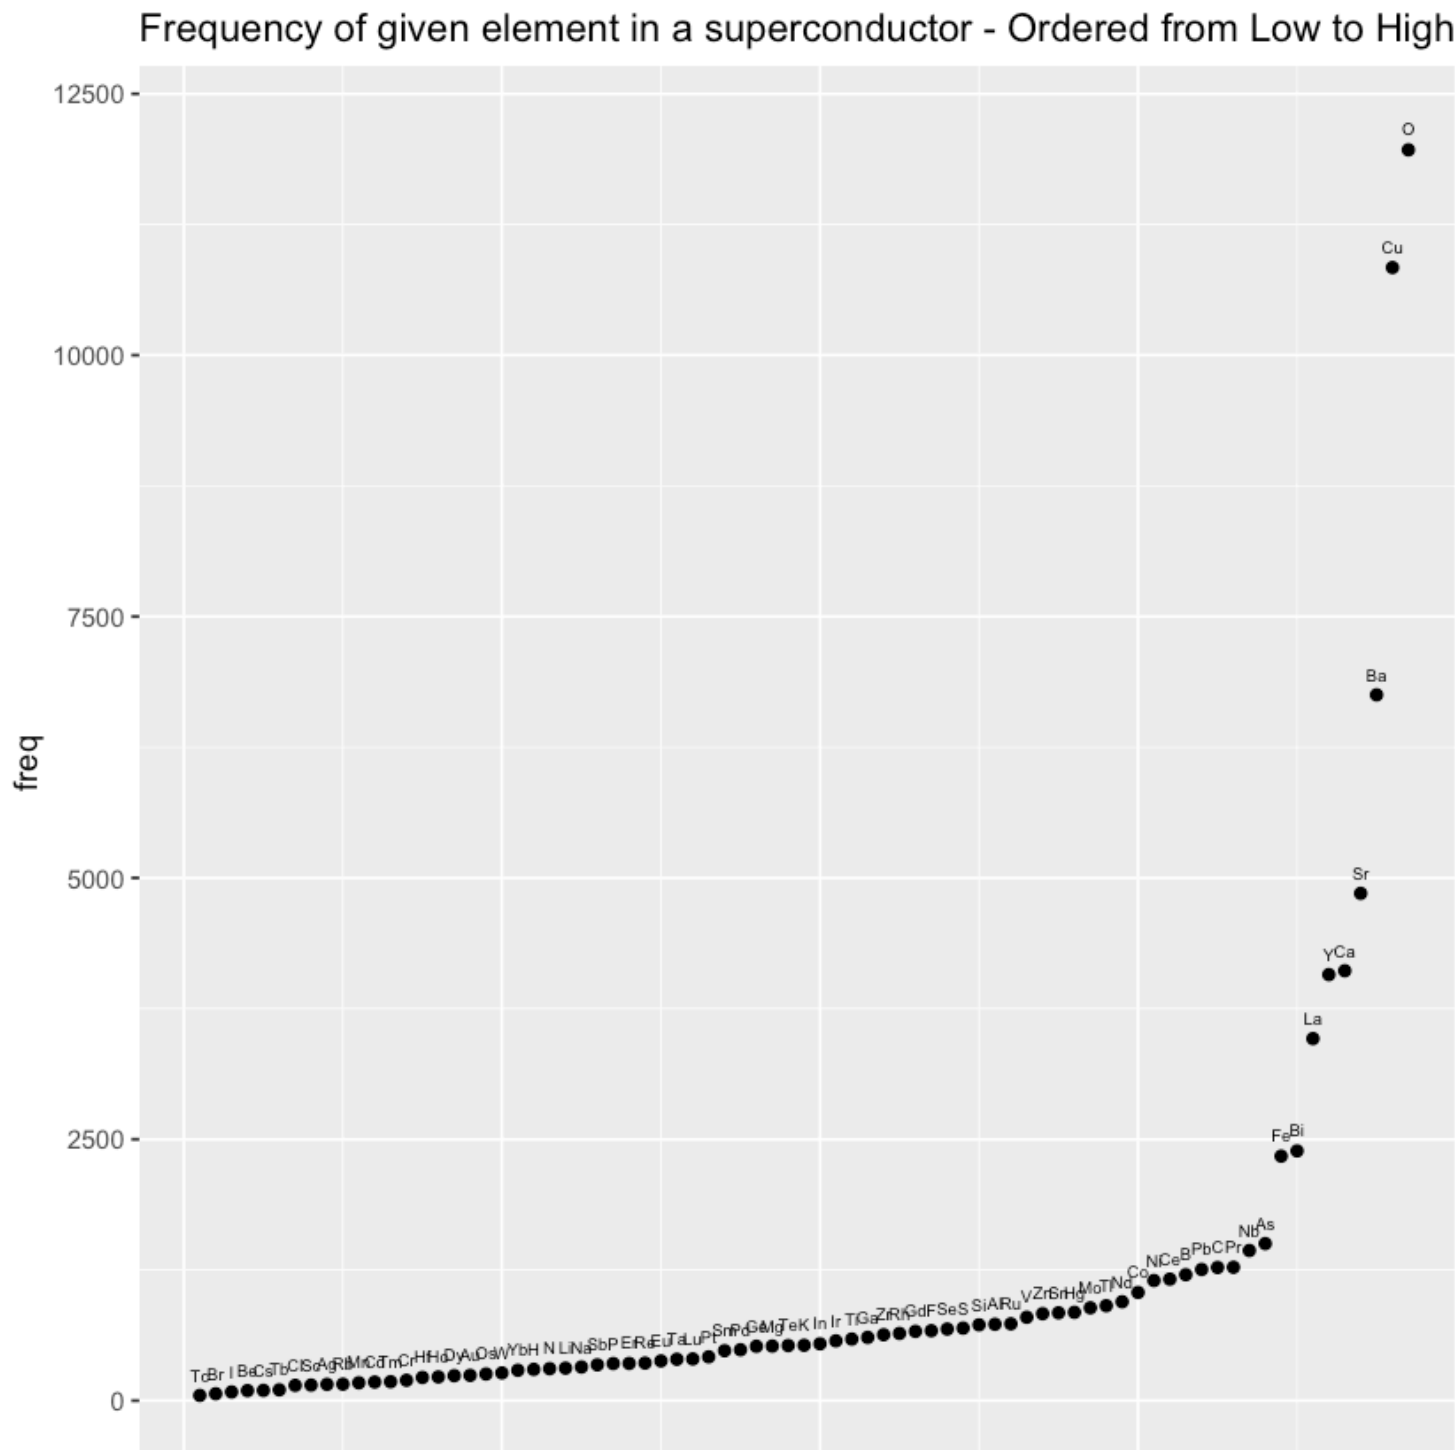
In [39]:

```
un3 <- data.frame(t(un2))
un4 <- tibble::rownames_to_column(un3, "VALUE")
colnames(un4)[1] <- "element"
colnames(un4)[2] <- "freq"
un4 <- un4[!grepl(21263, un4$freq),]
rownames(un4) <- NULL
#Ordering the dataframe by frequency of the material (ascending)
un5 <- un4[order(un4$freq),]
rownames(un5) <- NULL
un6<-un5
un6<- cbind(1:77,un5)
names(un6)[names(un6) == "1:77"] <- "sno"
```

Warning message:
“`list_len()` is deprecated as of rlang 0.2.0.
Please use `new_list()` instead.
This warning is displayed once per session.”

In [40]:

```
#Finally, plottting the dataframe
ggplot(un6, aes(sno,freq)) + geom_point() + geom_text(aes(label=element), position = position_nudge(y = 200),size=2) + theme(axis.title.x=element_blank(),
axis.text.x=element_blank(),
axis.ticks.x=element_blank()) + ggtitle("Frequency of given element in
a superconductor - Ordered from Low to High")
```



- We can see that Oxygen containing superconductors are the largest in number.
- Followed by Copper and Barium.

Scatter Plots between each Predictor and Target:

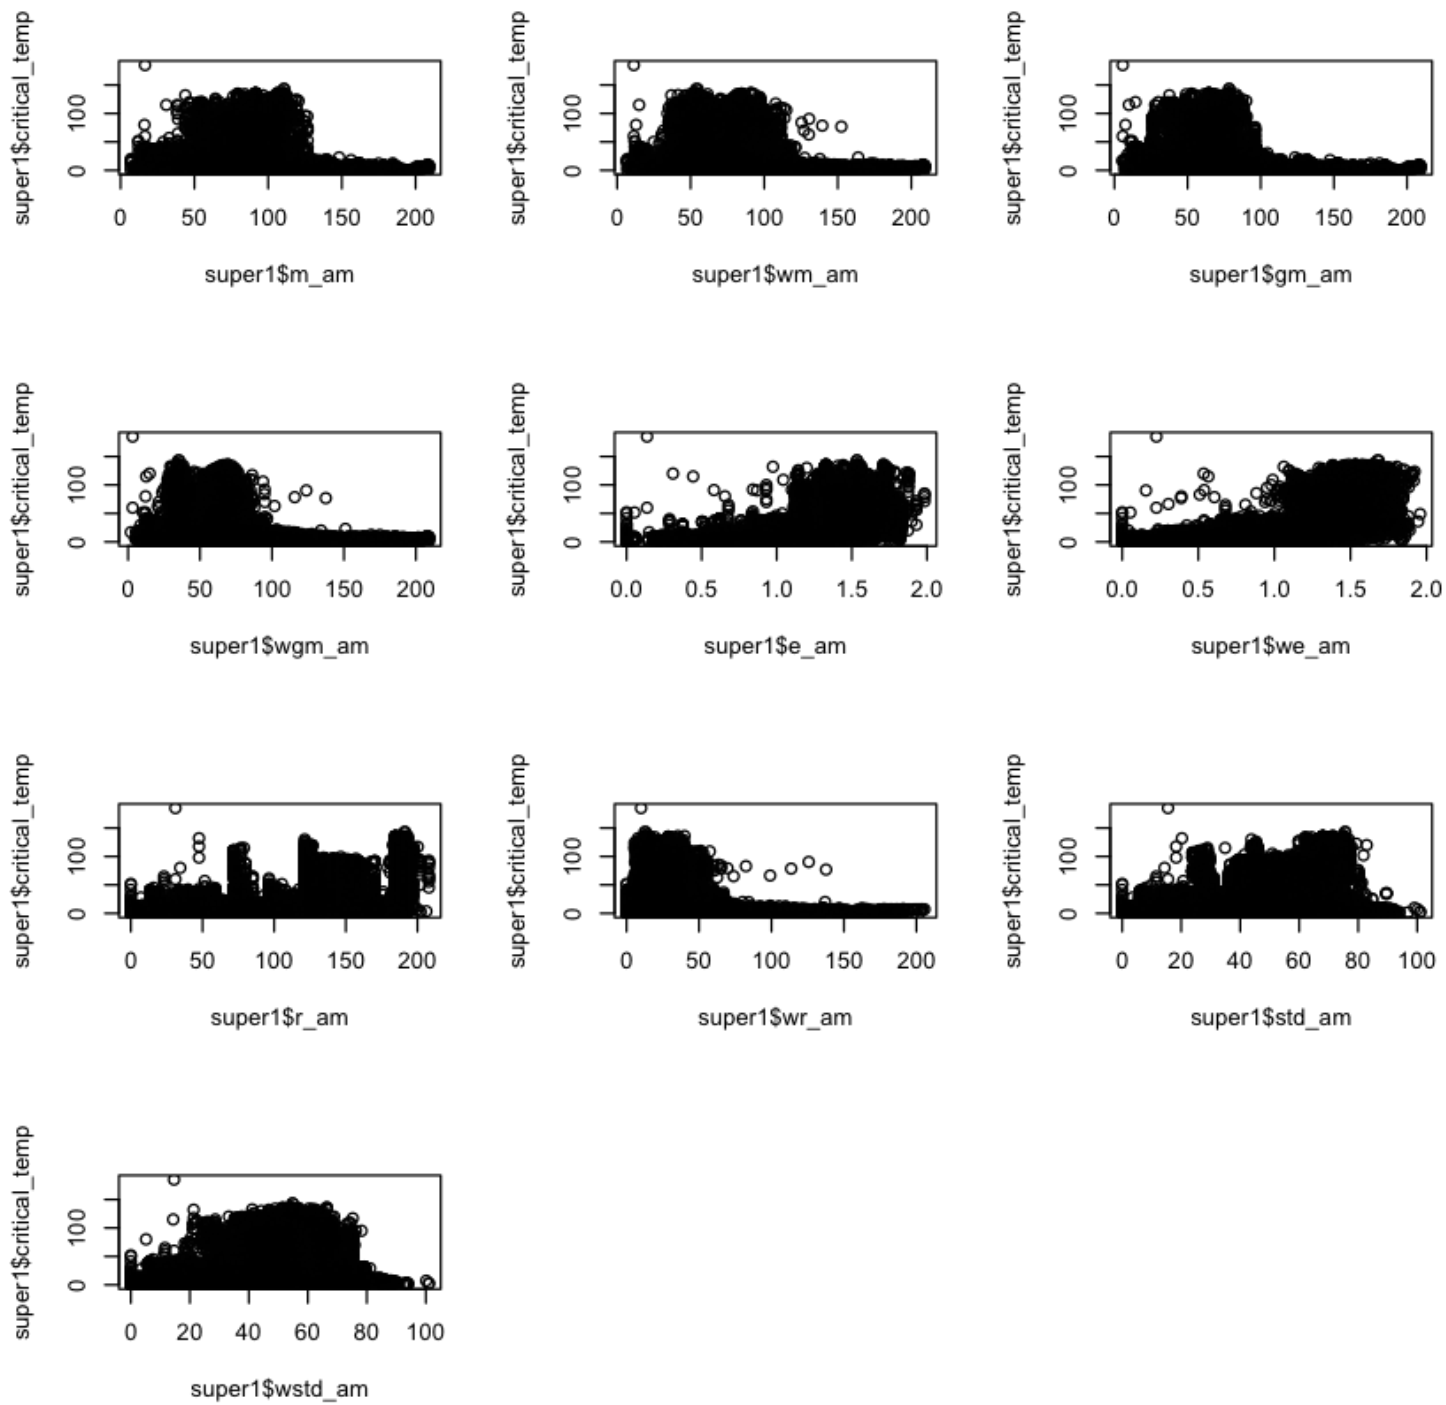
Scatter plots are a very good way to understand relationship between two variables. The intensity, strength, direction and disctribution of these plots gives a good evidence is drawing conclusions which could be useful for the model development stage.

We will be plotting all the variables with the target. Albeit this is tedious with such a large number of variables, it a good practice to do.

1. Atomic Mass vs Critical Temp Plots

In [42]:

```
#Scatter plots of Atomic Mass based variables:
par(mfrow = c(4,3))
#dev.new(width=5, height=4)
plot(super1$m_am, super1$critical_temp)
plot(super1$wm_am, super1$critical_temp)
plot(super1$gm_am, super1$critical_temp)
plot(super1$wgm_am, super1$critical_temp)
plot(super1$e_am, super1$critical_temp)
plot(super1$we_am, super1$critical_temp)
plot(super1$r_am, super1$critical_temp)
plot(super1$wr_am, super1$critical_temp)
plot(super1$std_am, super1$critical_temp)
plot(super1$wstd_am, super1$critical_temp)
```

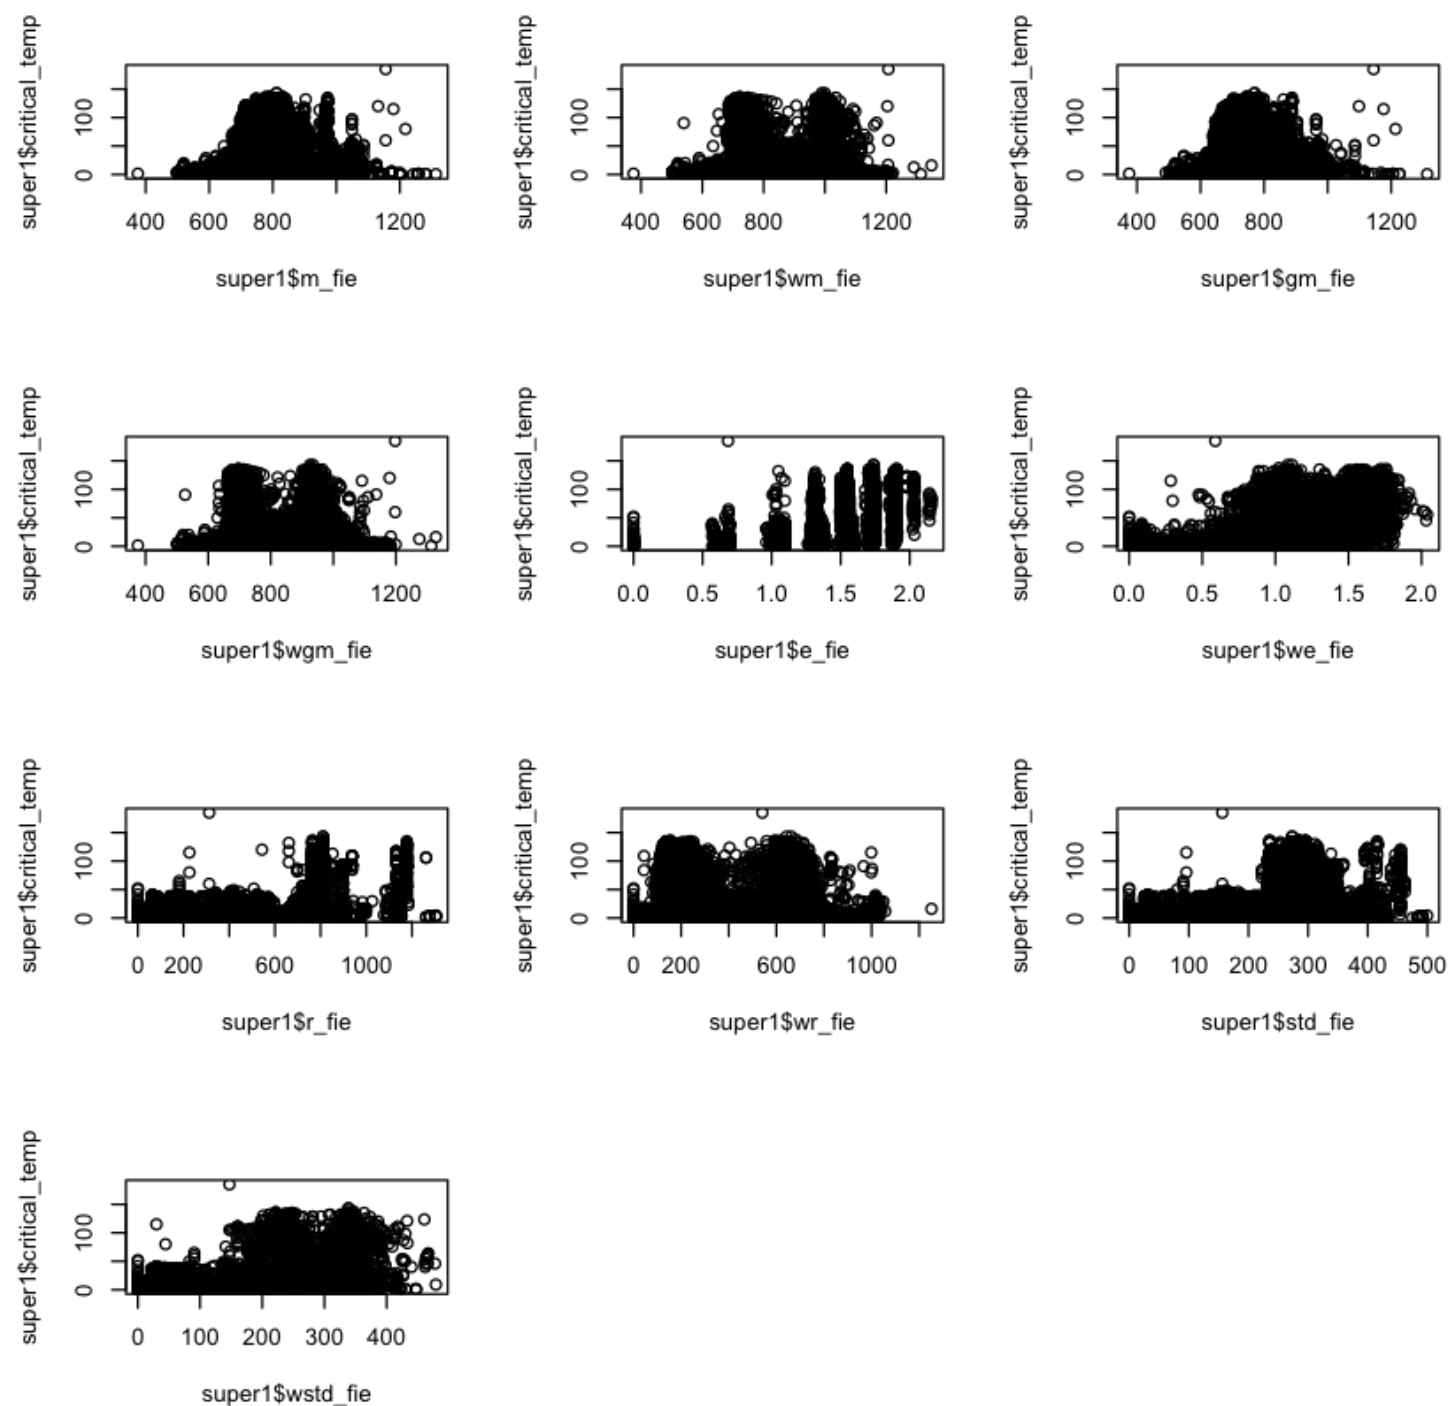


First Ionization Energy vs Critical Temp plots

In [43]:

```
#Scatter plots of First Ionization Energy based variables:
```

```
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_fie, super1$critical_temp)  
plot(super1$wm_fie, super1$critical_temp)  
plot(super1$gm_fie, super1$critical_temp)  
plot(super1$wgm_fie, super1$critical_temp)  
plot(super1$e_fie, super1$critical_temp)  
plot(super1$we_fie, super1$critical_temp)  
plot(super1$r_fie, super1$critical_temp)  
plot(super1$wr_fie, super1$critical_temp)  
plot(super1$std_fie, super1$critical_temp)  
plot(super1$wstd_fie, super1$critical_temp)
```

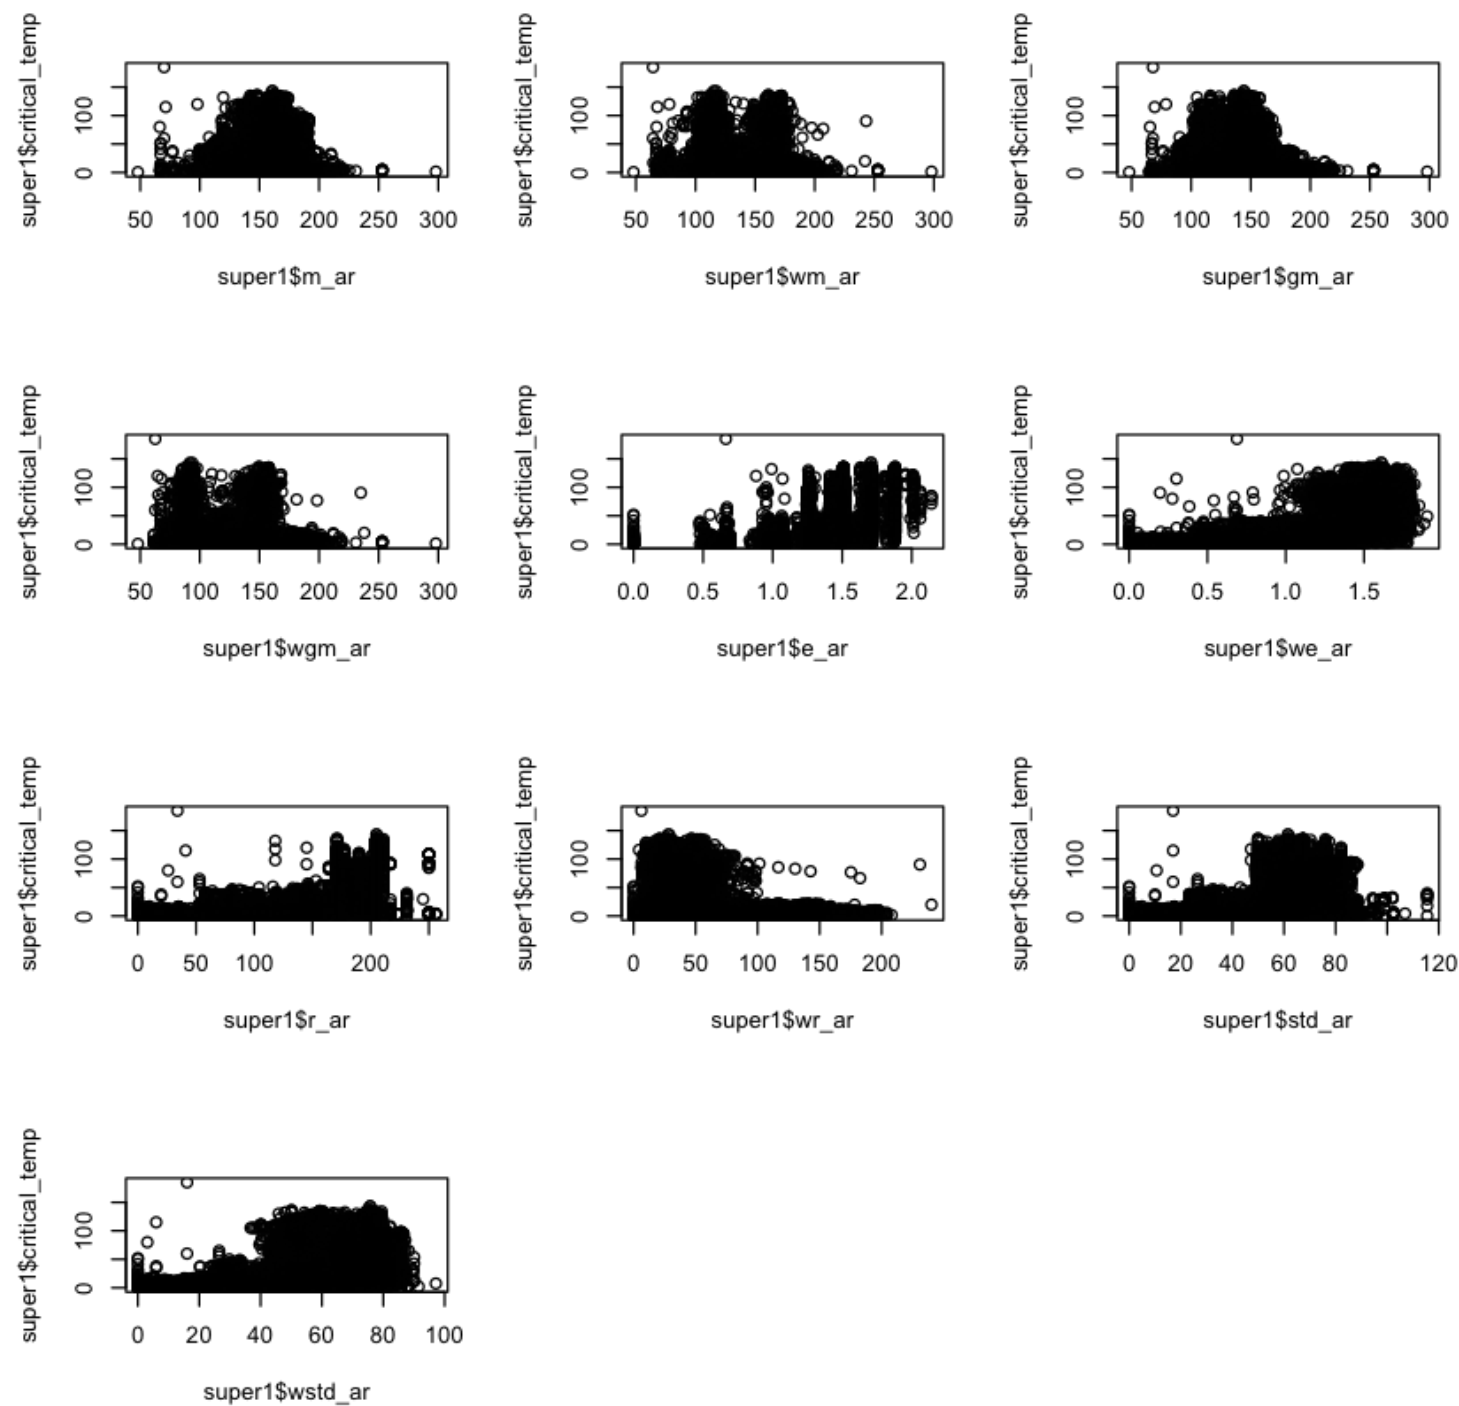


Atomic Radius vs Critical Temp Plots:

In [44]:

```
#Scatter plots of Atomic Mass based variables:
```

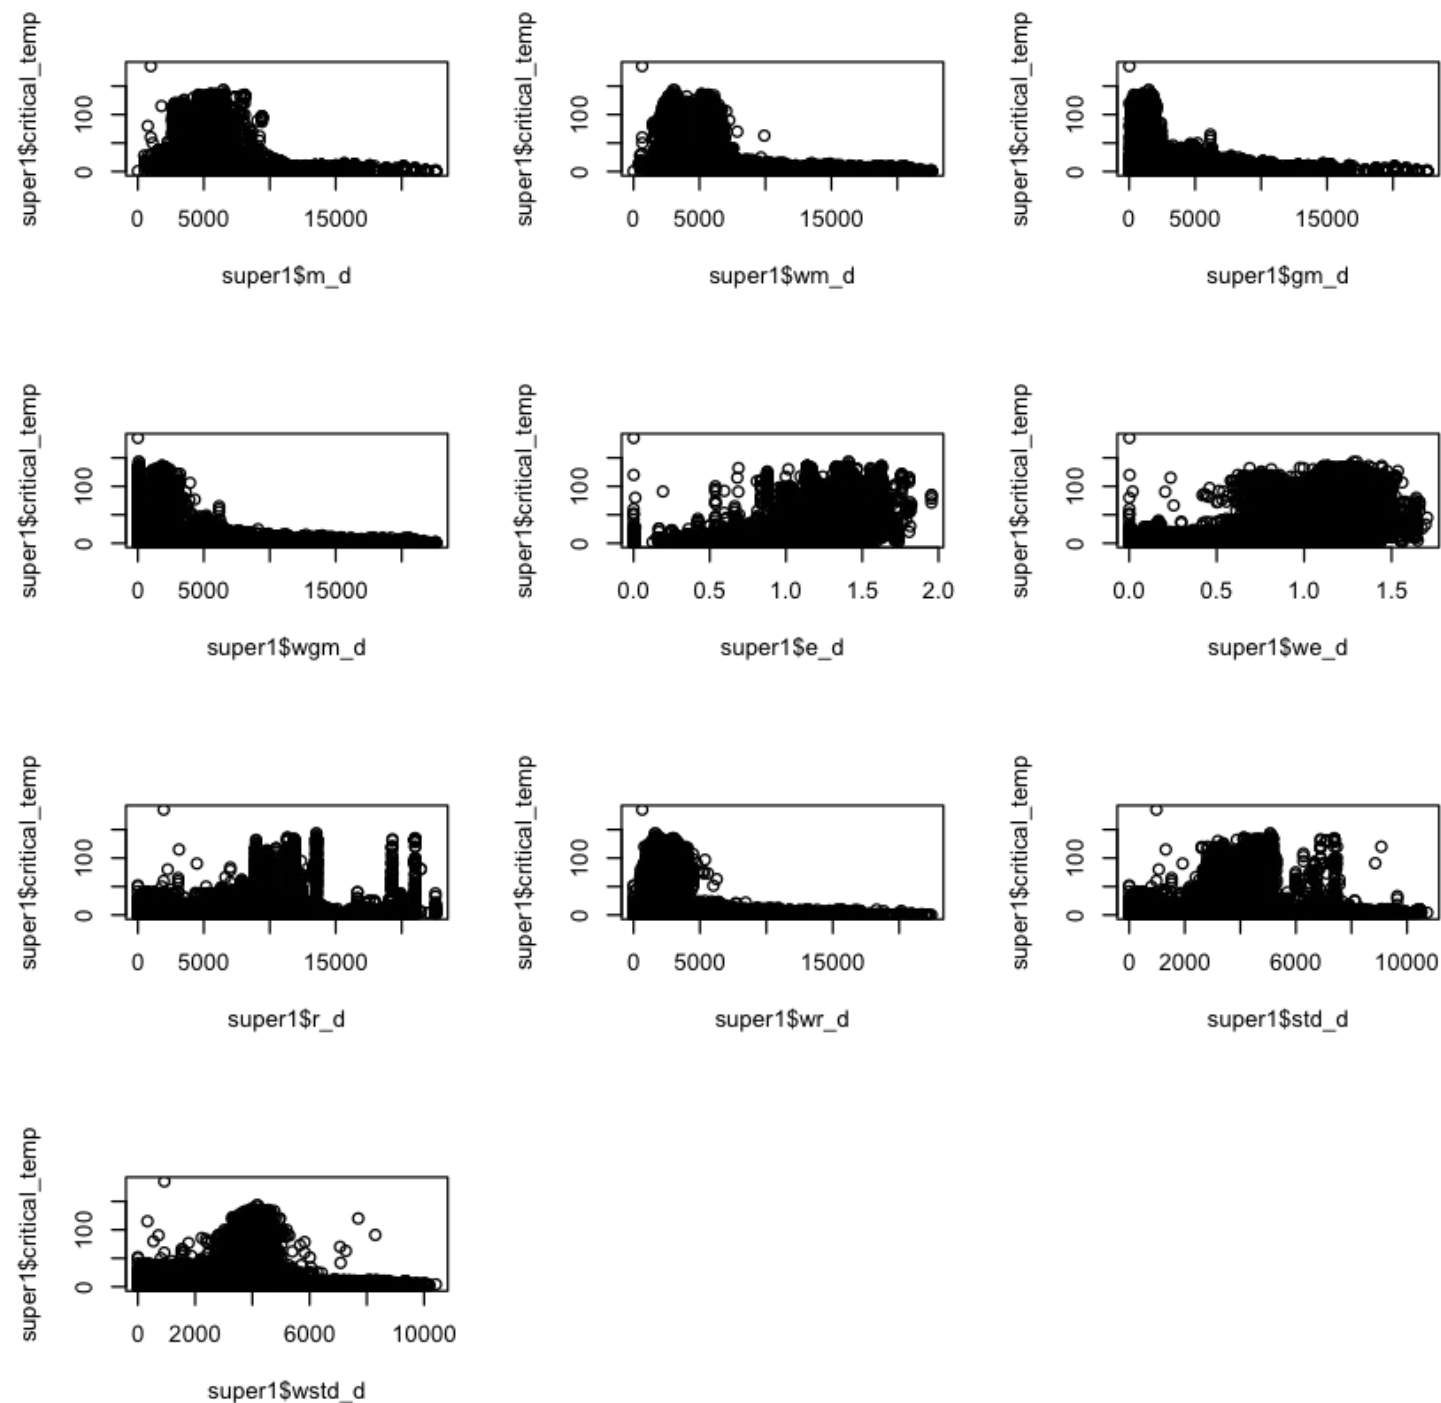
```
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_ar, super1$critical_temp)  
plot(super1$wm_ar, super1$critical_temp)  
plot(super1$gm_ar, super1$critical_temp)  
plot(super1$wgm_ar, super1$critical_temp)  
plot(super1$e_ar, super1$critical_temp)  
plot(super1$we_ar, super1$critical_temp)  
plot(super1$r_ar, super1$critical_temp)  
plot(super1$wr_ar, super1$critical_temp)  
plot(super1$std_ar, super1$critical_temp)  
plot(super1$wstd_ar, super1$critical_temp)
```



Density vs Critical Temp Plots:

In [45]:

```
#Scatter plots of Atomic Mass based variables:  
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_d, super1$critical_temp)  
plot(super1$wm_d, super1$critical_temp)  
plot(super1$gm_d, super1$critical_temp)  
plot(super1$wgm_d, super1$critical_temp)  
plot(super1$e_d, super1$critical_temp)  
plot(super1$we_d, super1$critical_temp)  
plot(super1$r_d, super1$critical_temp)  
plot(super1$wr_d, super1$critical_temp)  
plot(super1$std_d, super1$critical_temp)  
plot(super1$wstd_d, super1$critical_temp)
```



Electron Affinity vs Critical Temp Plots:

In [46]:

```
#Scatter plots of Atomic Mass based variables:
```

```
par(mfrow = c(4,3))
```

```
#dev.new(width=5, height=4)
```

```
plot(super1$m_ea, super1$critical_temp)
```

```
plot(super1$wm_ea, super1$critical_temp)
```

```
plot(super1$gm_ea, super1$critical_temp)
```

```
plot(super1$wgm_ea, super1$critical_temp)
```

```
plot(super1$e_ea, super1$critical_temp)
```

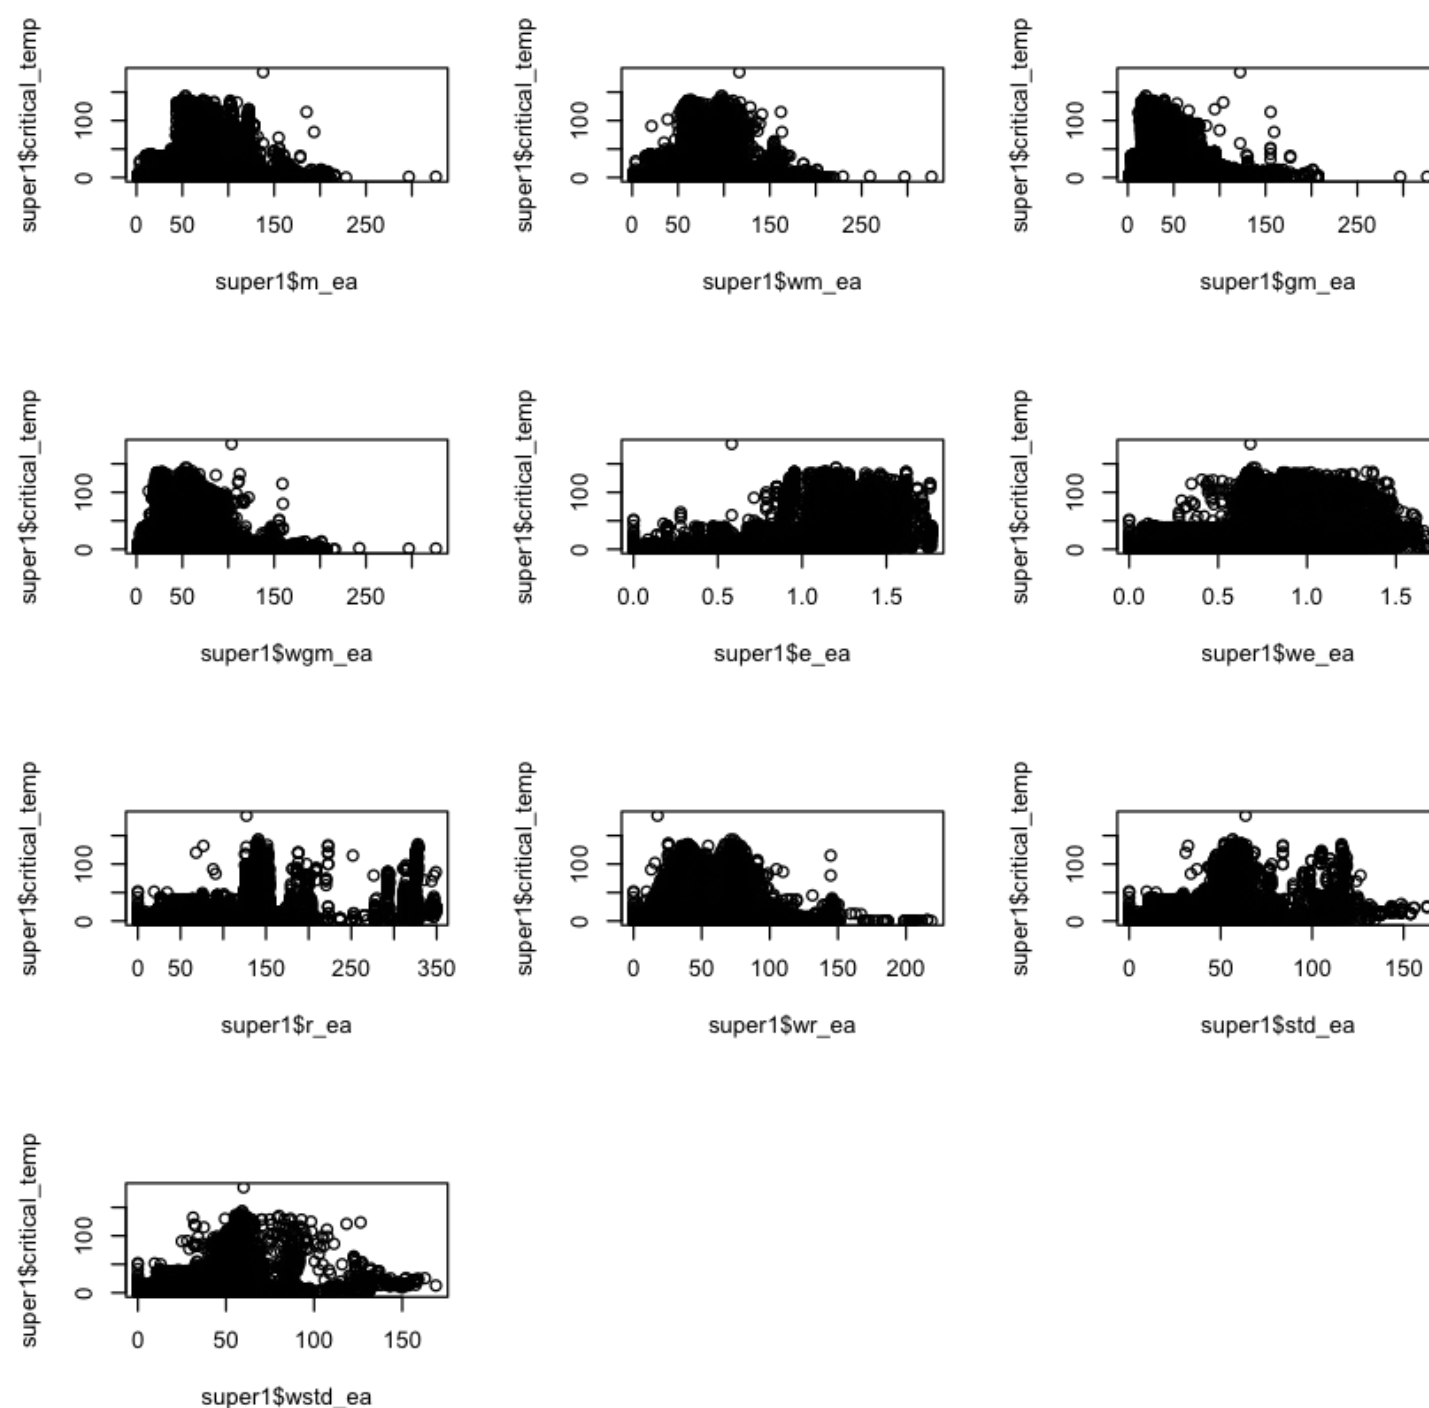
```
plot(super1$we_ea, super1$critical_temp)
```

```
plot(super1$r_ea, super1$critical_temp)
```

```
plot(super1$wr_ea, super1$critical_temp)
```

```
plot(super1$std_ea, super1$critical_temp)
```

```
plot(super1$wstd_ea, super1$critical_temp)
```

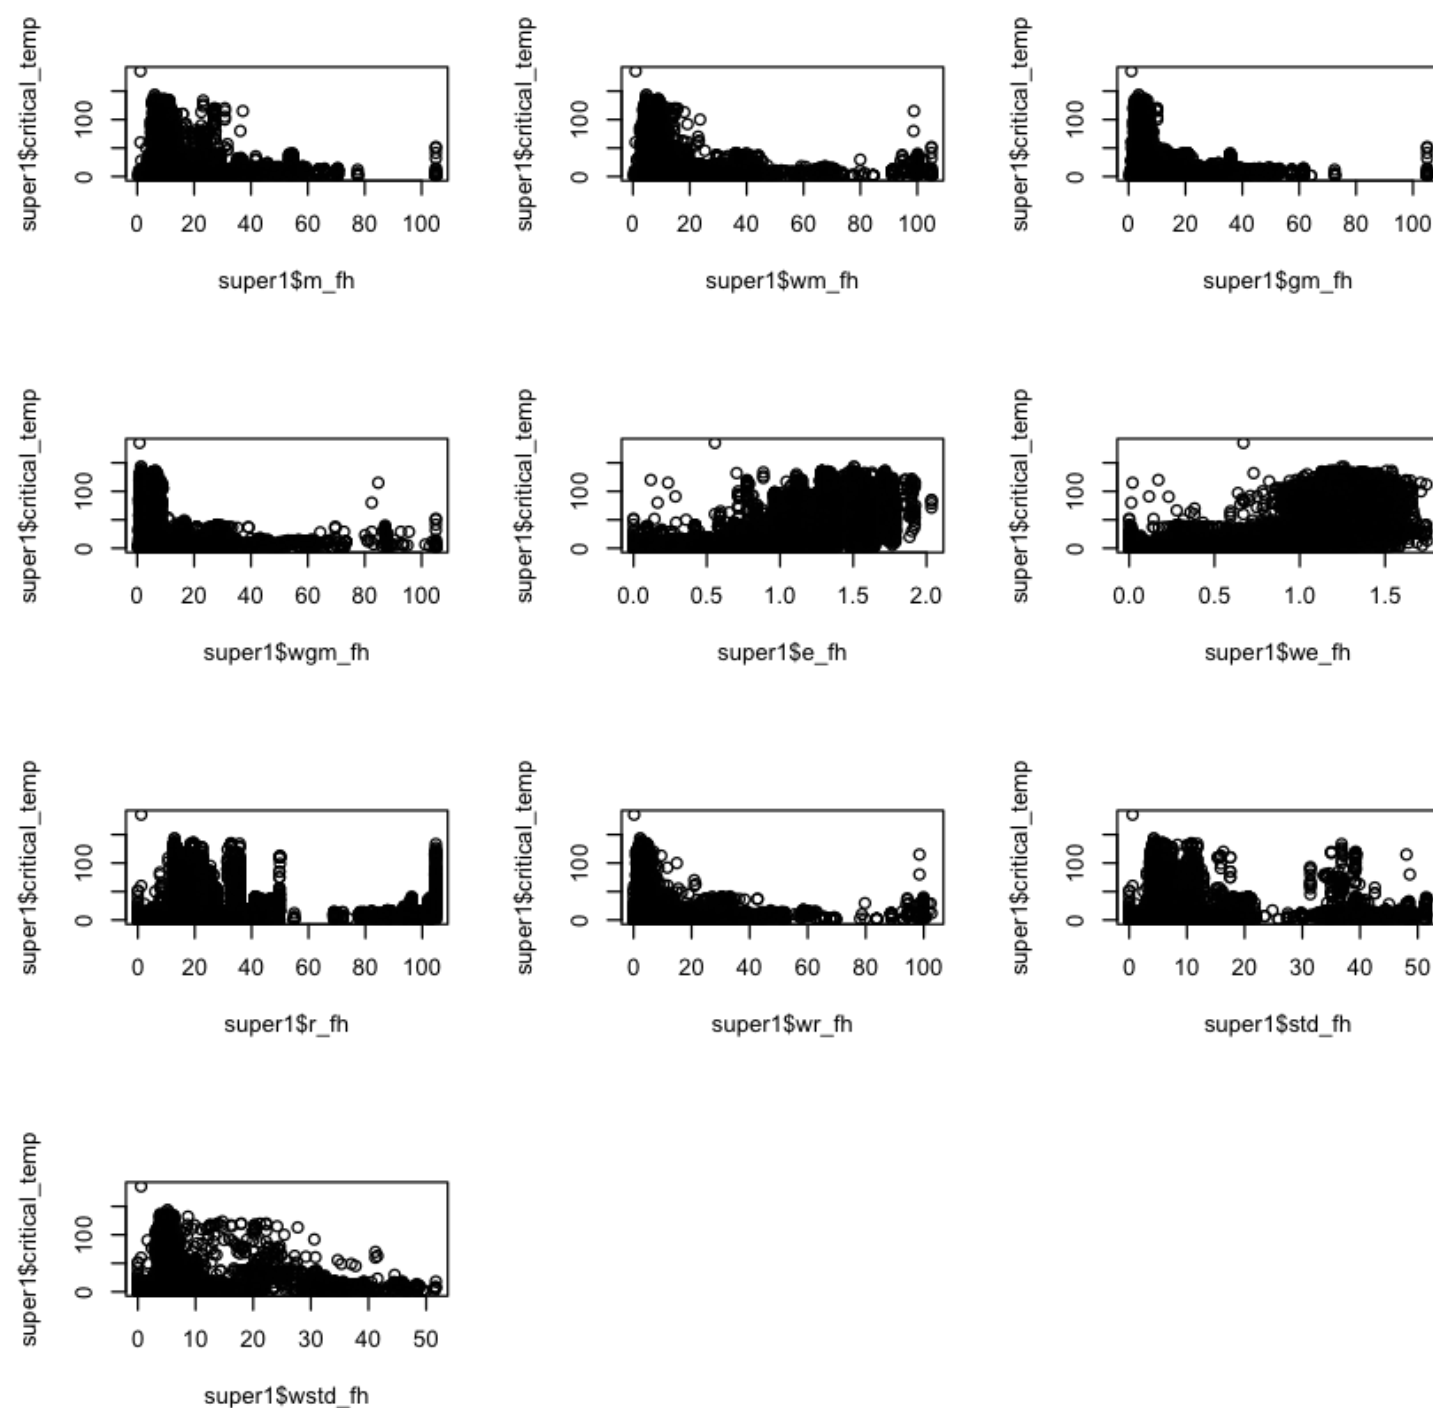


Fusion Heat vs Critical Temp Plots:

In [47]:

```
#Scatter plots of Atomic Mass based variables:
```

```
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_fh, super1$critical_temp)  
plot(super1$wm_fh, super1$critical_temp)  
plot(super1$gm_fh, super1$critical_temp)  
plot(super1$wgm_fh, super1$critical_temp)  
plot(super1$e_fh, super1$critical_temp)  
plot(super1$we_fh, super1$critical_temp)  
plot(super1$r_fh, super1$critical_temp)  
plot(super1$wr_fh, super1$critical_temp)  
plot(super1$std_fh, super1$critical_temp)  
plot(super1$wstd_fh, super1$critical_temp)
```

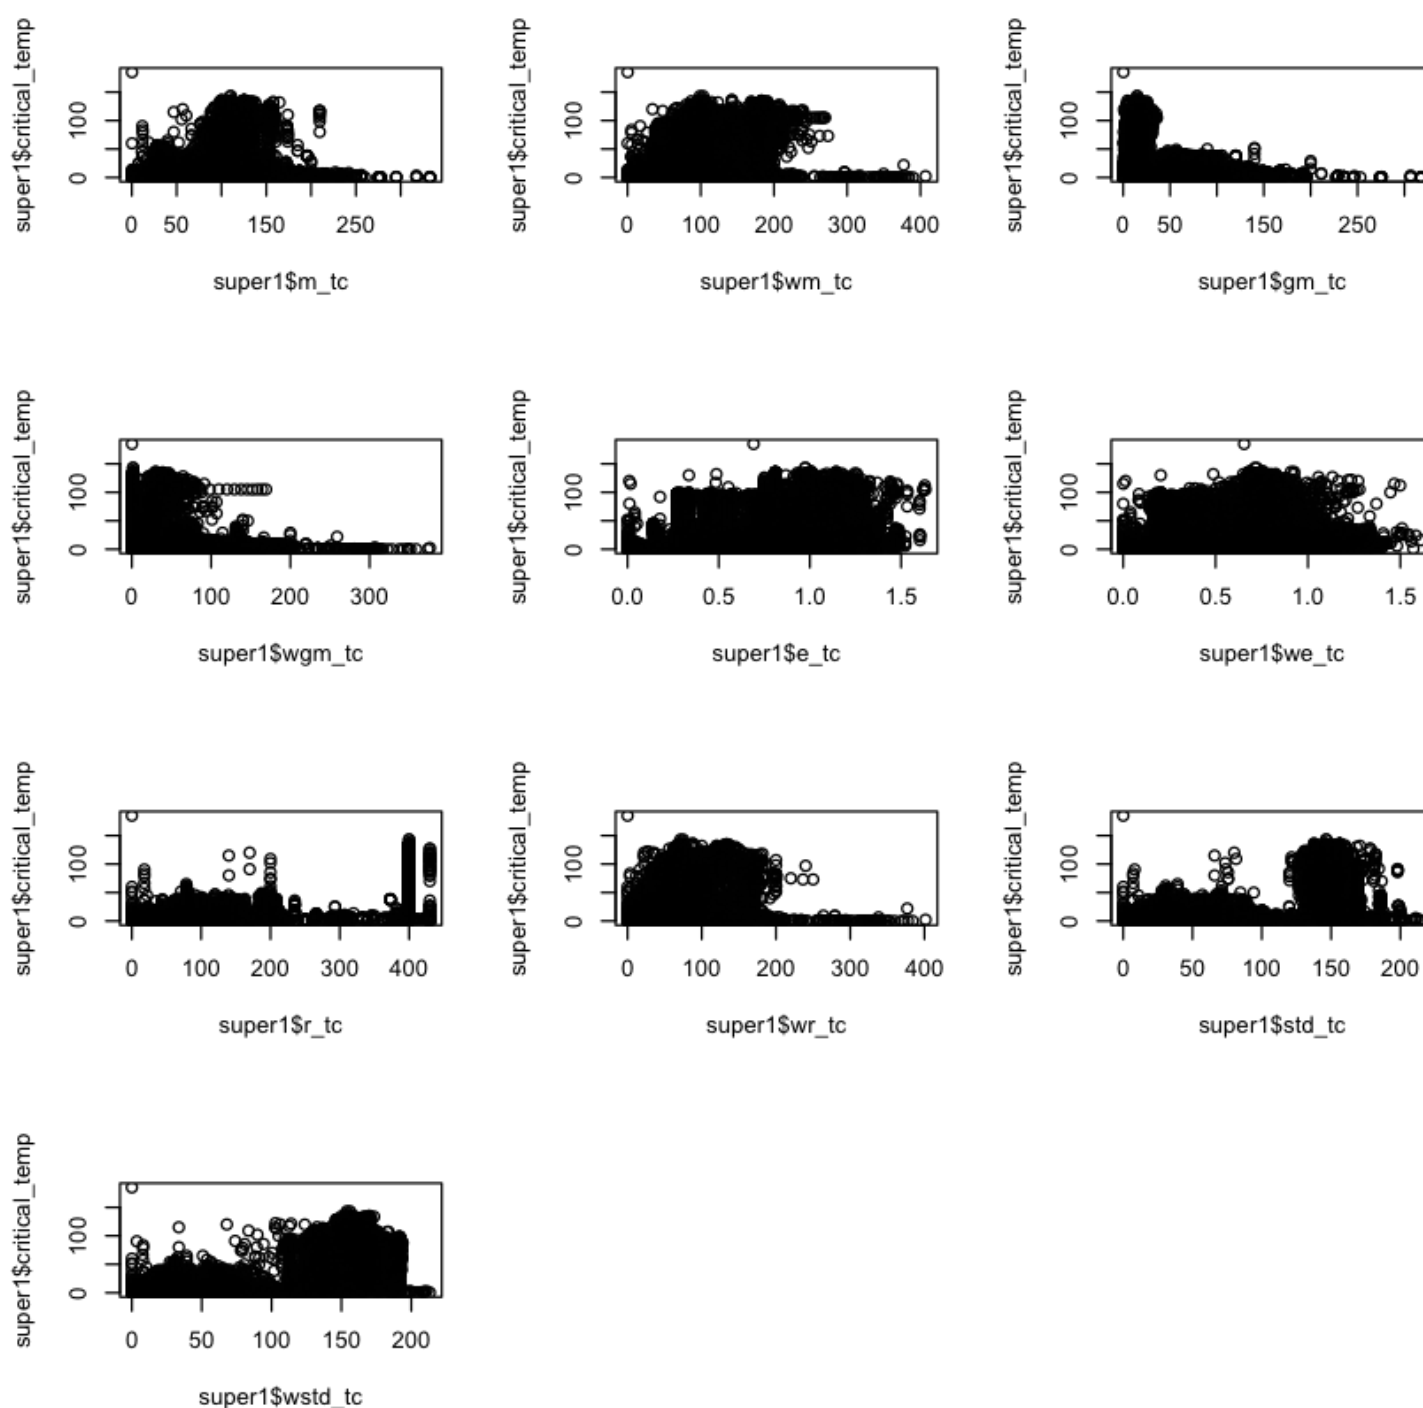


Thermal Conductivity vs Critical Temp Plots:

In [48]:

```
#Scatter plots of Atomic Mass based variables:
```

```
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_tc, super1$critical_temp)  
plot(super1$wm_tc, super1$critical_temp)  
plot(super1$gm_tc, super1$critical_temp)  
plot(super1$wgm_tc, super1$critical_temp)  
plot(super1$e_tc, super1$critical_temp)  
plot(super1$we_tc, super1$critical_temp)  
plot(super1$r_tc, super1$critical_temp)  
plot(super1$wr_tc, super1$critical_temp)  
plot(super1$std_tc, super1$critical_temp)  
plot(super1$wstd_tc, super1$critical_temp)
```

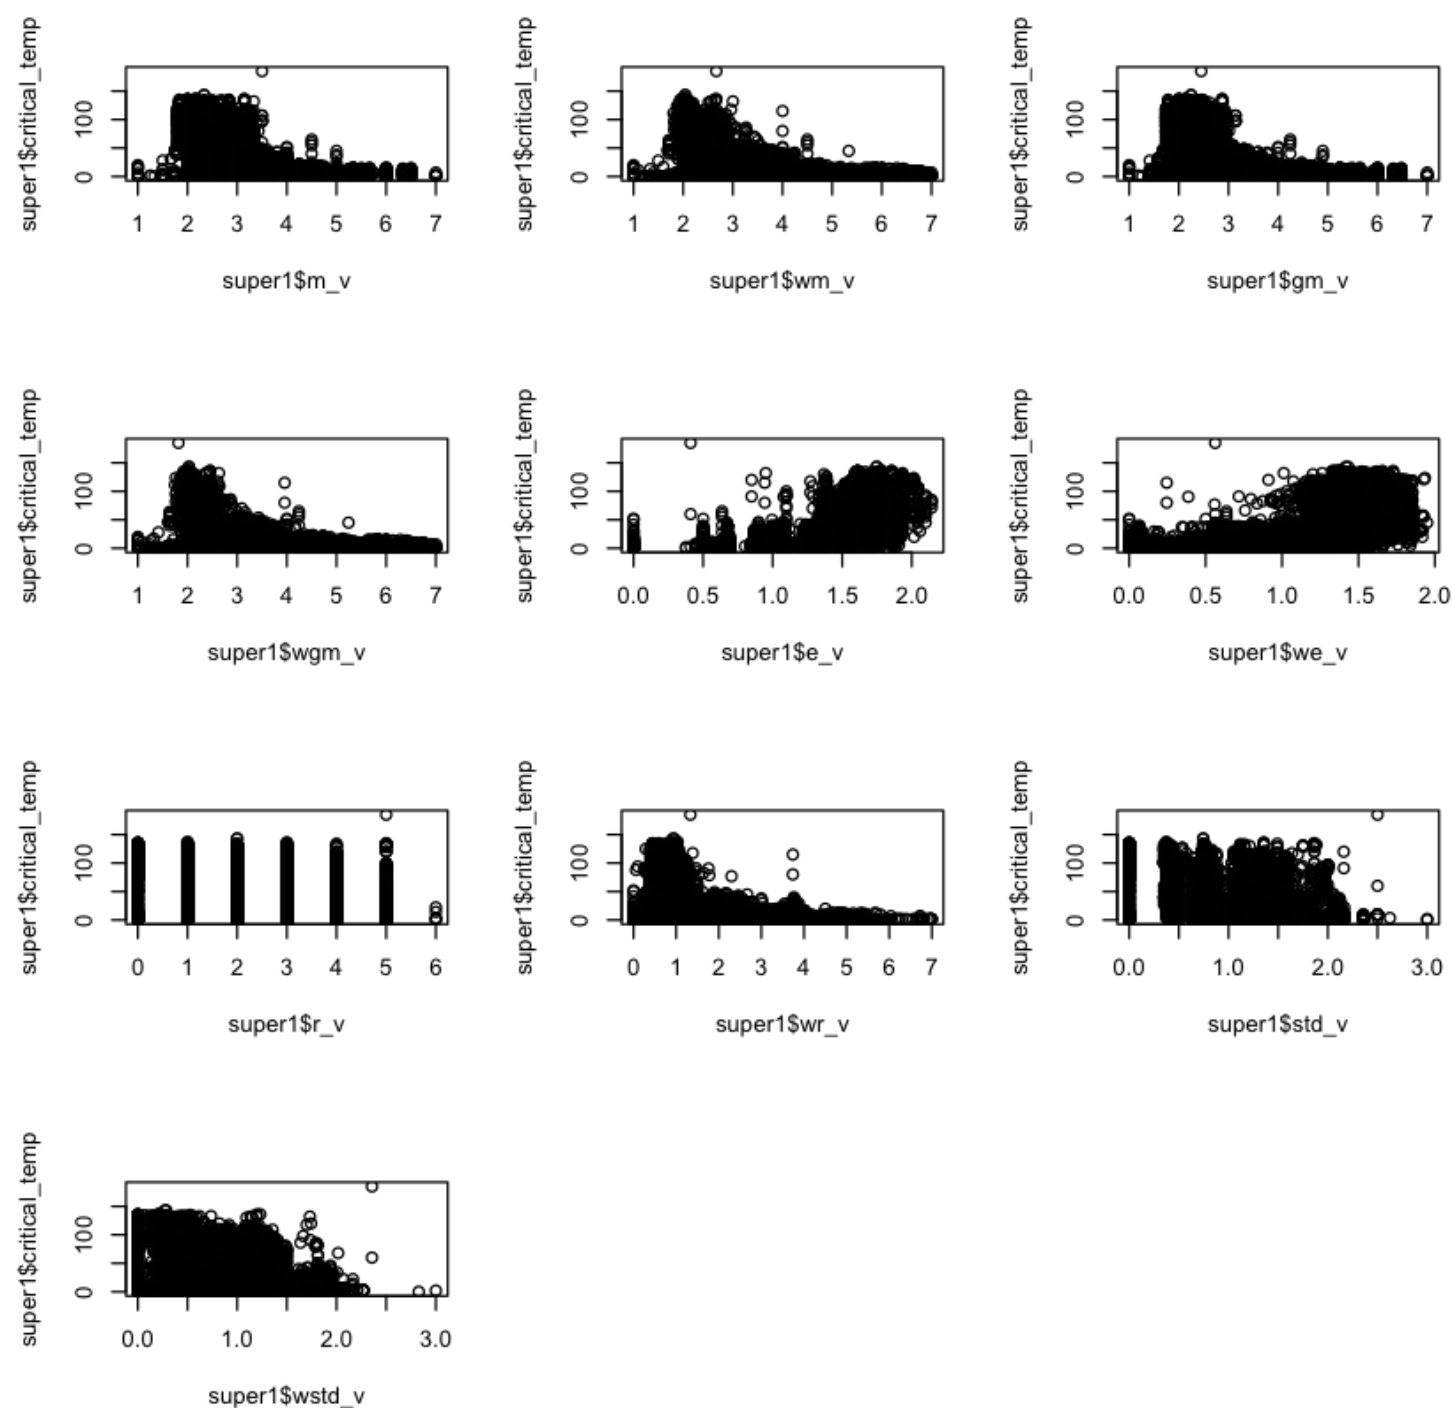


Valence vs Critical Temp Plots:

In [49]:

```
#Scatter plots of Atomic Mass based variables:
```

```
par(mfrow = c(4,3))  
#dev.new(width=5, height=4)  
plot(super1$m_v, super1$critical_temp)  
plot(super1$wm_v, super1$critical_temp)  
plot(super1$gm_v, super1$critical_temp)  
plot(super1$wgm_v, super1$critical_temp)  
plot(super1$e_v, super1$critical_temp)  
plot(super1$we_v, super1$critical_temp)  
plot(super1$r_v, super1$critical_temp)  
plot(super1$wr_v, super1$critical_temp)  
plot(super1$std_v, super1$critical_temp)  
plot(super1$wstd_v, super1$critical_temp)
```



range_valence or r_v -> an intriguing variable. Might have very low variance or zero variance.

In broad outline, EDA includes checks on data quality, calculation of summary statistics, plotting appropriate graphs as and when necessary and using complicated data-analytic techniques to derive information from the data. (Chatfield, 1986)

Since we have already seen the summary/descriptive statistics, plotted scatter plots between each variable and predictor, let us delve into the actual relationship between these predictors that can be explained by correlation between them.

Correlation between Variables:

In [50]:

```
#Computing the correlation between all the variables in the dataset
super.cor <- cor(super1)
super.cor
```

	number_of_elements	m_am	wm_am	gm_am	wgm_am
number_of_elements	1.000000000	-0.141922797	-0.353064435	-0.292968819	-0.454525138
m_am	-0.141922797	1.000000000	0.815977034	0.940298163	0.745840579
wm_am	-0.353064435	0.815977034	1.000000000	0.848241529	0.964085105
gm_am	-0.292968819	0.940298163	0.848241529	1.000000000	0.856975448
wgm_am	-0.454525138	0.745840579	0.964085105	0.856975448	1.000000000
e_am	0.939304058	-0.104000152	-0.308046011	-0.190213789	-0.37013789
we_am	0.881845150	-0.097609229	-0.412665507	-0.232183150	-0.484183150
r_am	0.682777066	0.125658599	-0.144029437	-0.175860637	-0.352175860
wr_am	-0.320293446	0.446224904	0.716623230	0.458473082	0.673082
std_am	0.513998153	0.196460491	-0.060739153	-0.121707976	-0.27407976
wstd_am	0.546391203	0.130675110	-0.089470766	-0.166042011	-0.331042011
m_fie	0.167450898	-0.285781955	-0.209296217	-0.367690186	-0.276690186
wm_fie	0.484444731	-0.222097406	-0.522595267	-0.354664335	-0.6124335
gm_fie	0.024228530	-0.240565224	-0.109490435	-0.286844286	-0.1544286
wgm_fie	0.424152224	-0.219381273	-0.508108815	-0.341584501	-0.5884501
e_fie	0.973195335	-0.166934898	-0.369772508	-0.316670438	-0.4710438
we_fie	0.719209167	-0.163564893	-0.129779058	-0.287701018	-0.22701018
r_fie	0.781226913	-0.255627735	-0.452303149	-0.431689462	-0.5750462

wr_fie	0.329623733	-0.080544826	-0.420457052	-0.155439062	-0.451
std_fie	0.674004751	-0.276560574	-0.459323462	-0.450045014	-0.578
wstd_fie	0.717830737	-0.222811916	-0.492250297	-0.390842997	-0.617
m_ar	-0.001388822	0.497663631	0.288451485	0.510867018	0.301
wm_ar	-0.422144163	0.376759801	0.660011051	0.488821849	0.720
gm_ar	-0.240444347	0.561060699	0.468457460	0.647559721	0.527
wgm_ar	-0.518255859	0.359894450	0.667112083	0.496461190	0.749
e_ar	0.972245245	-0.140034439	-0.345070536	-0.282047842	-0.441
we_ar	0.904120682	-0.147603805	-0.400482727	-0.311701386	-0.514
r_ar	0.768060022	-0.270694807	-0.524860727	-0.460197290	-0.645
wr_ar	-0.371349583	0.141100325	0.363882464	0.240295641	0.432
std_ar	0.624810298	-0.326402906	-0.551140623	-0.512841330	-0.665
wstd_ar	0.695088577	-0.280439751	-0.554820067	-0.462396711	-0.681
m_d	-0.418674700	0.756861118	0.749260525	0.779757034	0.740
wm_d	-0.507895375	0.608934913	0.842664701	0.677131018	0.852
gm_d	-0.630504092	0.596484643	0.712815440	0.728477184	0.789
wgm_d	-0.649882311	0.525588122	0.767010936	0.663642015	0.843
e_d	0.871831749	-0.043415596	-0.246377142	-0.125672127	-0.300
we_d	0.767078421	0.026324752	-0.195893891	-0.093880853	-0.273
r_d	0.413485633	0.198067091	-0.002868148	-0.024974646	-0.163
wr_d	-0.355389487	0.342390738	0.585686944	0.368142682	0.576
std_d	0.210723871	0.245042403	0.103156510	0.037866071	-0.048
wstd_d	0.334072073	0.180942527	0.009921269	-0.037299230	-0.174
m_ea	-0.119302572	0.088229536	0.147303238	0.079376432	0.119
wm_ea	0.195607767	0.061103034	-0.096426785	-0.006352581	-0.158
gm_ea	-0.356067248	0.189282435	0.272260819	0.219650894	0.274
wgm_ea	-0.052883973	0.134382340	0.021876579	0.111858038	-0.011
e_ea	0.877304020	-0.091539414	-0.290219664	-0.238002465	-0.395

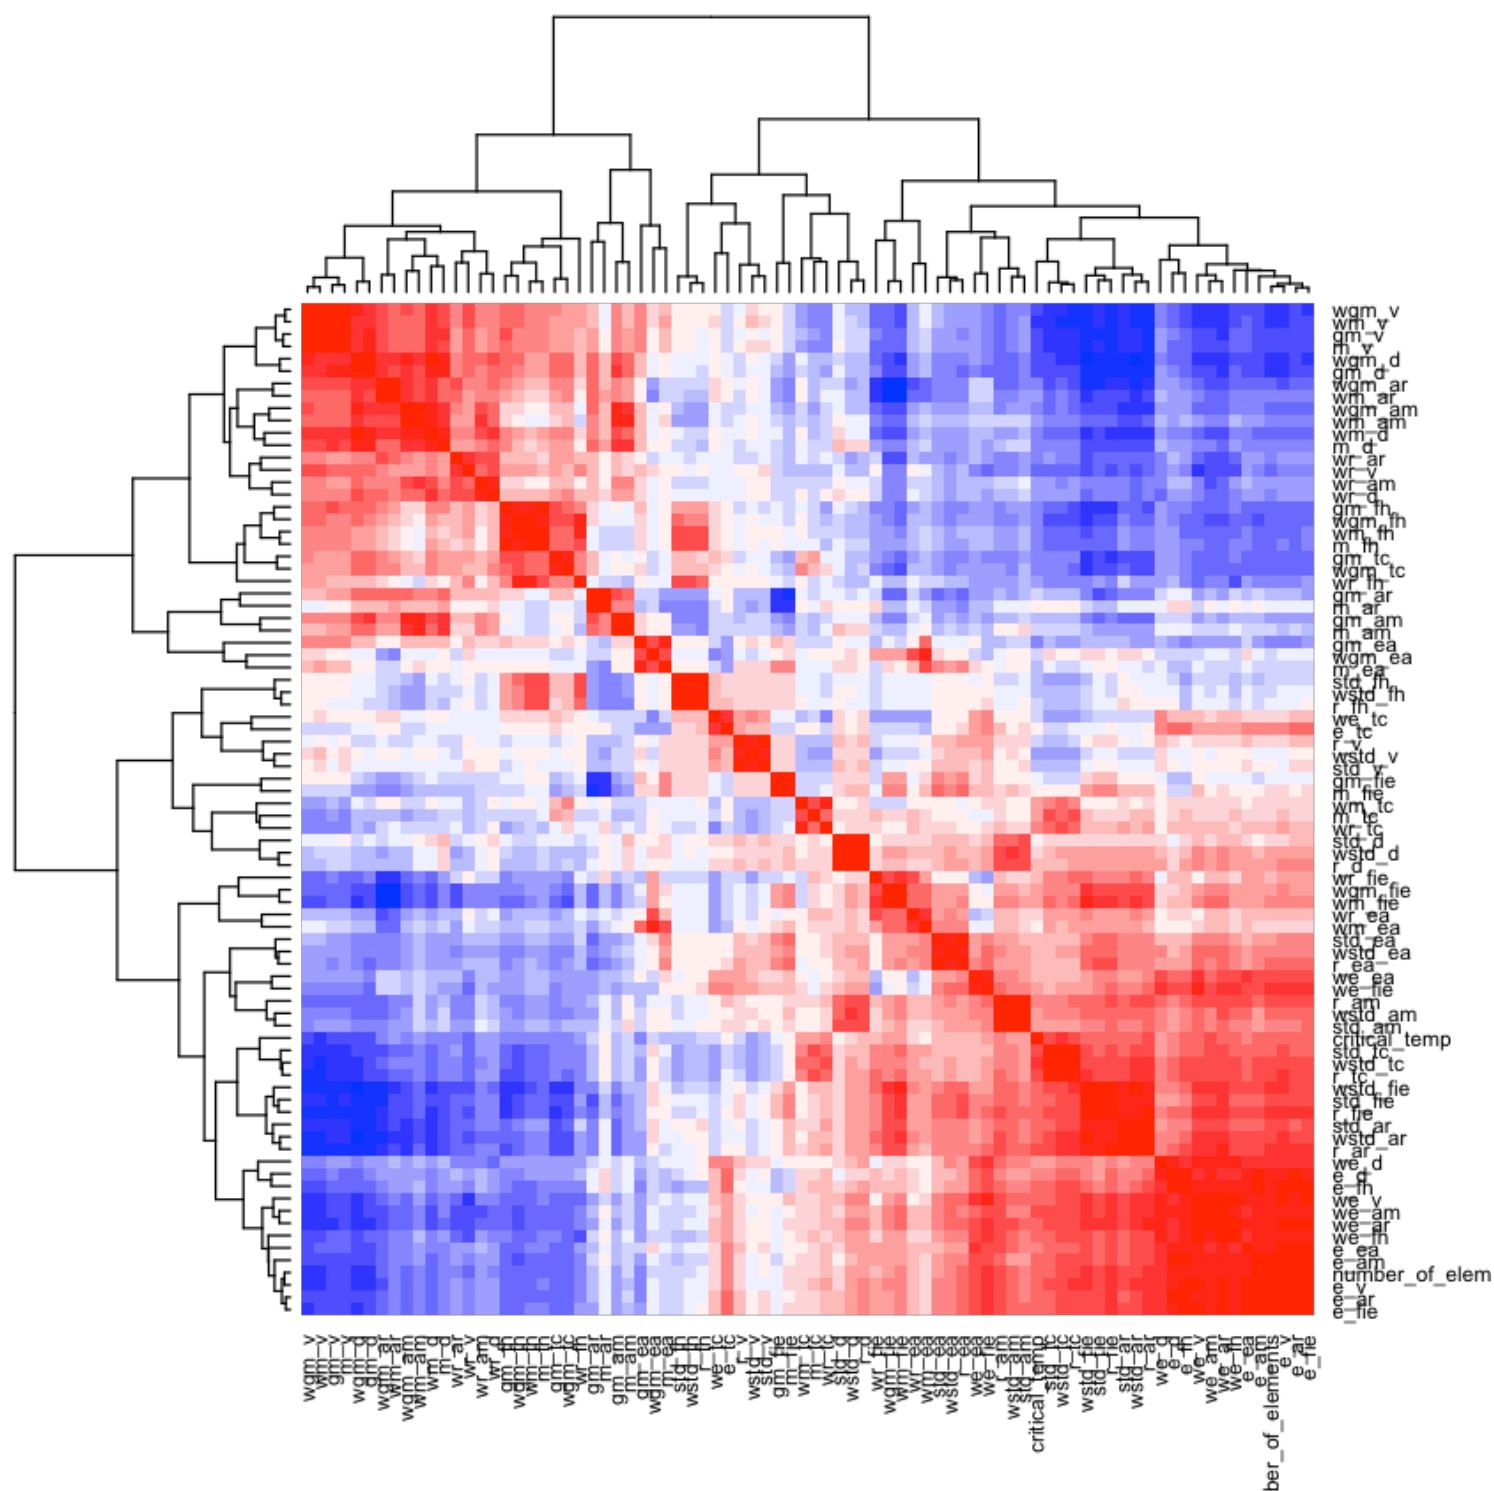
we_ea	0.625798404	-0.107650648	-0.093796393	-0.224756609	-0.194
r_ea	0.531540406	-0.187068587	-0.225890340	-0.284097947	-0.300
wr_ea	0.241411233	0.010235114	-0.204480413	-0.055316460	-0.246
std_ea	0.423738261	-0.164960192	-0.197728679	-0.249926458	-0.262
wstd_ea	0.480812544	-0.133100550	-0.210756959	-0.238027624	-0.291
m_fh	-0.437624455	-0.137668902	0.006730077	-0.092243615	0.054
wm_fh	-0.449271665	-0.135428565	0.014681004	-0.089138720	0.072
gm_fh	-0.514251754	0.014818495	0.164239431	0.086599422	0.219
wgm_fh	-0.519109340	-0.043002734	0.120043990	0.024199132	0.189
e_fh	0.900759345	-0.008498640	-0.225286701	-0.126798354	-0.313
we_fh	0.860478819	-0.028540944	-0.237218403	-0.171927681	-0.349
r_fh	0.005734172	-0.347581882	-0.283419738	-0.384838353	-0.279
wr_fh	-0.371787677	-0.167527922	-0.070411491	-0.128757736	-0.006
std_fh	-0.113361228	-0.337968801	-0.253910567	-0.361244227	-0.239
wstd_fh	-0.074796377	-0.335778336	-0.272805875	-0.372665857	-0.277
m_tc	0.227655893	-0.158265641	-0.236418136	-0.190936638	-0.248
wm_tc	0.206068511	-0.065988791	-0.058075067	-0.104939744	-0.056
gm_tc	-0.485323530	0.006004074	0.184990043	0.110768515	0.271
wgm_tc	-0.469206128	0.056394180	0.250226215	0.131642343	0.322
e_tc	0.501871343	-0.100076682	-0.076936360	-0.116455185	-0.076
we_tc	0.207065208	-0.098220524	0.025638296	-0.104643666	0.020
r_tc	0.696059989	-0.114537871	-0.376573091	-0.243464920	-0.464
wr_tc	0.316771509	-0.027789696	-0.108511530	-0.095661466	-0.129
std_tc	0.602018066	-0.110657807	-0.362511650	-0.233587261	-0.447
wstd_tc	0.665579856	-0.110856410	-0.350993285	-0.232079144	-0.431
m_v	-0.609412358	0.374098807	0.534450406	0.487020743	0.599

wm_v	-0.648550851	0.304683244	0.545587036	0.427960618	0.614
gm_v	-0.618512449	0.392152739	0.539780205	0.511507826	0.608
wgm_v	-0.659267738	0.321398940	0.548980987	0.450356645	0.623
e_v	0.967832451	-0.156786430	-0.375718236	-0.306246222	-0.477
we_v	0.892559475	-0.145610386	-0.331024627	-0.307662306	-0.448
r_v	0.231873976	-0.107449727	-0.039155397	-0.165010264	-0.078
wr_v	-0.447769825	0.168632512	0.330903813	0.272302621	0.409
std_v	0.105365152	-0.080278563	-0.003681475	-0.124626855	-0.033
wstd_v	0.035216288	-0.081252677	0.077322868	-0.117335886	0.030
critical_temp	0.601068571	-0.113523246	-0.312272020	-0.230345375	-0.369

In [51]:

```
palette = colorRampPalette(c("blue", "white", "red")) (20)
heatmap(x = super.cor, col = palette, symm = TRUE)

#Red -> high positive correlation
#Blue -> high negative correlation
#White -> little to now correlation
```



In the above heat map, RED indicates high positive correlation, BLUE indicates high negative correlation and WHITE indicates no correlation between the predictors themselves.

Now that we have performed significant exploration of the data, we can move to the model development stage where we will be looking at a few regression methods along with appropriate feature selection methodologies.

3. Model Development

"The aim of regression models is to model the variation of a quantitative response variable y in terms of the variation of one or several explanatory variables (x_1, \dots, x_p) " (Härdle W.K., Simar L., 2012)

In the steps that follow, we will be formulating a few models which could help us predict the Critical Temperature of a superconductor. Variable subset selection will be done independently for each model based on the type of model being used .

SPLITTING THE DATA:

- Randomly setting the seed to a number so that a random split is generated everytime this code is run. This ensures that the data is fully shuffled everytime for better prediction purposes.
- Then, `sample()` function is used to sample the data, in this case 1 sample containing all the data, but in a shuffled manner to ensure a good standard of practise.
- Finally, the shuffled data is assigned to **shuf_super** which will now be split into training and testing sets.

In [52]:

```
set.seed(62)
row <- sample(nrow(super1))
shuf_super <- super1[row,]
```

"At the end of the training process, the final model should predict correct outputs for the input samples from T , but it should also be able to generalize well to previously unseen data. Poor generalization can be characterized by over-training. If the model over-trains, it just memorizes the training examples and it is not able to give correct outputs also for patterns that were not in the training dataset." (Reitermanova, 2010)

Thus, splitting a dataset into 2, one for training and one for testing is extremely important to avoid the errors and reduce the "Bias Variance Tradeoff".

- We will split the data in a ratio of 4:1. i.e. **80% Training Data** and **20% Testing Data**.
- For this purpose, the `createDataPartition()` function is used.

In [53]:

```
#Determining the row to split the data on
split <- createDataPartition(shuf_super$critical_temp, p=0.8, list=FALSE)

#Creating the train dataset
train <- shuf_super[split,]

#Creating the test dataset
test <- shuf_super[-split,]
```

- Let us also divide the train and test data into a subset of the variables (all 81 of them) and the target variable (1 of them)

In [54]:

```
train.predictors <- train[,-82]
train.target <- train[,82]
test.predictors <- test[,-82]
test.target <- test[,82]
```

FEATURE SELECTION:

"Generally, features are characterized as:

1. Relevant: These are features which have an influence on the output and their role can not be assumed by the rest
2. Irrelevant: Irrelevant features are defined as those features not having any influence on the output, and whose values are generated at random for each example.
3. Redundant: A redundancy exists whenever a feature can take the role of another (perhaps the simplest way to model redundancy)." (Karagiannopoulos et. al, no date)

Thus, we will be using several methods while developing the models below which will help us select the best features for implementation. Feature selection can be done by: **Filter, Wrapper or Embedded** methods.

MODEL 1: MULTIPLE LINEAR REGRESSION WITH FILTER METHODS

"Multiple linear regression attempts to model the relationship between two or more explanatory variables and a response variable by fitting a linear equation to observed data. Every value of the independent variable x is associated with a value of the dependent variable y ."

Variable selection can be done for model development in many methods. For this linear regression model, we will be looking into some filter methods which will help us rank the predictors in order of relevance and helps us measure a particular feature's influence on the target, eg: correlation between variables and target. Irrelevancy can be eliminated by discarding the features that would have little to no influence on the model such as the ones with extremely low variance, as low variability affects regression models. Finally, redundancy can be eliminated by targeting multicollinearity within the features. We will look into this further:

Correlation Criteria (Relevance)

- **Correlation** is the measure which indicates the relationship between two variables. In our dataset, since we have a large number of predictor variables, it helps to identify the ones which are highly correlated to try to identify the type of relationship between these variables for better model development purposes.
- This will be checked by using the `cor()` function which returns the Pearson's correlation coefficient.

In [55]:

```
co <- data.frame(cor(train[1:81],train[82]))
co <- setDT(co, keep.rownames = TRUE)[ ]
new <- co[order(abs(critical_temp)) ]
#Displaying the top 30% of the features with the highest correlation
tail(new,24)
```

rn	critical_temp
e_am	0.5435378
std_fie	0.5438187
gm_d	-0.5439631
e_fh	0.5523997
e_ar	0.5580286
we_fh	0.5612418
std_ar	0.5623458
e_fie	0.5669584
gm_v	-0.5751778
wstd_fie	0.5824024
we_v	0.5903004
e_v	0.5978685
wstd_ar	0.6004421
number_of_elements	0.6005712
m_v	-0.6018994
r_fie	0.6021512
we_ar	0.6026045
wgm_v	-0.6178228
we_am	0.6267691
wm_v	-0.6343775
std_tc	0.6528679
r_ar	0.6556092
r_tc	0.6866783
wstd_tc	0.7228292

Here, we have the list of the top 30 % of the variables ordered by their correlation with the target variable. For the development of the initial model, we can try to go ahead with these features and perform more analysis as we go ahead:

In [56]:

```
vars<- tail(new$rn,24)
vars[25] = 'critical_temp'
#Function to extract the specific columns from the dataframe
extract_columns <- function(data,var) {
  extracted_data <- data %>%
    select_(.dots = var)
  return(extracted_data)
}

#Passing the 'train' dataframe with all variables and retrieving a 'train1' dataframe with selected 30 vars
train1 <- extract_columns(train,vars)
```

In [57]:

```
#head(train1)
dim(train1)
```

17011 25

Variance Criteria (Irrelevance)

- **Variance** of a predictor is one of the important factors to consider during the process of feature selection. Ideally, you would like all of your predictor values to be variant to provide the best and stable outcome for your prediction, but sometimes that is not the case.
- *"These so-called “near zero-variance predictors” can cause numerical problems during resampling for some models, such as linear regression"* (Kuhn, 2008). Since the first model we are computing is a multiple-linear regression model, it is safer to eliminate these predictors, if any.
- This will be done by using the `nearZeroVar()` function from the `caret` package.

In [58]:

```
nzv <- nearZeroVar(train1, saveMetrics= TRUE)
nzv[nzv$nzv,][1:5,]
```

	freqRatio	percentUnique	zeroVar	nzv
r_tc	24.33333	2.580683	FALSE	TRUE
NA	NA	NA	NA	NA
NA.1	NA	NA	NA	NA
NA.2	NA	NA	NA	NA
NA.3	NA	NA	NA	NA

We can see from the above output that there is only one predictor, r_tc (range_ThermalConductivity), which has a near zero variance. Let us remove this variable from our dataset and continue with the analysis.

In [59]:

```
nzv <- nearZeroVar(train1)
trains <- train1[, -nzv]
dim(trains)
```

```
17011  24
```

Redundancy Criteria (Multicollinearity)

In [60]:

```
cor(trains)
```

	e_am	std_fie	gm_d	e_fh	e_ar	we_fh
e_am	1.0000000	0.6031344	-0.5478821	0.9288372	0.9721652	0.8451658
std_fie	0.6031344	1.0000000	-0.8089333	0.4775156	0.6418863	0.5589844
gm_d	-0.5478821	-0.8089333	1.0000000	-0.4493800	-0.5954943	-0.5028169
e_fh	0.9288372	0.4775156	-0.4493800	1.0000000	0.9312397	0.8819167
e_ar	0.9721652	0.6418863	-0.5954943	0.9312397	1.0000000	0.8679841
we_fh	0.8451658	0.5589844	-0.5028169	0.8819167	0.8679841	1.0000000
std_ar	0.5692775	0.8760483	-0.8031595	0.4852203	0.6127398	0.5350553
e_fie	0.9644950	0.6733612	-0.6247972	0.9176415	0.9977341	0.8647613
gm_v	-0.5573122	-0.7402440	0.8260959	-0.4951590	-0.5931249	-0.5066060
wstd_fie	0.6611082	0.9343778	-0.7951378	0.5753409	0.6855607	0.6525401
we_v	0.8616050	0.6929071	-0.6229998	0.8246655	0.8996565	0.9083106
e_v	0.9633610	0.6918555	-0.6353083	0.9221954	0.9895608	0.8664223
wstd_ar	0.6481837	0.8590990	-0.7979871	0.5799674	0.6804241	0.6327714
number_of_elements	0.9389885	0.6750714	-0.6311550	0.9016500	0.9720279	0.8606433
m_v	-0.5538916	-0.7215544	0.8115204	-0.5046794	-0.5815610	-0.5070151
r_fie	0.7045358	0.9818244	-0.8157000	0.5961189	0.7412939	0.6578987
we_ar	0.8799557	0.7218116	-0.6538538	0.8415550	0.9142093	0.9077977
wgm_v	-0.5953249	-0.7138098	0.7953745	-0.5633628	-0.6256183	-0.6003254
we_am	0.8894220	0.6626746	-0.5965650	0.8350306	0.8902726	0.8733723
wm_v	-0.5888600	-0.6973042	0.7786380	-0.5612617	-0.6122889	-0.5906187
std_tc	0.5989179	0.6114040	-0.5868615	0.5990330	0.5961358	0.5963150
r_ar	0.7098016	0.8819359	-0.8155271	0.6432901	0.7461311	0.6689321
wstd_tc	0.6519329	0.6257707	-0.6041828	0.6536395	0.6456540	0.6392261
critical_temp	0.5435378	0.5438187	-0.5439631	0.5523997	0.5580286	0.5612418

We can see from the above correlation matrix that the variables are highly correlated with each other which could be a cause of concern as multicollinearity is not a property you would like to have in your dataset. Let us check what we can do in order to explain or repair this outcome.

In [61]:

```
#Setting an extremely high arbitrary threshold value to check the vif all variables.  
vifstep(trains[1:24],th=2000)
```

No variable from the 24 input variables has collinearity problem.

The linear correlation coefficients ranges between:

min correlation (e_fh ~ gm_d): -0.4467544

max correlation (e_fie ~ e_ar): 0.9977541

----- VIFs of the remained variables -----

	Variables	VIF
1	e_am	32.251656
2	std_fie	163.056012
3	gm_d	5.367897
4	e_fh	22.568991
5	e_ar	1150.609920
6	we_fh	12.860893
7	std_ar	100.893969
8	e_fie	1691.536609
9	gm_v	925.503303
10	wstd_fie	53.674596
11	we_v	28.372194
12	e_v	986.960003
13	wstd_ar	48.371481
14	number_of_elements	41.017526
15	m_v	912.935437
16	r_fie	173.297078
17	we_ar	62.611072
18	wgm_v	611.933738
19	we_am	25.695426
20	wm_v	609.087256
21	std_tc	17.356693
22	r_ar	110.953949
23	wstd_tc	21.202031
24	critical_temp	2.819749

VIF values are said to be high if they exceed >5. But here, the VIF values seem to be completely out of bounds. This high correlation/multicollinearity could be explained by the fact that the variables have been derived from 8 basic properties.

"The fact that some or all predictor variables are correlated among themselves does not, in general, inhibit our ability to obtain a good fit nor does it tend to affect inferences about mean responses or predictions of new observations." —Applied Linear Statistical Models, p289, 4th Edition.

Thus, we could try examining the future model by eliminating variables which have VIF>150 in our dataset, just out of intuition.

In [62]:

```
v <-vifstep(trains[1:24],th=150)
v
```

5 variables from the 24 input variables have collinearity problem:

e_fie gm_v wgm_v r_fie e_ar

After excluding the collinear variables, the linear correlation coefficients ranges between:

min correlation (e_fh ~ gm_d): -0.4537581

max correlation (number_of_elements ~ e_v): 0.9691109

----- VIFs of the remained variables -----

	Variables	VIF
1	e_am	22.441076
2	std_fie	32.187058
3	gm_d	5.130122
4	e_fh	18.156723
5	we_fh	12.723701
6	std_ar	70.237477
7	wstd_fie	47.288675
8	we_v	24.944791
9	e_v	52.915476
10	wstd_ar	47.261214
11	number_of_elements	25.788545
12	m_v	16.780625
13	we_ar	43.187501
14	we_am	20.315218
15	wm_v	16.026122
16	std_tc	15.525964
17	r_ar	81.321476
18	wstd_tc	18.871100
19	critical_temp	2.717609

Thus, now we are left with 18 variables and one target in this dataset, by arbitrarily setting the limit as 150. The mean VIF has also drastically reduced as it is a stepwise function and the removal of one variables affects the VIF of another. We are going to exclude the variables with the collinearity problem from the `trains` dataset and proceed to compute a linear regression model.

In [63]:

```
trains <- exclude(trains,v)
```

In [64]:

```
#fitting the linear regression model
lm.mod <- lm(critical_temp~., trains)
```

In [65]:

```
summary(lm.mod)
```

Call:
lm(formula = critical_temp ~ ., data = trains)

Residuals:

Min	1Q	Median	3Q	Max
-77.846	-13.251	0.291	12.953	178.801

Coefficients:

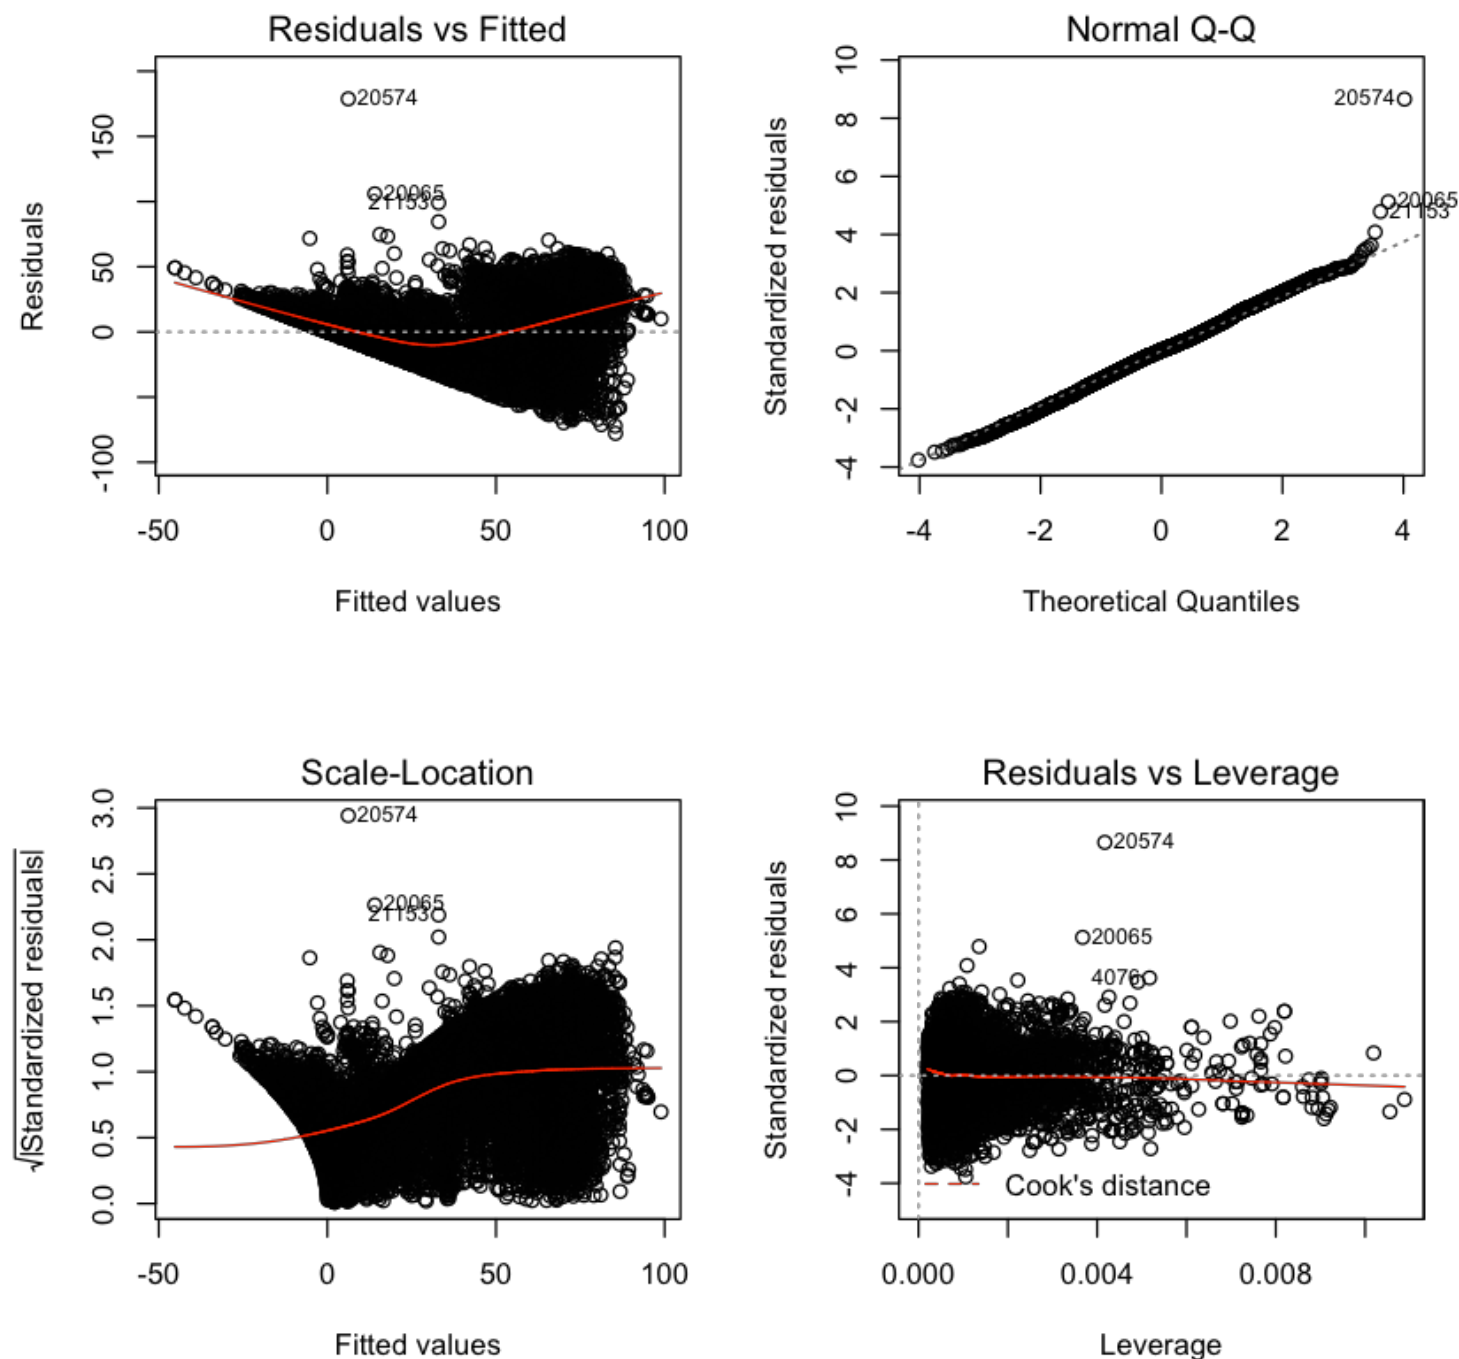
	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.029e+01	1.490e+00	20.331	< 2e-16	***
e_am	-5.771e+01	2.074e+00	-27.820	< 2e-16	***
std_fie	9.925e-02	8.224e-03	12.069	< 2e-16	***
gm_d	-5.458e-04	9.742e-05	-5.603	2.14e-08	***
e_fh	1.434e+00	1.771e+00	0.810	0.41798	
we_fh	1.842e+01	1.515e+00	12.159	< 2e-16	***
std_ar	-1.652e+00	5.701e-02	-28.985	< 2e-16	***
wstd_fie	-1.115e-01	8.586e-03	-12.993	< 2e-16	***
we_v	-1.639e+01	2.099e+00	-7.810	6.07e-15	***
e_v	3.146e+01	2.934e+00	10.725	< 2e-16	***
wstd_ar	2.341e-01	4.325e-02	5.412	6.30e-08	***
number_of_elements	-1.516e+00	5.522e-01	-2.745	0.00606	**
m_v	-5.679e+00	6.268e-01	-9.061	< 2e-16	***
we_ar	-3.605e+01	2.574e+00	-14.008	< 2e-16	***
we_am	5.351e+01	1.771e+00	30.222	< 2e-16	***
wm_v	2.135e+00	5.413e-01	3.944	8.05e-05	***
std_tc	-1.859e-01	1.031e-02	-18.032	< 2e-16	***
r_ar	6.543e-01	2.058e-02	31.796	< 2e-16	***
wstd_tc	3.814e-01	1.040e-02	36.690	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 20.7 on 16992 degrees of freedom
Multiple R-squared: 0.6357, Adjusted R-squared: 0.6353
F-statistic: 1647 on 18 and 16992 DF, p-value: < 2.2e-16

In [66]:

```
par(mfrow=c(2,2))
plot(lm.mod)
```



In [67]:

```
ggplot(data=trainss, aes(lm.mod$residuals)) +
  geom_histogram(binwidth = 1, color = "black", fill = "purple4") +
  theme(panel.background = element_rect(fill = "white"),
        axis.line.x=element_line(),
        axis.line.y=element_line()) +
  ggtitle("Histogram for Model Residuals")
```

Error in ggplot(data = trainss, aes(lm.mod\$residuals)): object 'trainss' not found

Traceback:

```
1. ggplot(data = trainss, aes(lm.mod$residuals))
```

The relationship between the predictors and the target variable is not linearly explainable as the Residuals v/s Fitted plot indicates.

Furthermore, upon inspection of the predictor variables, we can see that std_tc and wstd_tc, two variables which we have identified to have **bimodal** distribution are present in this subset of variables. This can also be a reason why a simple multiple linear regression model cannot explain all of the data.

In [68]:

```
# Predicting the critical_temp of the test predictors
mod.target <- predict(lm.mod, new = test.predictors)

# Computing the MSE
lm.model_mse <- mean((test.target - mod.target)^2)
print(paste('Mean Squared Error of Multiple Linear Regression:', lm.model_mse))

# Computing the R-Squared Value
lm.model_rsv <- cor(test.target, mod.target)^2
print(paste('R-Squared for Multiple Linear Regression:', lm.model_rsv))

[1] "Mean Squared Error of Multiple Linear Regression: 435.4987306
49402"
[1] "R-Squared for Multiple Linear Regression: 0.626257247849855"
```

MODEL 2: MULTIPLE LINEAR REGRESSION WITH WRAPPER METHODS

- Wrapper methods are like a hit-and-a-miss scenario where they check the performance of a model by adding/removing each predictor and evaluating the change it brings into the model.
- There are different approaches to wrapper methods namely: **Sequential Selection and Heuristic Search Algorithms**
- "Wrapper methods use the predictor as a black box and the predictor performance as the objective function to evaluate the variable subset. Since evaluating 2^N subsets becomes a NP-hard problem, suboptimal subsets are found by employing search algorithms which find a subset heuristically." (Chandrashekar et. al, 2013)

In [71]:

```
# Set seed for reproducibility
set.seed(123)
# Set up repeated k-fold cross-validation for 10 folds
train.control <- trainControl(method = "cv", number = 10)
# Train the model
step.model <- train(critical_temp ~., data = train,
                    method = "leapSeq", #This method uses both forwards and backwards selection and eliminates the need for greed
                    tuneGrid = data.frame(nvmax = 1:81),
                    trControl = train.control
                    )
step.model$results
#step.model$bestTune
```

	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
	1	23.68916	0.5227404	18.20678	0.4069329	0.011887677	0.2407592
	2	22.72123	0.5609134	17.63992	0.4265662	0.013243107	0.2858098
	3	21.81559	0.5952298	17.18490	0.3495601	0.009700402	0.2261947
	4	21.36543	0.6117857	16.76538	0.4129202	0.011317030	0.2269454
	5	20.89348	0.6287547	16.45187	0.3617765	0.011126575	0.2164826
	6	20.57948	0.6398353	16.14364	0.3678355	0.010915427	0.2178291
	7	20.33289	0.6484015	15.72183	0.4111112	0.011469591	0.2414392
	8	20.09061	0.6567433	15.47778	0.4279212	0.011746159	0.2497992
	9	19.83752	0.6653455	15.42911	0.3457477	0.008130484	0.2007563
	10	20.16865	0.6514220	15.68044	1.6233801	0.066505469	1.3642993
	11	22.41792	0.5660468	17.58547	3.0654539	0.114512794	2.5262213
	12	21.25030	0.6099809	16.60399	2.8483853	0.108029852	2.3198174
	13	19.82388	0.6638406	15.37016	2.0370101	0.067735780	1.6687459
	14	19.16262	0.6877243	14.80532	0.2736966	0.004956687	0.2463729
	15	19.01752	0.6924112	14.65064	0.3059947	0.006303751	0.2404426
	16	19.44300	0.6765084	15.01911	1.5724466	0.057981638	1.4242728
	17	18.88007	0.6968536	14.51555	0.2904394	0.005724171	0.2236763
	18	18.78607	0.6998470	14.41649	0.2915485	0.005835299	0.2336664
	19	20.60940	0.6342794	16.02621	2.4607281	0.087086668	2.2145209
	20	18.69161	0.7028562	14.35599	0.2738353	0.005626557	0.2313863
	21	18.57924	0.7064556	14.24645	0.3014835	0.005920535	0.2716255
	22	19.37666	0.6774710	14.94110	1.8993880	0.069766237	1.5882238
	23	18.40918	0.7117715	14.09914	0.2548139	0.006345504	0.2059777
	24	18.38581	0.7125252	14.08202	0.2728976	0.006628166	0.2091420
	25	19.08543	0.6881440	14.64717	1.5801252	0.055514261	1.3354431

26	18.63174	0.7039005	14.27983	1.1067421	0.036645866	0.8086115
27	19.21596	0.6843011	14.67815	1.6283279	0.052372671	1.1817011
28	18.24507	0.7169125	13.94784	0.2694993	0.006151948	0.2307212
29	18.21141	0.7179406	13.90853	0.2777830	0.006390650	0.2149014
30	18.43840	0.7106606	14.12119	1.0174455	0.026788844	0.7862241
31	18.64797	0.7031817	14.25729	1.1872418	0.039339938	0.9027445
32	18.09313	0.7215803	13.79492	0.3049385	0.007592059	0.2065819
33	19.01011	0.6914811	14.54367	1.0998648	0.038255199	0.9027605
34	18.04784	0.7229609	13.75418	0.3044595	0.007433787	0.2297705
35	18.26334	0.7158542	13.92888	0.8583813	0.026949655	0.6827564
36	18.21129	0.7175463	13.86347	0.8408911	0.025948468	0.6643917
37	18.33017	0.7135792	13.94022	0.8107829	0.028364157	0.6507498
38	17.90110	0.7274873	13.59181	0.2901231	0.006887490	0.2269319
39	17.89368	0.7277133	13.57576	0.3031369	0.007202308	0.2358013
40	17.86368	0.7286186	13.56269	0.2930121	0.007161437	0.2133945
41	18.05465	0.7222897	13.72239	0.8662924	0.027236845	0.6723537
42	18.30250	0.7142444	13.88876	1.0767502	0.035580775	0.8624708
43	18.04033	0.7228254	13.66925	0.7241600	0.023379375	0.5008489
44	17.80136	0.7305028	13.48966	0.2834799	0.006980234	0.2275015
45	17.77710	0.7312435	13.46677	0.2796734	0.006662921	0.2214543
46	18.15720	0.7188282	13.77894	1.0023664	0.032844118	0.8558430
47	17.72646	0.7327716	13.43971	0.2760439	0.006190935	0.2358137
48	17.70410	0.7334460	13.42231	0.2694166	0.006124425	0.2227117
49	17.84761	0.7288561	13.54338	0.6678392	0.019693658	0.5894378
50	17.69405	0.7337535	13.40881	0.2927363	0.006223073	0.2306924
51	17.80895	0.7300921	13.50491	0.5936678	0.016996158	0.5317792
52	17.84387	0.7289717	13.53064	0.6549630	0.019364024	0.4990973
53	17.82716	0.7295885	13.50328	0.4887244	0.014393049	0.3741216
54	17.63949	0.7353934	13.34324	0.2906538	0.005933659	0.2354656
55	18.05306	0.7225897	13.69892	0.7230776	0.020236011	0.6610805
56	17.93107	0.7263307	13.59783	0.5850526	0.018125085	0.4904180
57	17.73788	0.7323421	13.42153	0.4359578	0.011800055	0.4062169
58	18.02683	0.7233215	13.67513	0.6826092	0.020108730	0.6530140
59	17.86901	0.7282432	13.53504	0.5938014	0.017745486	0.5233816
60	17.86166	0.7285571	13.54700	0.7431412	0.018483893	0.6489782
61	17.59010	0.7368740	13.29717	0.2722493	0.005676918	0.2422476
62	17.79634	0.7306577	13.50951	0.5890504	0.013190131	0.5685358
63	17.68683	0.7339000	13.39955	0.3872490	0.010827912	0.3350161

64	17.57850	0.7372227	13.29196	0.2752096	0.005573310	0.2413526
65	17.58316	0.7370853	13.29592	0.2794671	0.005699298	0.2459209
66	17.63704	0.7354503	13.33618	0.3417941	0.007988569	0.2772516
67	17.65080	0.7350909	13.34174	0.3438630	0.005901890	0.2912669
68	17.68291	0.7340194	13.36900	0.2555359	0.008317623	0.2356220
69	17.65382	0.7349253	13.35224	0.3157337	0.008289990	0.2738759
70	17.63094	0.7356248	13.34773	0.2620633	0.006528547	0.2204127
71	17.66944	0.7345113	13.37430	0.3123297	0.005784910	0.2583230
72	17.62957	0.7356590	13.33717	0.3440330	0.008223652	0.2605901
73	17.67362	0.7343543	13.37414	0.3309930	0.007837515	0.2748610
74	17.72753	0.7327076	13.40194	0.2672438	0.007458989	0.2067941
75	17.56799	0.7375415	13.30167	0.2738096	0.005565643	0.2393507
76	17.72245	0.7328912	13.39675	0.2284759	0.005137378	0.2065644
77	17.61301	0.7361910	13.31895	0.2765239	0.005722663	0.2348059
78	17.61472	0.7361476	13.32542	0.2362516	0.004294750	0.2164531
79	17.63038	0.7356645	13.32622	0.2943008	0.006586147	0.2307190
80	17.61461	0.7361425	13.30967	0.2428038	0.004963498	0.2219284
81	17.56611	0.7376001	13.30393	0.2745499	0.005563751	0.2388539

- From the above output, we can see that we can get a healthy r-squared of 0.709 at n=21, thus let's run the command again with the nvmax set to 21 variables only.
- Furthermore, the rmse is also considerably low at n=21.

In [72]:

```
# Set seed for reproducibility
set.seed(123)
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model
step.model <- train(critical_temp ~., data = train,
                    method = "leapSeq",
                    tuneGrid = data.frame(nvmax = 1:21),
                    trControl = train.control
                    )
step.model$results
```

nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
1	23.68916	0.5227404	18.20678	0.4069329	0.011887677	0.2407592
2	22.72123	0.5609134	17.63992	0.4265662	0.013243107	0.2858098
3	21.81559	0.5952298	17.18490	0.3495601	0.009700402	0.2261947
4	21.36543	0.6117857	16.76538	0.4129202	0.011317030	0.2269454
5	20.89348	0.6287547	16.45187	0.3617765	0.011126575	0.2164826
6	20.57948	0.6398353	16.14364	0.3678355	0.010915427	0.2178291
7	20.33289	0.6484015	15.72183	0.4111112	0.011469591	0.2414392
8	20.09061	0.6567433	15.47778	0.4279212	0.011746159	0.2497992
9	19.83752	0.6653455	15.42911	0.3457477	0.008130484	0.2007563
10	20.16865	0.6514220	15.68044	1.6233801	0.066505469	1.3642993
11	22.41792	0.5660468	17.58547	3.0654539	0.114512794	2.5262213
12	21.25030	0.6099809	16.60399	2.8483853	0.108029852	2.3198174
13	19.82388	0.6638406	15.37016	2.0370101	0.067735780	1.6687459
14	19.16262	0.6877243	14.80532	0.2736966	0.004956687	0.2463729
15	19.01752	0.6924112	14.65064	0.3059947	0.006303751	0.2404426
16	19.44300	0.6765084	15.01911	1.5724466	0.057981638	1.4242728
17	18.88007	0.6968536	14.51555	0.2904394	0.005724171	0.2236763
18	18.78607	0.6998470	14.41649	0.2915485	0.005835299	0.2336664
19	20.60940	0.6342794	16.02621	2.4607281	0.087086668	2.2145209
20	18.69161	0.7028562	14.35599	0.2738353	0.005626557	0.2313863
21	18.57924	0.7064556	14.24645	0.3014835	0.005920535	0.2716255

In [73]:

```
#Extracting the dataset
a <- coef(step.model$finalModel, 21)
a <- data.frame(a)
variables <- as.vector(rownames(a))
variables <- variables[2:21]
variables[21] <- 'critical_temp'
variables
```

```
'e_am' 'r_am' 'wstd_am' 'm_fie' 'r_ar' 'std_ar' 'r_d' 'm_ea' 'wm_ea'
'gm_ea' 'wgm_ea' 'r_ea' 'std_ea' 'wstd_ea' 'r_fh' 'gm_tc' 'we_tc' 'r_tc'
'wstd_tc' 'r_v' 'critical_temp'
```

Now we have extracted the dataset with variables that have been selected using the wrapper methods' **stepwise functionality**.

In [74]:

```
#Extracting the data for only the selected columns
trainstep <- extract_columns(train,variables)
dim(trainstep)
```

17011 21

In [75]:

```
#Fitting the linear model
step.mod <- lm(critical_temp~., trainstep)
```

In [76]:

```
#Summary of the fit
summary(step.mod)
```

Call:

```
lm(formula = critical_temp ~ ., data = trainstep)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-82.943	-10.675	0.404	11.644	174.797

Coefficients:

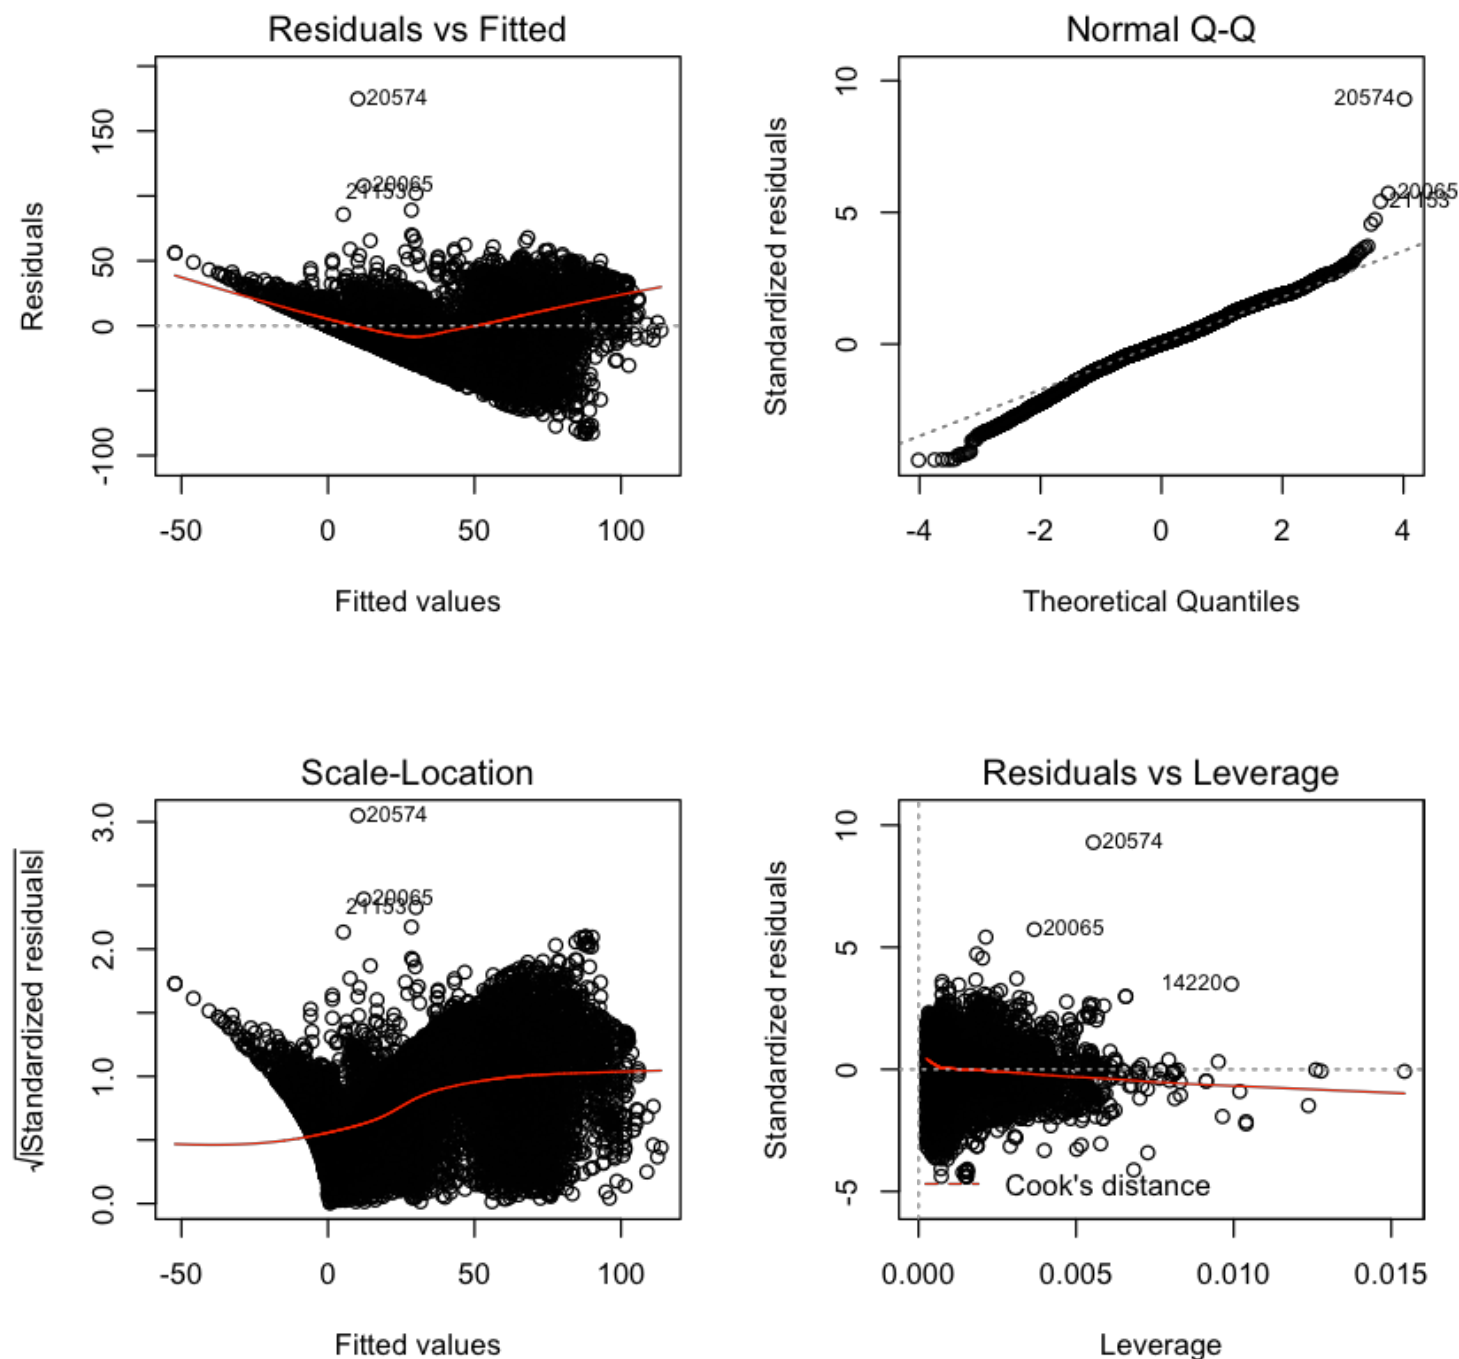
	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-5.600e+00	2.093e+00	-2.676	0.00747	**
e_am	-1.570e+01	1.068e+00	-14.692	< 2e-16	***
r_am	2.375e-01	8.684e-03	27.355	< 2e-16	***
wstd_am	-5.472e-01	2.171e-02	-25.211	< 2e-16	***
m_fie	2.168e-02	2.503e-03	8.662	< 2e-16	***
r_ar	4.951e-01	1.729e-02	28.634	< 2e-16	***
std_ar	-8.494e-01	4.113e-02	-20.651	< 2e-16	***
r_d	-5.414e-04	5.325e-05	-10.167	< 2e-16	***
m_ea	-7.132e-01	4.086e-02	-17.457	< 2e-16	***
wm_ea	1.369e+00	3.978e-02	34.404	< 2e-16	***
gm_ea	7.002e-01	3.687e-02	18.992	< 2e-16	***
wgm_ea	-1.405e+00	3.916e-02	-35.885	< 2e-16	***
r_ea	-2.986e-01	1.470e-02	-20.308	< 2e-16	***
std_ea	1.505e+00	5.037e-02	29.876	< 2e-16	***
wstd_ea	-1.119e+00	3.401e-02	-32.909	< 2e-16	***
r_fh	-1.440e-01	8.890e-03	-16.197	< 2e-16	***
gm_tc	-6.210e-02	6.498e-03	-9.556	< 2e-16	***
we_tc	1.679e+01	8.017e-01	20.940	< 2e-16	***
r_tc	-9.350e-02	3.939e-03	-23.736	< 2e-16	***
wstd_tc	4.560e-01	1.020e-02	44.692	< 2e-16	***
r_v	-2.132e+00	1.535e-01	-13.888	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.85 on 16990 degrees of freedom
Multiple R-squared: 0.698, Adjusted R-squared: 0.6976
F-statistic: 1963 on 20 and 16990 DF, p-value: < 2.2e-16

In [77]:

```
par(mfrow=c(2,2))
plot(step.mod)
```



In [78]:

```
# Predicting the critical_temp of the test predictors
mod.target <- predict(step.mod, new = test.predictors)

# Computing the MSE
lm.mod2_mse <- mean((test.target - mod.target)^2)
print(paste('Mean Squared Error of Multiple Linear Regression:', lm.mod2_mse))

# Computing the R-Squared Value
lm.mod2_rsv <- cor(test.target, mod.target)^2
print(paste('R-Squared for Multiple Linear Regression:', lm.mod2_rsv))
```

```
[1] "Mean Squared Error of Multiple Linear Regression: 364.5878937
75746"
```

```
[1] "R-Squared for Multiple Linear Regression: 0.687142499053063"
```

We can see from the above output that :

1. **The R-squared** value is higher than the previous multiple linear regression model. Which could mean better prediction accuracy
2. **The F-statistic** has also jumped up a notch.

In [79]:

```
cor(trainstep[1:20], trainstep[21])
```

	critical_temp
e_am	0.5435378
r_am	0.4914067
wstd_am	0.3588237
m_fie	0.1010390
r_ar	0.6556092
std_ar	0.5623458
r_d	0.2637002
m_ea	-0.1959834
wm_ea	0.1088088
gm_ea	-0.3811771
wgm_ea	-0.1095284
r_ea	0.2789385
std_ea	0.2616284
wstd_ea	0.3124665
r_fh	-0.1431193
gm_tc	-0.3883755
we_tc	-0.1195392
r_tc	0.6866783
wstd_tc	0.7228292
r_v	-0.1448442

An interesting observation is that: the variables that have been chosen by the step function do not seem to be correlated with the target variable at all. This is something that could be a huge **red flag** during prediction.

Although the **R squared** seemed to have improved, the relationship between the regressors and the target does not seem to be strong and the evidence from the correlation matrix suggests against it.

Thus, we can safely forgo taking this model into consideration and move on to other methods.

MODEL 3 : Extreme Gradient Boost (XGBoost)

In the paper that is the basis of this assignment, the authors seemed to stress on the fact that the XGBoost model was able to give them more prediction accuracy in comparison to the multiple regression model developed by them. Intuitively, I have also decided to explore the boost in accuracy that this model brings to the table.

- Upon a quick research, it has also been found that XGBoost is a gradient boosting library and has some inbuilt **regularisation methods** which are immune to multi-collinearity, a phenomenon which seemed to have a significant affect on the linear models developed before.
- It is also enables with some internal **cross validation** techniques, which can be used during model development processes.

In [80]:

```
#install.packages('xgboost')
```

XGBoost takes in data in the form of a matrix and thus, we will be converting the training and testing dataframes into matrices.

In [82]:

```
#preparing matrices for parsing  
xgbtrain <- xgb.DMatrix(data = as.matrix(train.predictors),label = as.matrix(t  
rain.target))  
xgbtest <- xgb.DMatrix(data = as.matrix(test.predictors),label=as.matrix(test.  
target))
```

In [83]:

```
#Generating default parameters to build the first model on  
params <- list(booster = "gbtree", objective = "reg:linear", eta=0.3, gamma=0,  
max_depth=6, min_child_weight=1, subsample=1, colsample_bytree=  
1)
```

"Using the inbuilt xgb.cv function, let's calculate the best nround for this model. In addition, this function also returns CV error, which is an estimate of test error."

In [84]:

```
xgbcv <- xgb.cv( params = params, data = xgbtrain, nrounds = 100, showsd= T,  
                nfold = 10, print.every.n = 10)
```

Warning message:

"'print.every.n' is deprecated.

Use 'print_every_n' instead.

See help("Deprecated") and help("xgboost-deprecated")."

```
[1]      train-rmse:35.471250+0.093575      test-rmse:35.551220+0.5855  
04  
[11]      train-rmse:10.076469+0.094814      test-rmse:11.478412+0.3104  
85  
[21]      train-rmse:8.736347+0.086450      test-rmse:10.770055+0.3712  
25  
[31]      train-rmse:8.048167+0.074983      test-rmse:10.495155+0.3961  
97  
[41]      train-rmse:7.553851+0.055001      test-rmse:10.317308+0.4191  
18  
[51]      train-rmse:7.147134+0.078783      test-rmse:10.197543+0.3978  
92  
[61]      train-rmse:6.805248+0.073109      test-rmse:10.104826+0.4047  
22  
[71]      train-rmse:6.507882+0.054745      test-rmse:10.029500+0.3969  
94  
[81]      train-rmse:6.255416+0.055524      test-rmse:9.969836+0.40539  
1  
[91]      train-rmse:6.051318+0.053126      test-rmse:9.948029+0.41043  
9  
[100]     train-rmse:5.889784+0.038735      test-rmse:9.926119+0.42014  
9
```

In [85]:

```
#first default - model training  
xgb.mod <- xgb.train(params = params,  
                    data = xgb.DMatrix(data = as.matrix(train.predictors),  
label = as.matrix(train.target)),  
                    nrounds = 3000, nfold = 10, showsd = T, stratified = T  
,  
                    print_every_n = 10,early_stopping_rounds = 100,  
                    watchlist = list(test = xgb.DMatrix(data = as.matrix(t  
est.predictors),label = as.matrix(test.target))))
```



```
[1]      test-rmse:35.276165
Will train until test_rmse hasn't improved in 100 rounds.
```

```
[11]      test-rmse:11.352916
[21]      test-rmse:10.667747
[31]      test-rmse:10.356871
[41]      test-rmse:10.166113
[51]      test-rmse:10.034712
[61]      test-rmse:9.980517
[71]      test-rmse:9.910322
[81]      test-rmse:9.872714
[91]      test-rmse:9.834014
[101]     test-rmse:9.815414
[111]     test-rmse:9.812716
[121]     test-rmse:9.785287
[131]     test-rmse:9.759849
[141]     test-rmse:9.723497
[151]     test-rmse:9.711094
[161]     test-rmse:9.684143
[171]     test-rmse:9.693506
[181]     test-rmse:9.680573
[191]     test-rmse:9.682490
[201]     test-rmse:9.680367
[211]     test-rmse:9.680950
[221]     test-rmse:9.680128
[231]     test-rmse:9.669400
[241]     test-rmse:9.662138
[251]     test-rmse:9.669702
[261]     test-rmse:9.684066
[271]     test-rmse:9.682046
[281]     test-rmse:9.686298
[291]     test-rmse:9.690759
[301]     test-rmse:9.686247
[311]     test-rmse:9.692360
[321]     test-rmse:9.692547
[331]     test-rmse:9.695355
[341]     test-rmse:9.697485
Stopping. Best iteration:
[241]     test-rmse:9.662138
```

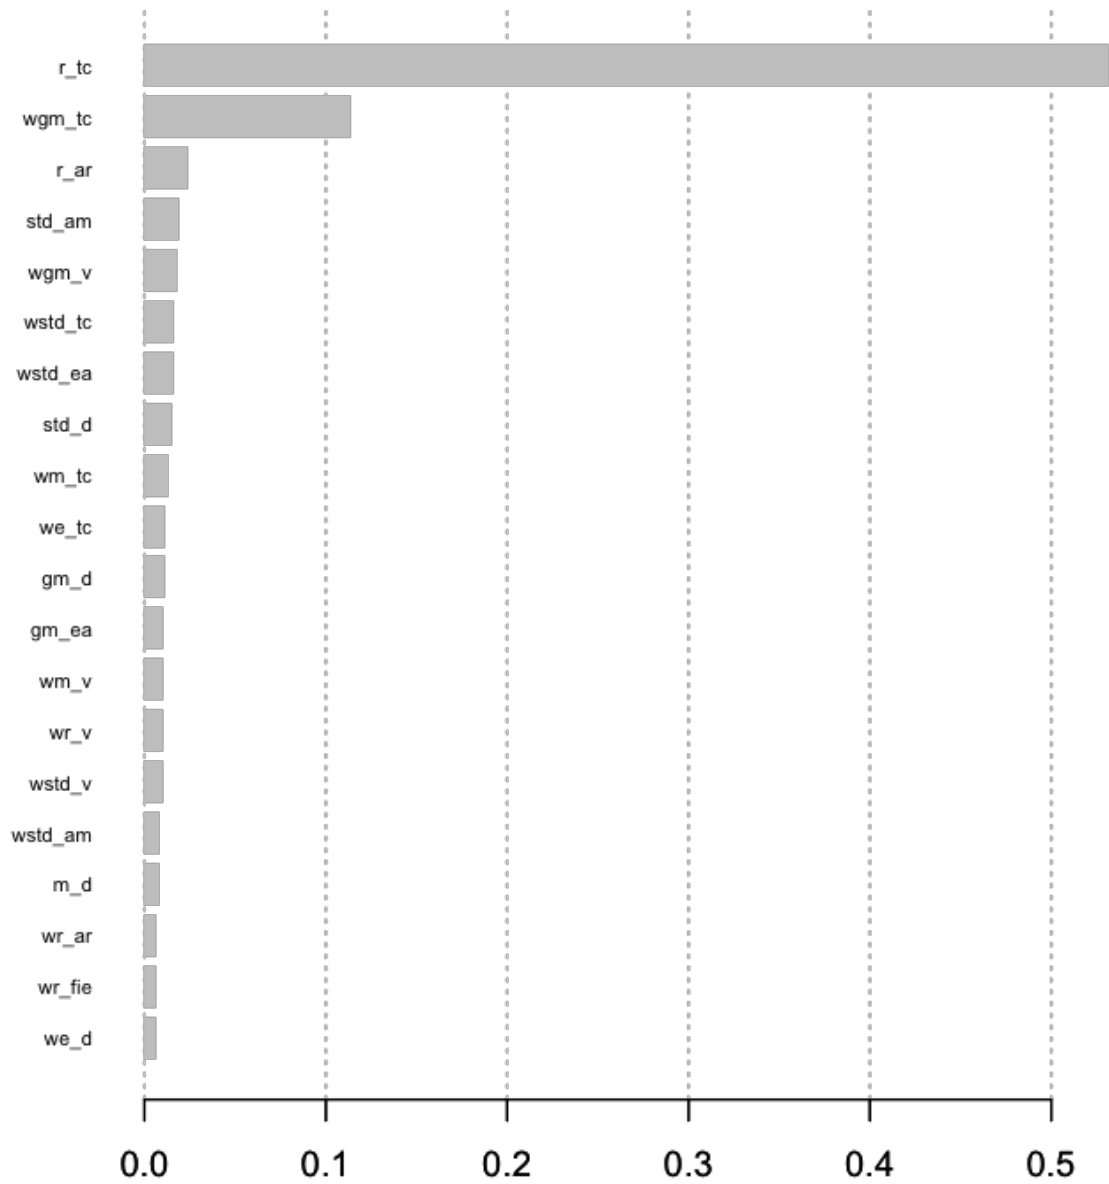
- From this model, let us print out the top 20 variables with the **highest gain**.

In [87]:

```
high.gain <- xgb.importance(colnames(train.predictors), model = xgb.mod)

xgb.plot.importance(importance_matrix = high.gain[1:20])
print(high.gain[1:20])
```

	Feature	Gain	Cover	Frequency
1:	r_tc	0.531166561	0.005682358	0.001995596
2:	wgm_tc	0.113363507	0.010460162	0.012248830
3:	r_ar	0.023875344	0.003304206	0.002614919
4:	std_am	0.019208324	0.009162641	0.008601707
5:	wgm_v	0.017703838	0.018990899	0.009358657
6:	wstd_tc	0.016252701	0.036438644	0.022777319
7:	wstd_ea	0.016223963	0.028113847	0.016652904
8:	std_d	0.014967749	0.007928180	0.005780347
9:	wm_tc	0.013093914	0.019978405	0.016446463
10:	we_tc	0.011455211	0.020896899	0.020919350
11:	gm_d	0.011314336	0.007545330	0.005917974
12:	gm_ea	0.010425610	0.010107176	0.005367465
13:	wm_v	0.010021468	0.029707991	0.013969171
14:	wr_v	0.009878199	0.016949034	0.015001376
15:	wstd_v	0.009873890	0.007009513	0.012042389
16:	wstd_am	0.008527059	0.011123855	0.020850537
17:	m_d	0.007878599	0.008160793	0.006537297
18:	wr_ar	0.006869601	0.017325244	0.018786127
19:	wr_fie	0.006833986	0.027753100	0.020231214
20:	we_d	0.006698233	0.024610887	0.020093587



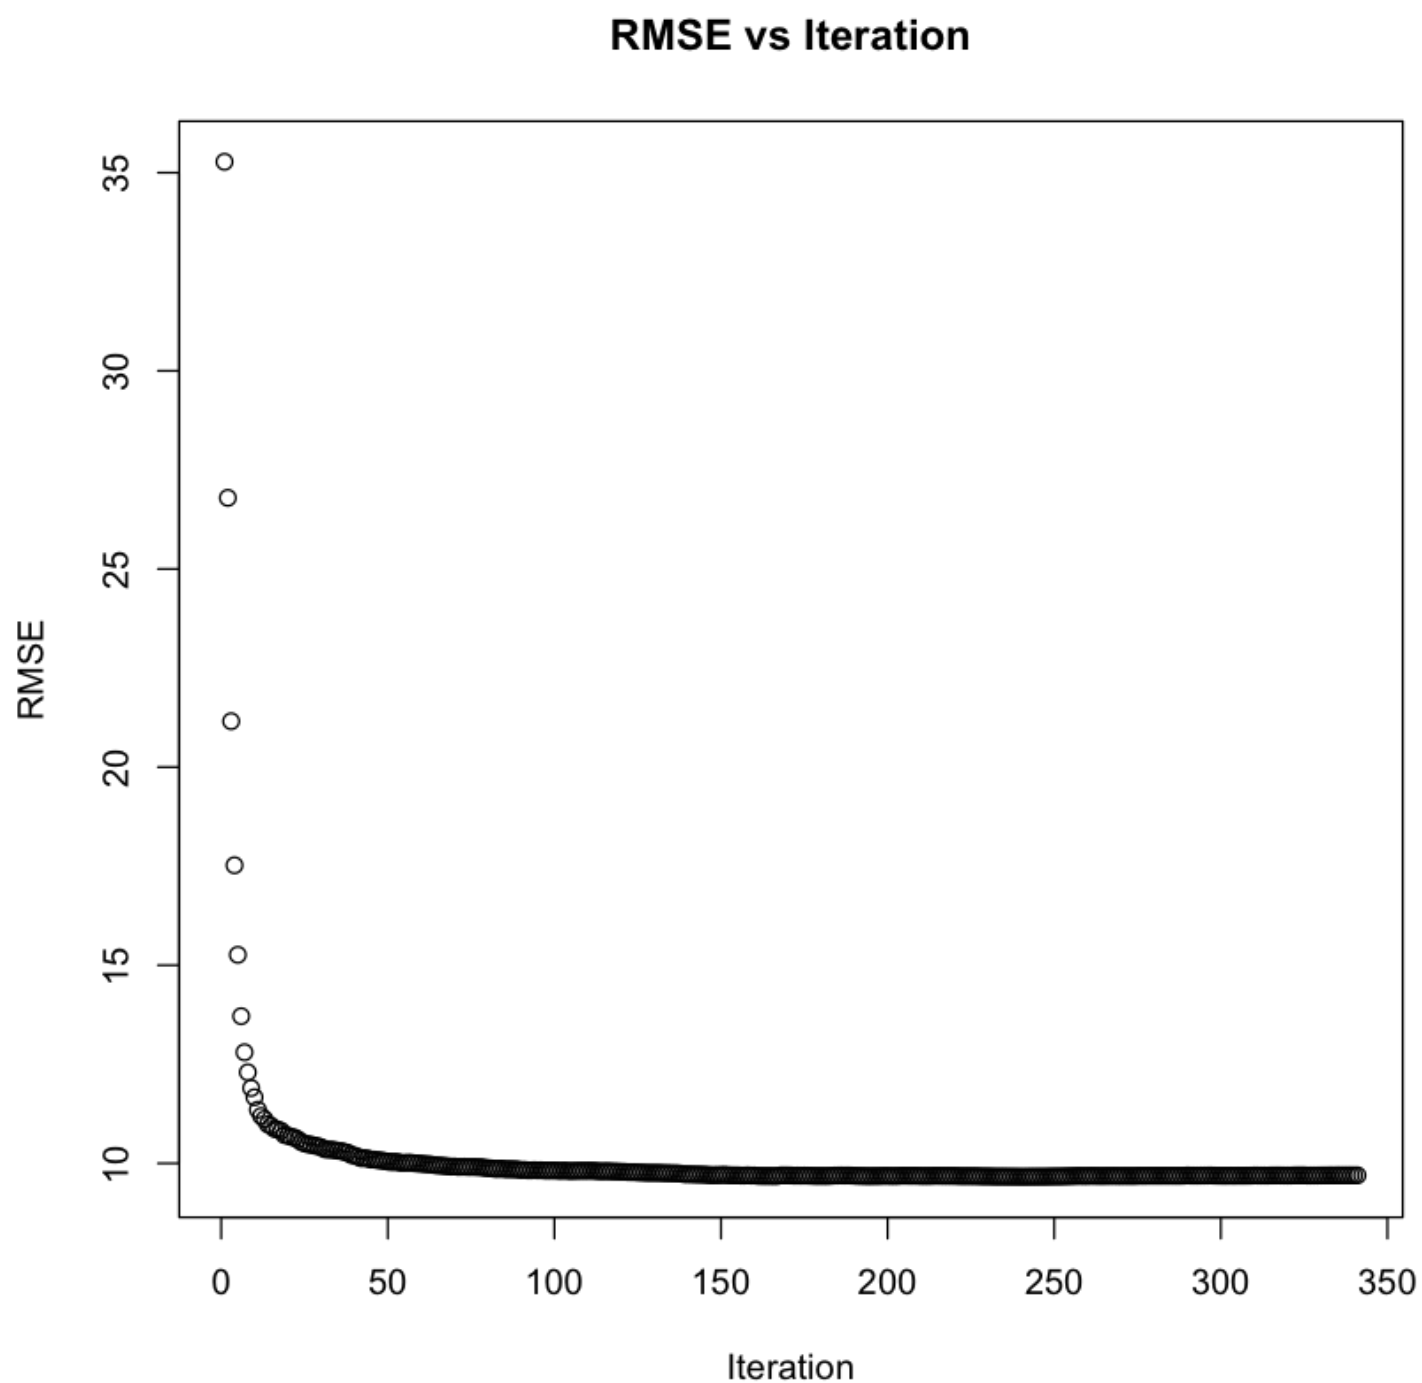
In [88]:

```
xgb.log <- data.frame(xgb.mod$evaluation_log)
head(xgb.log)
```

iter	test_rmse
1	35.27616
2	26.79476
3	21.15707
4	17.52185
5	15.26232
6	13.71331

In [89]:

```
plot(xgb.log, main= 'RMSE vs Iteration',xlab = 'Iteration', ylab = 'RMSE')
```



- It can be observed from the above plot that the RMSE reduces drastically from the first iteration to about the 20th. After this, there is a slow decrease and the RMSE reaches a point after which the change is insignificant.

In [91]:

```
# Test the model on test data
test_pred <- predict(xgb.mod, newdata = xgbtest)

# Minimum RMSE
min_rmse <- min(xgb.log$test_rmse)
print(paste('Lowest RMSE with XGBoost :', xgb.log[xgb.log$test_rmse == min_rmse, 'test_rmse']))

# Checking R-Squared Value
rsq_XGB <- cor(test.target, test_pred)^2
print(paste('R-Squared for XGB :', rsq_XGB))
```

```
[1] "Lowest RMSE with XGBoost : 9.662138"
[1] "R-Squared for XGB : 0.919962613273422"
```

4. Model Comparsion

Broadly, models can be compared by using certain tests of statistics. Or evaluated based on their **R squared** values. For the models that have been generated above, we will be using the Rsquared and RMSE values that have been calculated based on the predictions made by these models. Furthermore, the significance values or p-values of the coefficients can also be assessed in order to compare.

10 fold cross validation for 1st model:

In [92]:

```
# Define train control for k fold cross validation
train_control <- trainControl(method="cv", number=10)
# Fit Linear Regression with critical temp and the 'trains' dataset which contains all the extracted regressors
modell1 <- train(critical_temp~., data=trains, trControl=train_control, method="lm")
# Summarise Results
print(modell1)
```

Linear Regression

17011 samples
18 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 15310, 15310, 15309, 15309, 15311, 15310,
...

Resampling results:

RMSE	Rsquared	MAE
20.70998	0.6352009	16.20647

Tuning parameter 'intercept' was held constant at a value of TRUE

10 fold cross validation for 2nd model:

In [93]:

```
# Define train control for k fold cross validation
train_control <- trainControl(method="cv", number=10)
# Fit Linear Regression with critical temp and the 'trainstep' dataset which contains all the extracted regressors
model2 <- train(critical_temp~., data=trainstep, trControl=train_control, method="lm")
# Summarise Results
print(model2)
```

Linear Regression

17011 samples
20 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 15310, 15309, 15309, 15309, 15311, 15310,
...

Resampling results:

RMSE	Rsquared	MAE
18.85998	0.6974165	14.43609

Tuning parameter 'intercept' was held constant at a value of TRUE

10 fold cross validation for 2nd model:

In [94]:

```
xgbcv <- xgb.cv( params = params, data = xgbtrain, nrounds = 100, showsd= T,  
               nfold = 10, print.every.n = 10)  
print(paste('R-Squared for XGB :',rsq_XGB))
```

Warning message:

“'print.every.n' is deprecated.

Use 'print_every_n' instead.

See help("Deprecated") and help("xgboost-deprecated").”

```
[1]      train-rmse:35.488320+0.058075      test-rmse:35.545168+0.3795  
46  
[11]      train-rmse:10.093161+0.083885      test-rmse:11.534892+0.5059  
18  
[21]      train-rmse:8.765811+0.116767      test-rmse:10.771117+0.4860  
24  
[31]      train-rmse:8.089108+0.147080      test-rmse:10.474384+0.4946  
50  
[41]      train-rmse:7.551393+0.128546      test-rmse:10.282844+0.4876  
25  
[51]      train-rmse:7.141851+0.116996      test-rmse:10.162015+0.5115  
47  
[61]      train-rmse:6.825651+0.108048      test-rmse:10.060993+0.5056  
91  
[71]      train-rmse:6.562148+0.101006      test-rmse:9.994928+0.51549  
6  
[81]      train-rmse:6.309626+0.073465      test-rmse:9.942285+0.51502  
4  
[91]      train-rmse:6.092490+0.064584      test-rmse:9.910442+0.51161  
8  
[100]     train-rmse:5.929964+0.059036      test-rmse:9.867889+0.50787  
4  
[1] "R-Squared for XGB : 0.919962613273422"
```

- After cross validation on all the three models we can see that:
- Model 1 has an R squared of 0.635 and an RMSE of 20.71.
- Model 2 has an R squared of 0.699 and an RMSE of 18.79.
- And Model 3, has the lowest test RMSE and an R squared of 0.91, making it the model that BEST predicts the critical temperature of a semi conductor.

5. Variable Identification and Explanation

NOTE:

Since the models have been developed with highly independent feature selection mechanisms for each model respectively, I believe that the reason to explain the identified variables doesn't exist in this case. Meaning, step by step feature selection was NOT DONE to develop and enhance a single model. Each model has its own feature selection methods involved.

Furthermore, reasoning for feature selection/removal has been provided in the model development stage for every model and thus, this section becomes redundant.

6. Conclusion

"Among the machine learning methods used in practice, gradient tree boosting is one technique that shines in many applications. Tree boosting has been shown to give state-of-the-art results on many standard classification benchmarks" (Li, 2010)

As we have seen above, the RMSE and out-of-sample R squared that have been obtained by implementing the XGBoost method have had phenomenal in comparison to the stepwise linear regression models. Furthermore, the XGBoost model has also eliminated the effect of multicollinearity that has been found in the predictors which has had a drastic affect on the previously developed models, thereby restricting their R squared values.

Thus, XGBoost has yet again proved to be one of the best Machine Learning model building methods.

7. References

1. Hamidieh, K. (2018), "A data-driven statistical model for predicting the critical temperature of a superconductor".
2. Chris Chatfield (1986), European Journal of Operational Research 23 (1986) 5-13 North-Holland.
3. Reitermanova, Z. (2010), "Data splitting", WDS's 10 Proceedings of Contributed Papers, Part, pp. 31-36.
4. Härdle W.K., Simar L. (2012), Regression Models. In: Applied Multivariate Statistical Analysis. Springer, Berlin, Heidelberg
5. Max Kuhn (2008), "Building Predictive Models in R Using the caret Package", Journal of Statistical Software, November 2008, Volume 28, Issue 5.
6. M. Karagiannopoulos, D. Anyfantis, S. B. Kotsiantis, and P. E. Pintelas, "Feature Selection for Regression Problems" [http://www.math.upatras.gr/ \(http://www.math.upatras.gr/~dany/Downloads/hercma07.pdf\)](http://www.math.upatras.gr/~dany/Downloads/hercma07.pdf). (last accessed: 2015-Oct-22).
7. Beginners Tutorial on XGBoost and Parameter Tuning in R, <https://www.hackerearth.com/practice/machine-learning/machine-learning-algorithms/beginners-tutorial-on-xgboost-parameter-tuning-r/tutorial/> (<https://www.hackerearth.com/practice/machine-learning/machine-learning-algorithms/beginners-tutorial-on-xgboost-parameter-tuning-r/tutorial/>)
8. Chandrashekar G., Sahin F. (2013), "A survey on feature selection methods", www.elsevier.com/locate/compeleceng.
9. Li, P. (2010). Robust Logitboost and adaptive base class (ABC) Logitboost. In Proceedings of the Twenty-Sixth Conference Annual Conference on Uncertainty in Artificial Intelligence (UAI'10), pages 302–311.

In []:

In []: